



Machine Learning for Pixel and Object Classification

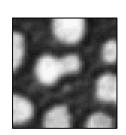
Robert Haase

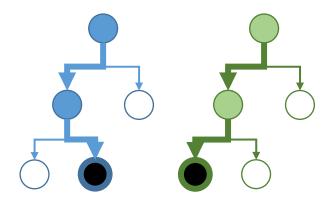
With Material from
Deborah Schmidt, Jug Lab, MPI CBG
Uwe Schmidt, Myers Lab, MPI CBG
Martin Weigert, EPFL
Ignacio Arganda-Carreras, Universidad del Pais Vasco
Carsen Stringer, HHMI Janelia
Wei Ouyang, KTH Royal Institute of Technology, Stockholm and
The Scikit-Learn community



Today

- Machine learning for Pixel and Object Classification
 - Random Forest Classifiers
 - Algorithm evaluation
 - Outlook: Deep learning
- Fiji
 - Trainable Weka Segmentation
 - Labkit
 - CLIJ object classifier
- Python
 - scikit-learn / napari





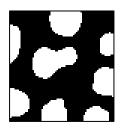


Image segmentation using thresholding



Recap: Finding the right workflow towards a good segmentation takes time

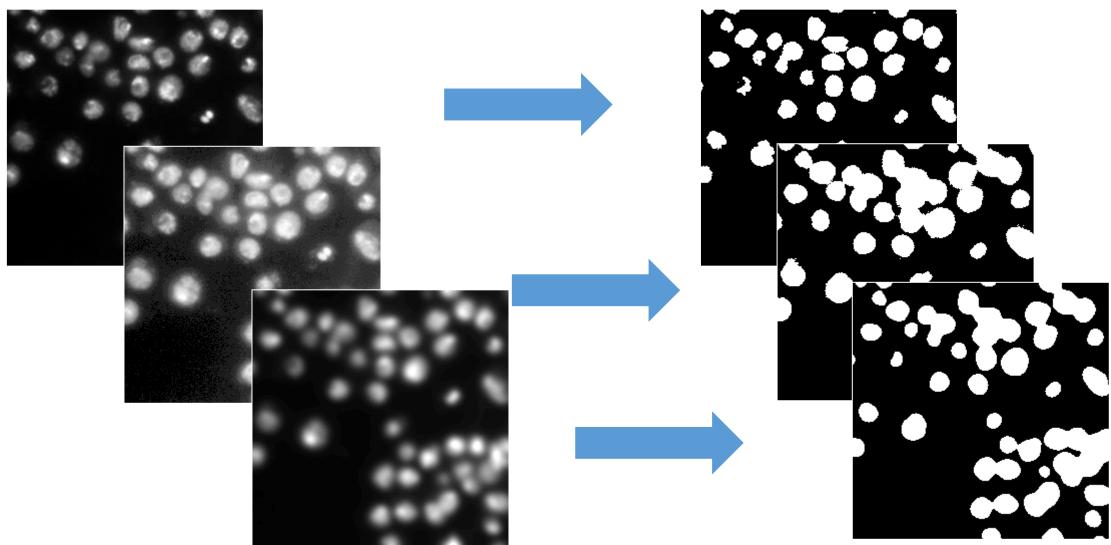


Image segmentation using thresholding



Recap: Combining images, e.g. using Difference of Gaussian (DoG)

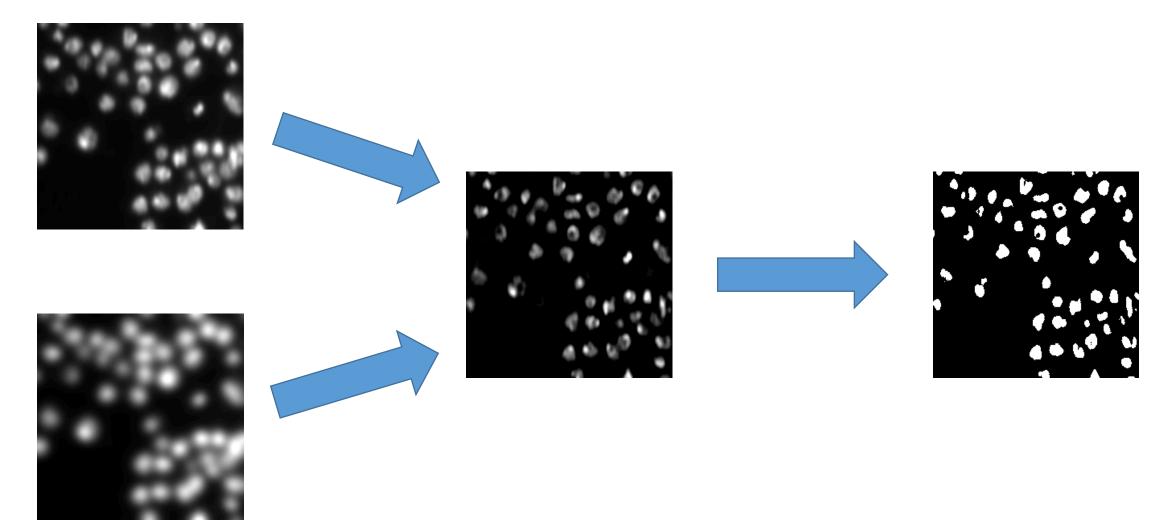
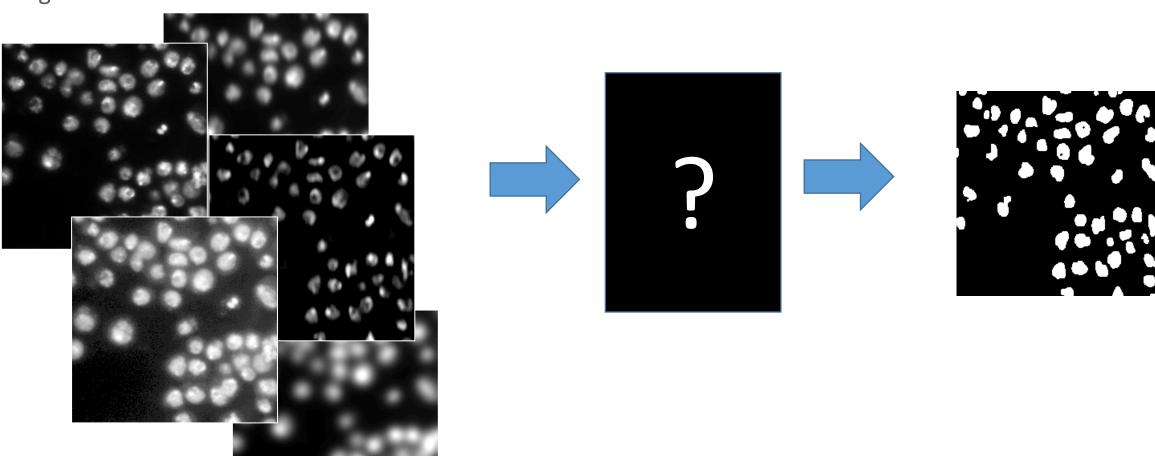


Image segmentation using thresholding



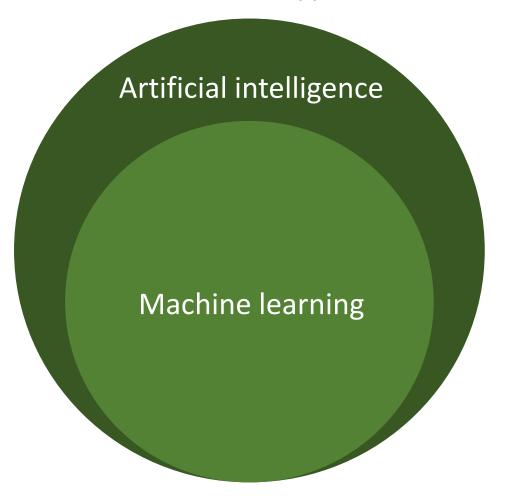
 Might there be a technology for optimization which combination of images can be used to get the best segmentation result?

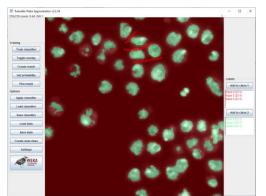


Machine learning

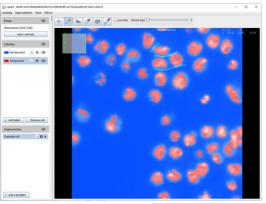
PoL
Physics of Life
TU Dresden

- A research field in computer science
- Finds more and more applications, also in life sciences.



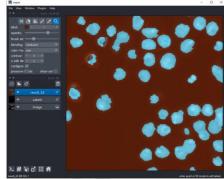


Trainable Weka Segmentation https://imagej.net/plugins/tws/



LabKit
https://imagej.net/
plugins/labkit/

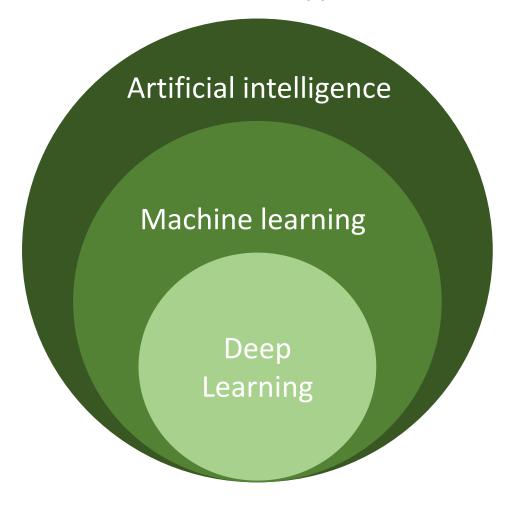
Python / scikit-learn / napari / apoc

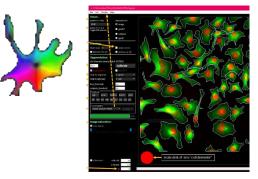


Machine learning

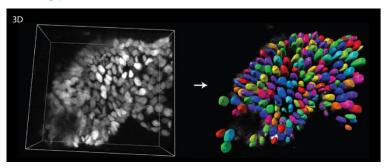
PoL
Physics of Life
TU Dresden

- A research field in computer science
- Finds more and more applications, also in life sciences.

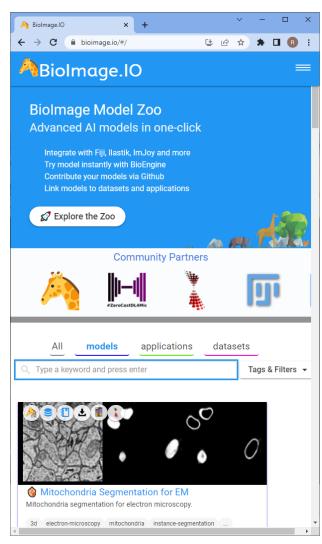








https://github.com/stardist/stardist



https://bioimage.io/



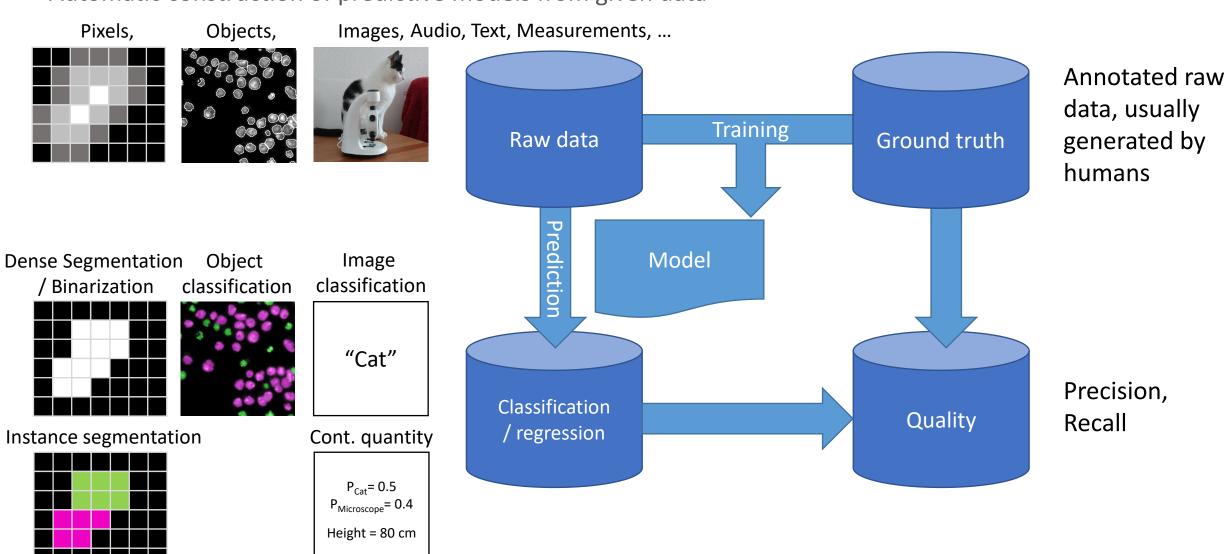
Logos and screenshots are taken from the github repositories / websites provided under BSD and MIT licenses.

Machine learning

@haesleinhuepf



Automatic construction of predictive models from given data

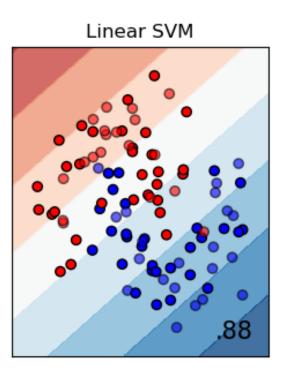


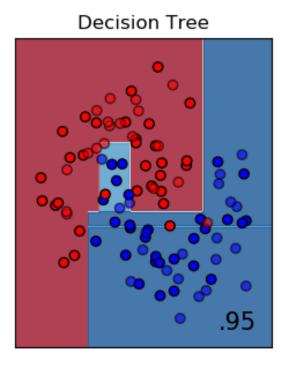
Bio-Image Analysis May 2022

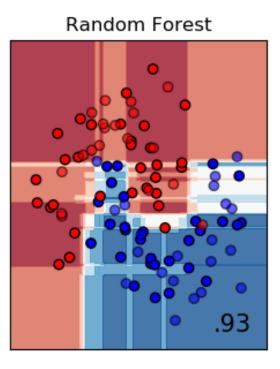


• Guess classification (color) from position of a sample in parameter space.

Input data



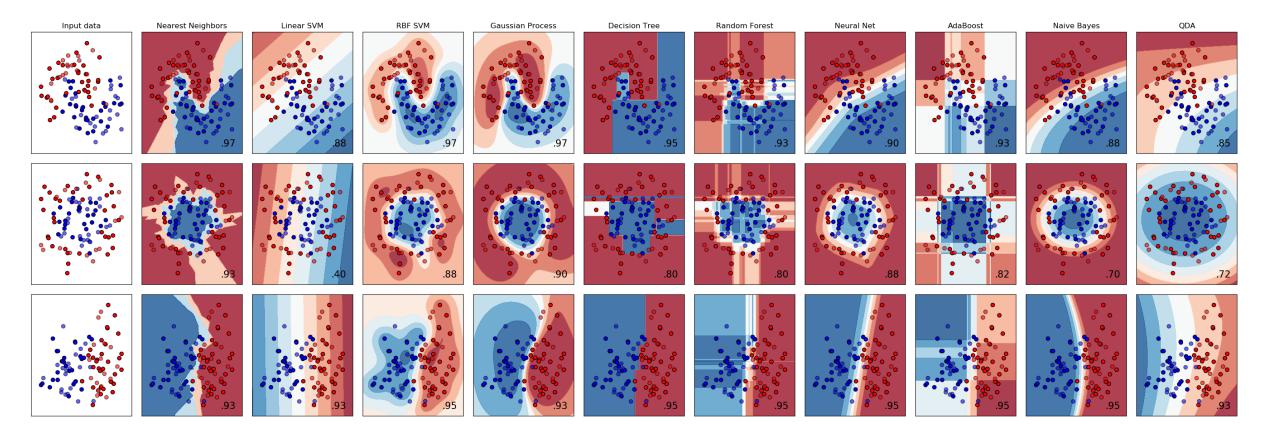




Approaches



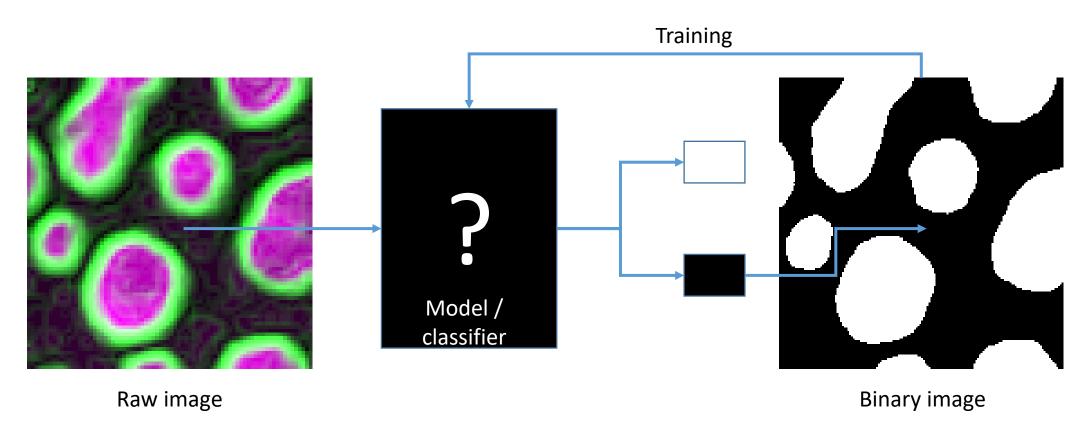
• The right approach depends on data, computational resources and desired quality



Machine learning for image segmentation



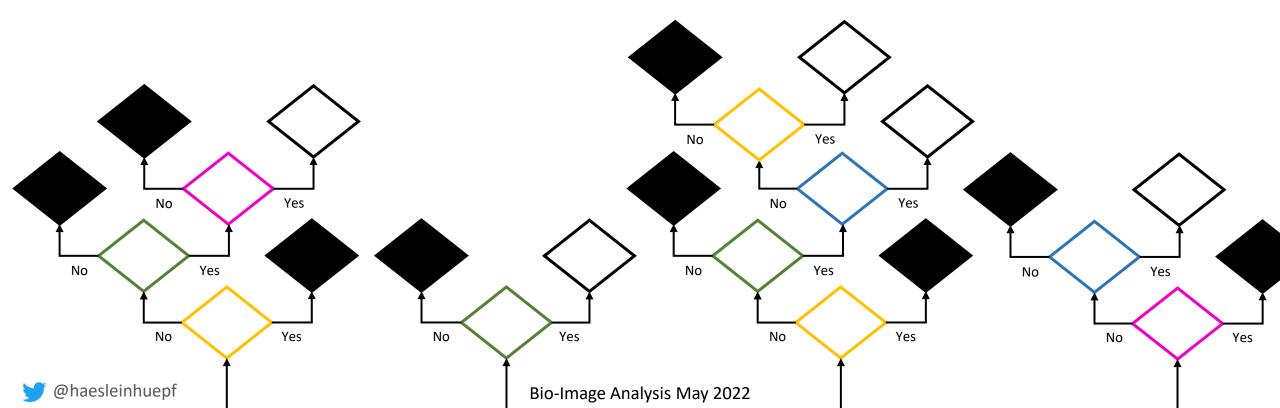
- Supervised machine learning: We give the computer some ground truth to learn from
- The computer derives a *model* or a *classifier* which can judge if a pixel should be foreground (white) or background (black)
- Example: Binary classifier



Random forest based image segmentation



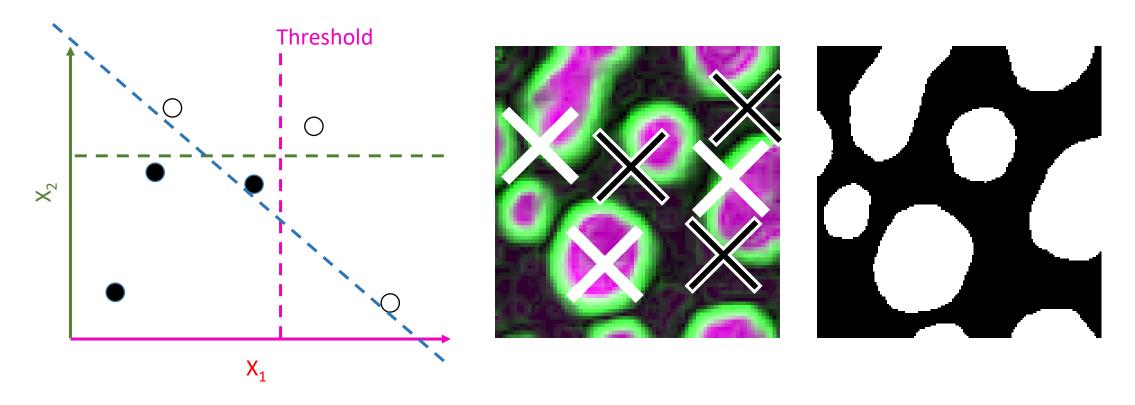
- Decision trees are classifiers, they decide if a pixel should be white or black
- Random decision trees are randomly initialized, afterwards evaluated and selected
- Random forests consist of many random decision trees
- Example: Random forest of binary decision trees



Deriving random decision trees



- For efficient processing, we randomly *sample* our data set
 - Individual pixels, their intensity and their classification

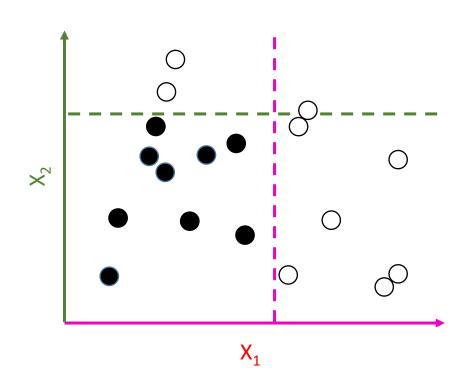


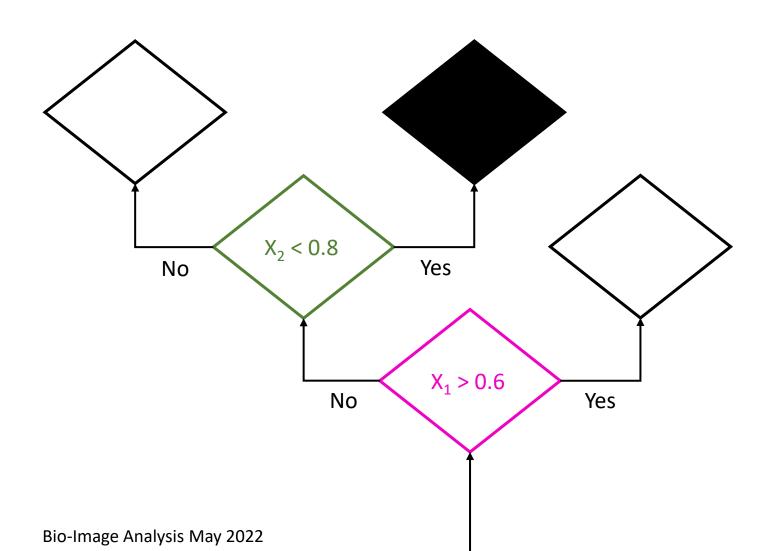
Note: You cannot use a single threshold to make the decision correctly

Deriving random decision trees



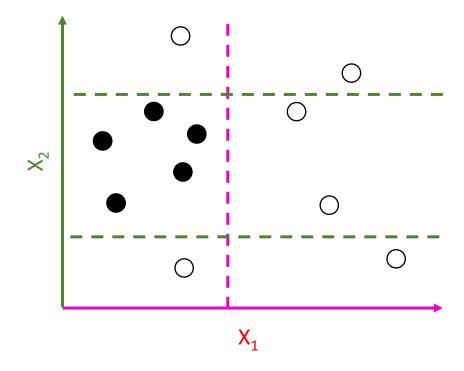
 Decision trees combine several thresholds on several parameters

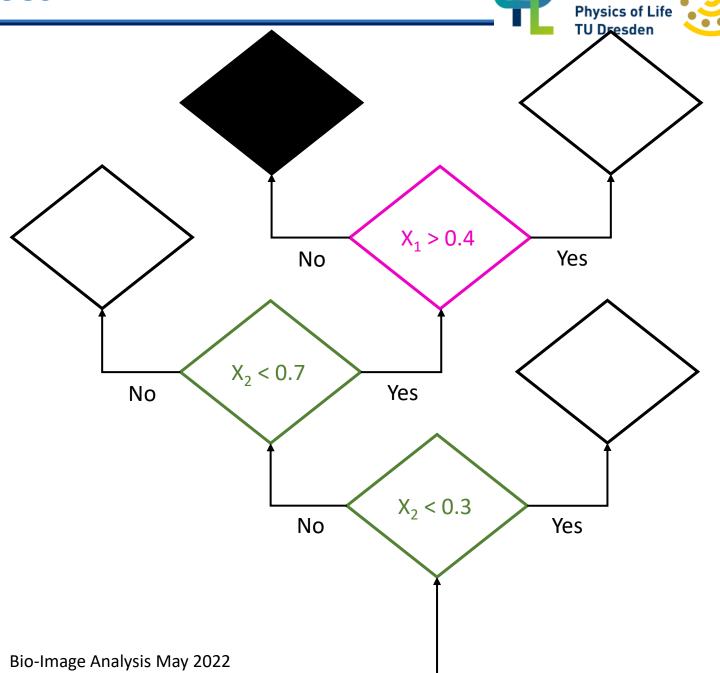




Deriving random decision trees

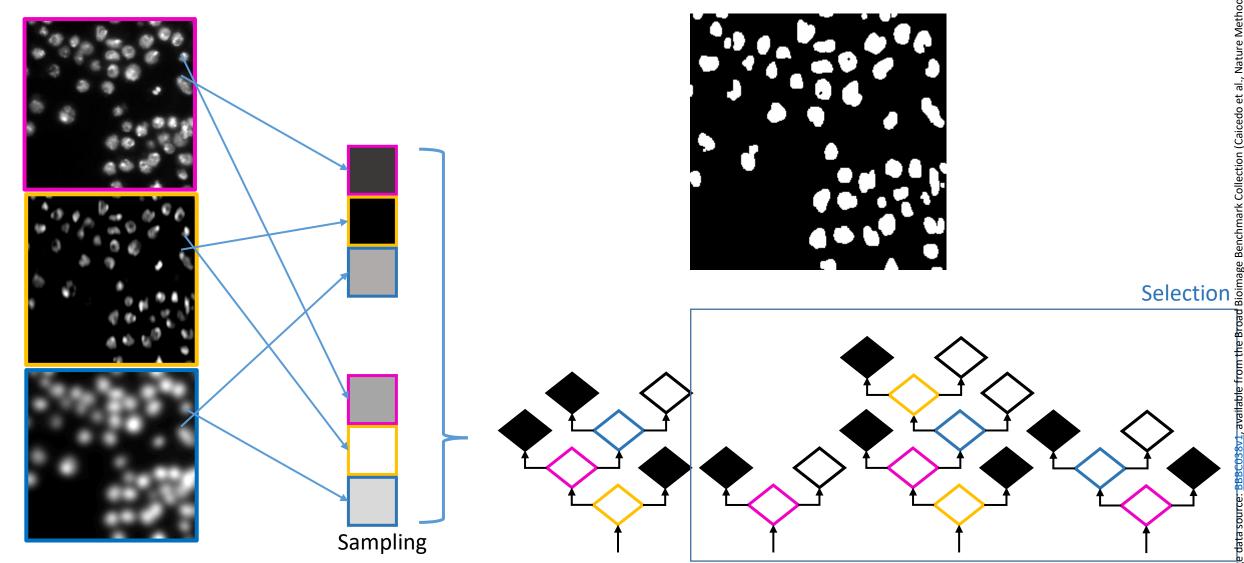
Depending on sampling, the decision trees are different





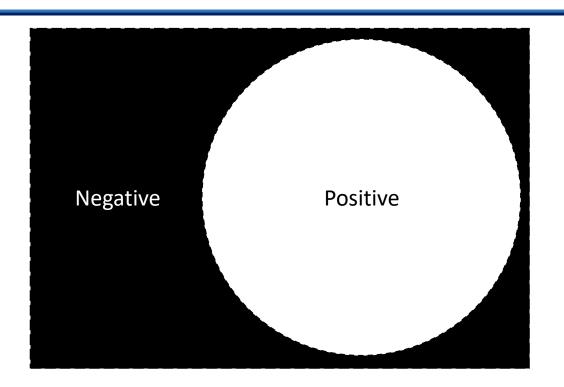


By comparing performance of individual decision trees, good ones can be selected, bad ones excluded.



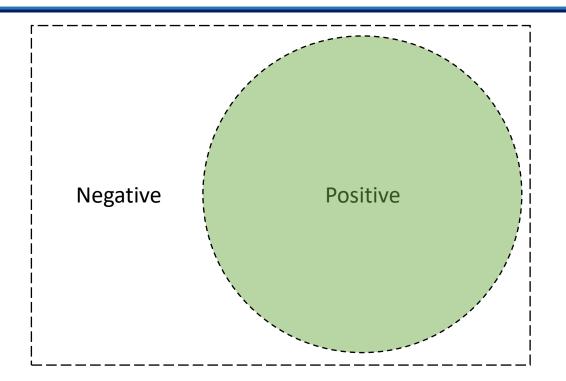


- In general
 - Define what's positive and what's negative.



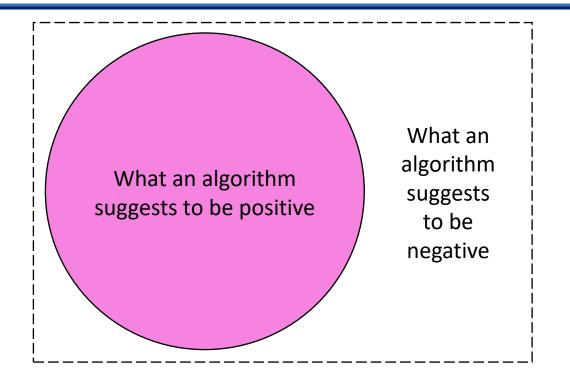


- In general
 - Define what's positive and what's negative.



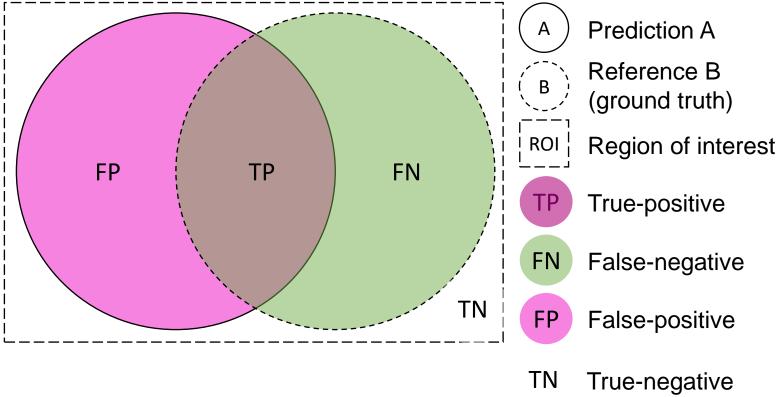


- In general
 - Define what's positive and what's negative.



Pol Physics of Life TU Dresden

- In general
 - Define what's positive and what's negative.
 - Compare with a reference to figure out what was true and false
 - Welcome to the Theory of Sets



Precision
$$\frac{TP}{TP + FP}$$

What fraction of points that were predicted as positives were really positive?

$$\frac{TP}{TP + FN}$$

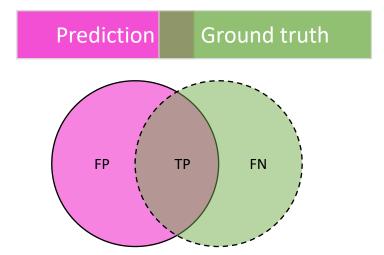
What fraction of positives points were predicted as positives?

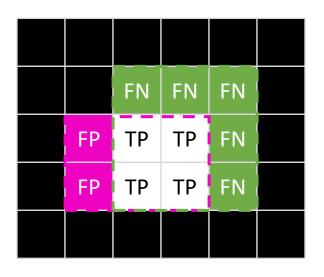
Pixel-wise versus Object-wise evaluation













True-positive: 4

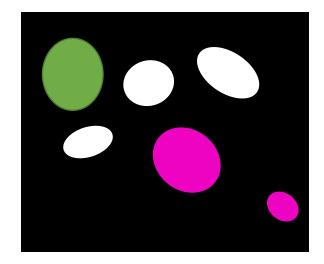
False-negative: 5

False-positive: 2

Precision: 4/6 = 66%

Recall: 4/9 = 44%

Object wise: Detection quality



True-positive: 3

False-negative: 1

False-positive: 2

Precision: 3/4 = 75%

Recall: 3/5 = 60%

Precision

TPTP + FP

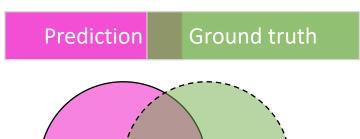
Recall (a.k.a. sensitivity)

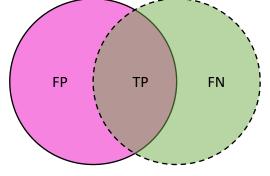
TPTP + FN

Pixel-wise versus Object-wise evaluation





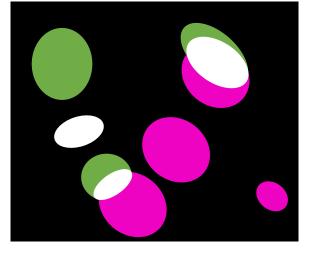




$$\frac{-TP}{TP + FN + FP}$$

$$\frac{TP}{TP + FP}$$

$$\frac{TP}{TP + FN}$$



Objects with at least 50% pixel-wise overlap between P and GT

True positive: 2

False positives: 3

False negatives: 2

Precision: 2/5 = 40%

Recall: 2/4 = 50%

Bio-Image Analysis May 2022

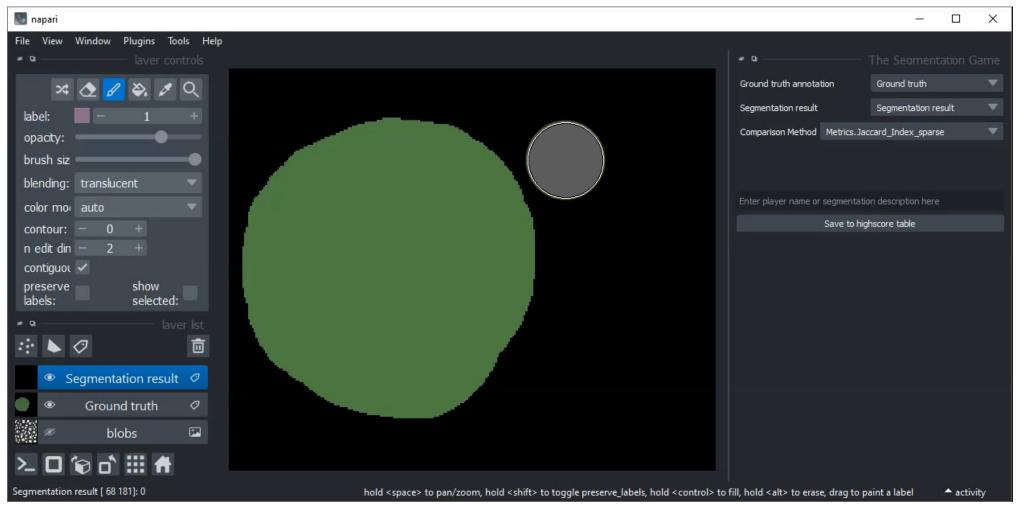
Average Overlap for all ground-truth objects:

$$(0+1+0.5+0.2)/4$$

Pixel-wise versus Object-wise evaluation



- Average Overlap for all ground-truth objects
- https://github.com/haesleinhuepf/the-segmentation-game





Voxel-wise Youden-Index

$$YI = p_{TP} + p_{TN} - 1$$

Volume error

$$\Delta_V = V_A - V_B$$

$$\delta_V = \frac{\Delta_V}{V_B}$$

Dice Index

$$DI(A,B) = \frac{2|A \cap B|}{|A| + |B|}$$

Jaccard Index

$$JI(A,B) = \frac{|A \cap B|}{|A \cup B|} = \frac{DI}{2 - DI}$$

Contour distance

$$d_{e,min}(a, B) = \min(d_e(a, b)|b \in B)$$
$$\bar{d}_c(A, B) = \frac{\sum_{\forall a \in C(A)} d_{e,min}(a, C(B))}{|C(A)|}$$
$$\bar{d}_{bil,c}(A, B) = \frac{\bar{d}_c(A, B) + \bar{d}_c(B, A)}{2}$$

Hausdorff distance

$$d_H(A, B) = \max(d_{e,min}(a, B) | a \in A)$$
$$d_{bil,H}(A, B) = \max(d_H(A, B), d_H(B, A))$$

• Simplified Hausdorff distance $d_H(A, B) = \max(d_{e,min}(a, C(B)) | a \in C(A))$

Volume standard deviation

$$\delta_{\bar{V}} = 2 \frac{|V_A - V_B|}{|V_A + V_B|}$$

Classification error

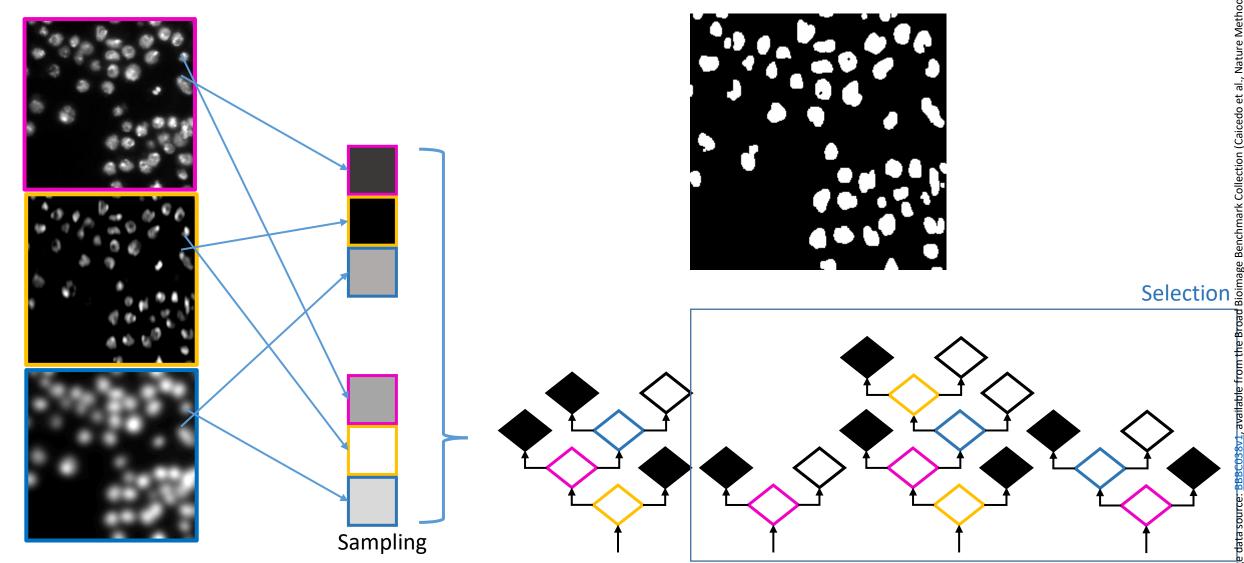
$$e_{Class} = \frac{H}{|TP| + |FN|}$$

Hamming distance

$$d_h = |A \cup B| - |A \cap B|$$
$$= |FP| + |FN|$$

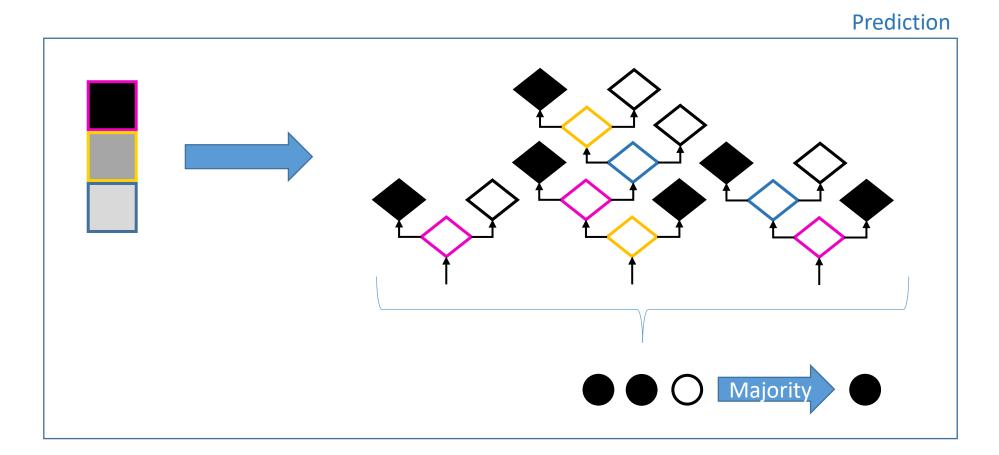


By comparing performance of individual decision trees, good ones can be selected, bad ones excluded.





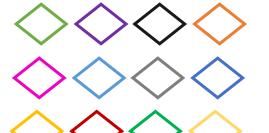
Combination of individual tree decisions by voting or max / mean





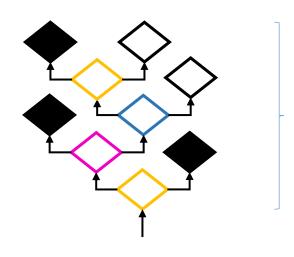
Typical numbers for pixel classifiers in microscopy

Available features: > 20



- Gaussian blur image
- DoG image
- LoG image
- Hessian

Selected features: <= depth



Depth: 4

Number of trees: > 100

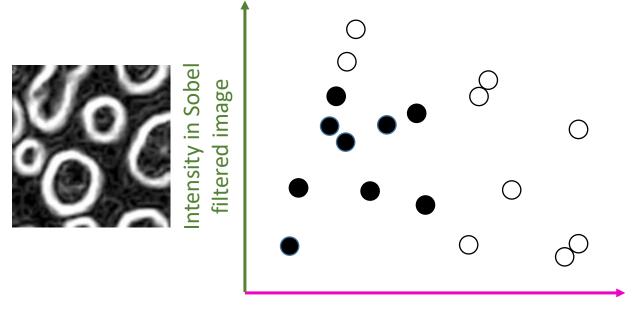


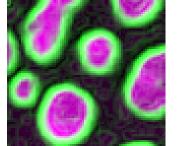
Object classification



What if we exchange pixel features with object features?

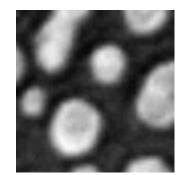
Pixel classification



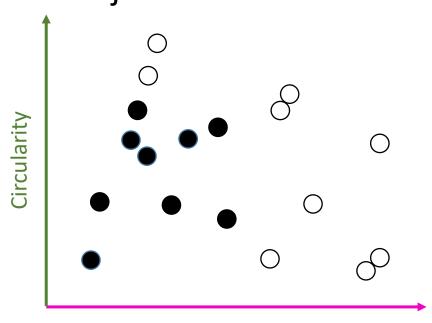


@haesleinhuepf

Intensity



Object classification



Aspect ratio

- The algorithms work the same but with different
 - Features
 - Number of features
 - Tree / forest parameters
 - Selection criteria

Model validation



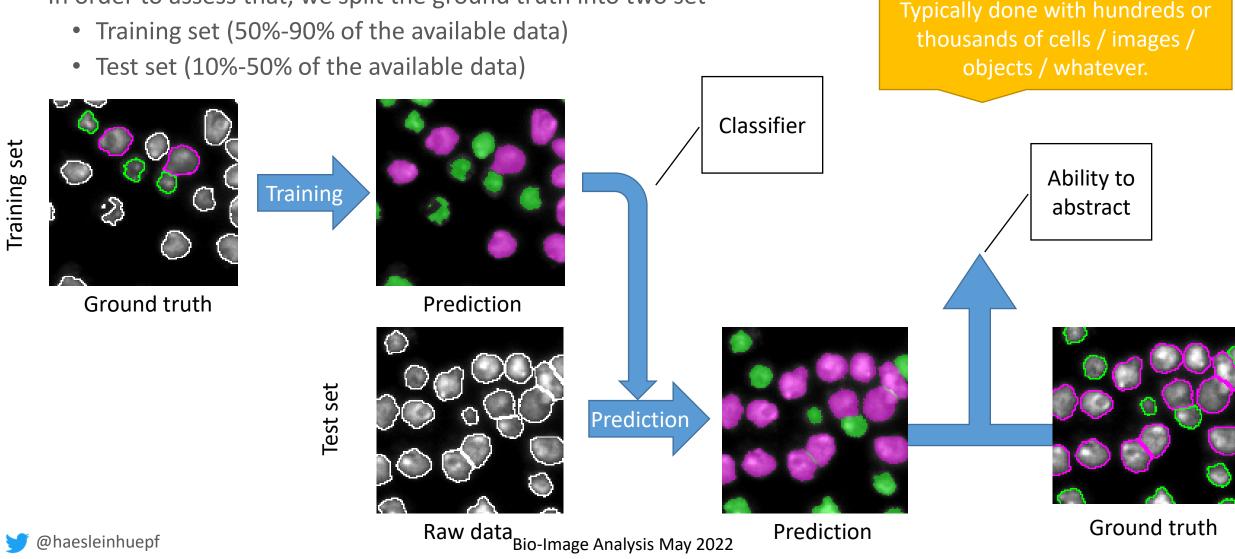
- Underfitting
 - A trained model that is not even able to properly process the data it was trained on
- Overfitting
 - A model that is able to process data it was trained on well
 - It processes other data poorly

Model validation



• A good classifier is trained on a hand full of datasets and works on thousands similarly well.

• In order to assess that, we split the ground truth into two set



@haesleinhuepf





Pixel and object classification in Python using scikit-learn

Robert Haase



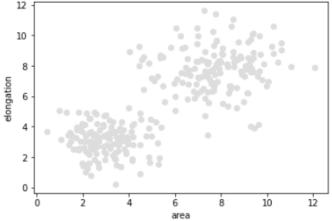
Tabular object classification



Classify objects starting from feature vectors (table columns)

Raw data

	area	elongation
0	3.950088	2.848643
1	4.955912	3.390093
2	7.469852	5.575289
3	2.544467	3.017479
4	3.465662	1.463756
5	3.156507	3.232181
6	9.978705	6.676372
7	6.001683	5.047063
8	2.457139	3.416050
9	3.672295	3.407462
10	9.413702	7.598608

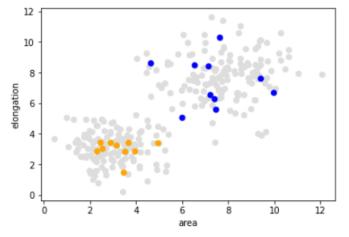


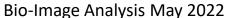
"Ground truth" annotation

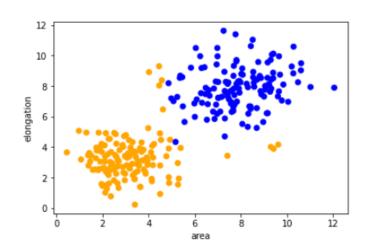
Classifier prediction

Classifier training

result = classifier.predict(validation_data)









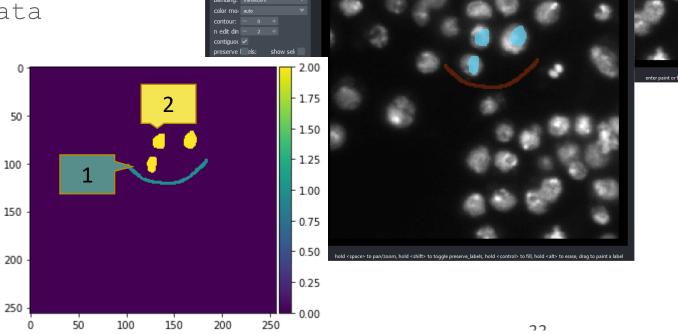


Prepare an empty layer for annotations and keep a reference

```
labels = viewer.add_labels(
    np.zeros(image.shape).astype(int))
```

Read annotations

manual annotations = labels.data





- Pixel classification using scikit-learn
 - Expects one-dimensional arrays for

annotation

- every feature individually
- ground truth

train classifier

prediction



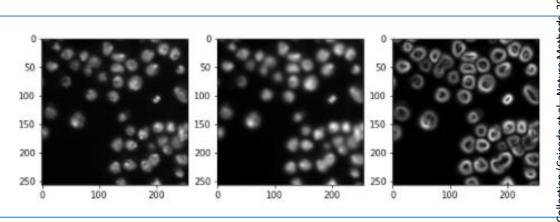
- Pixel classification using scikit-learn
 - Expects one-dimensional arrays for
 - every feature individually
 - ground truth

```
# for training, we need to generate features
feature_stack = generate_feature_stack(image)

X, y = format data(feature stack, manual annotations)
```

train classifier

```
from sklearn.ensemble import RandomForestClassifier
classifier = RandomForestClassifier(max_depth=2, random_state=0)
classifier.fit(X, y)
```

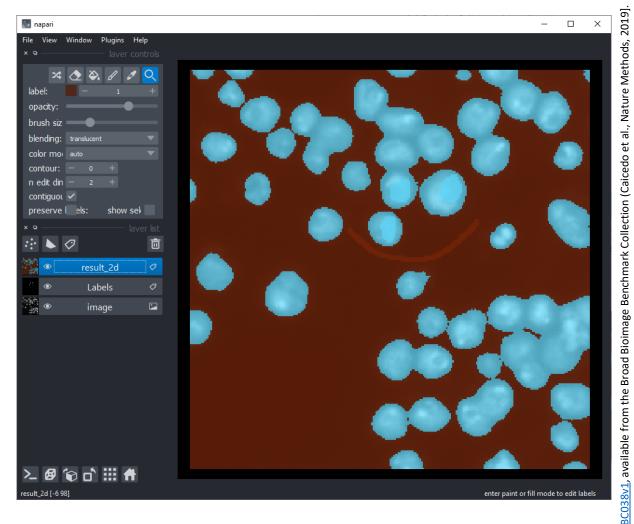




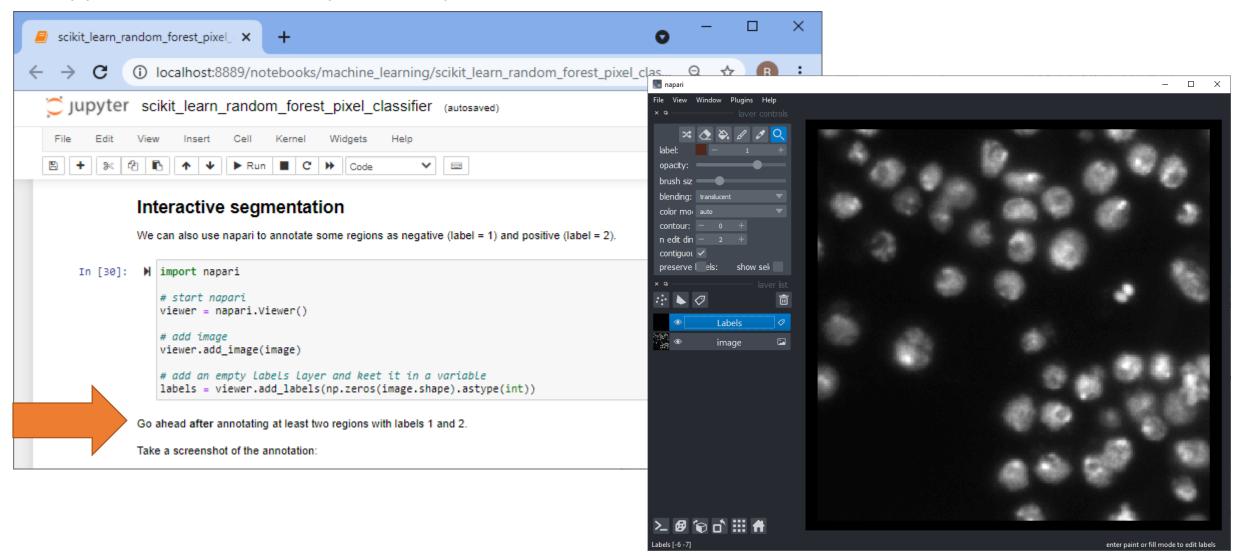


Pixel classification using scikit-learn

```
# process the whole image and show result
result_1d = classifier.predict(feature_stack.T)
result_2d = result_1d.reshape(image.shape)
viewer.add_labels(result_2d)
```



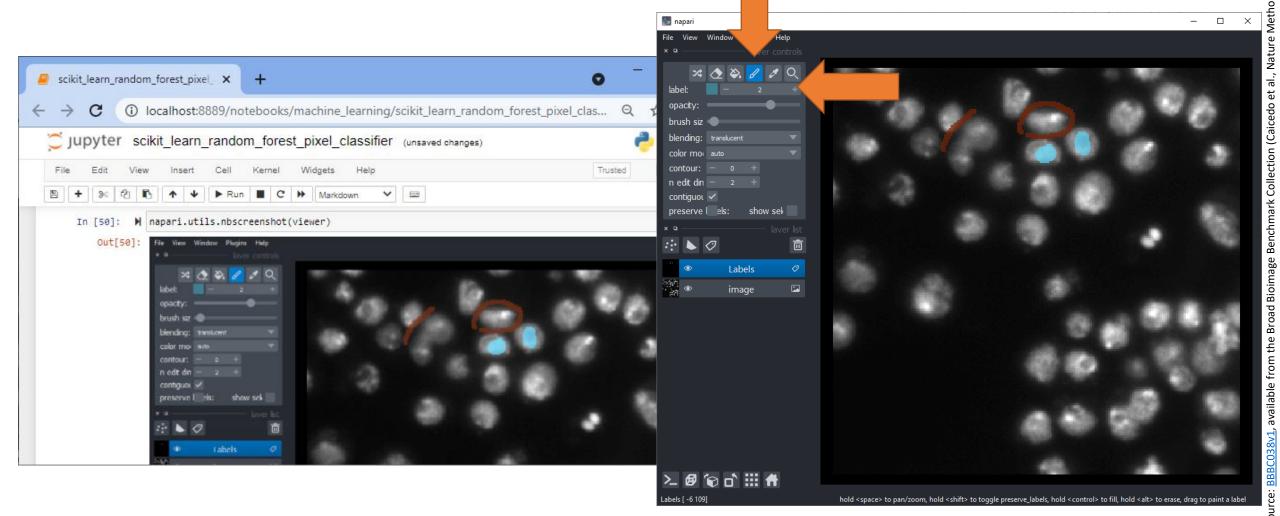
Jupyter notebooks and napari side-by-side



Interactive pixel classification



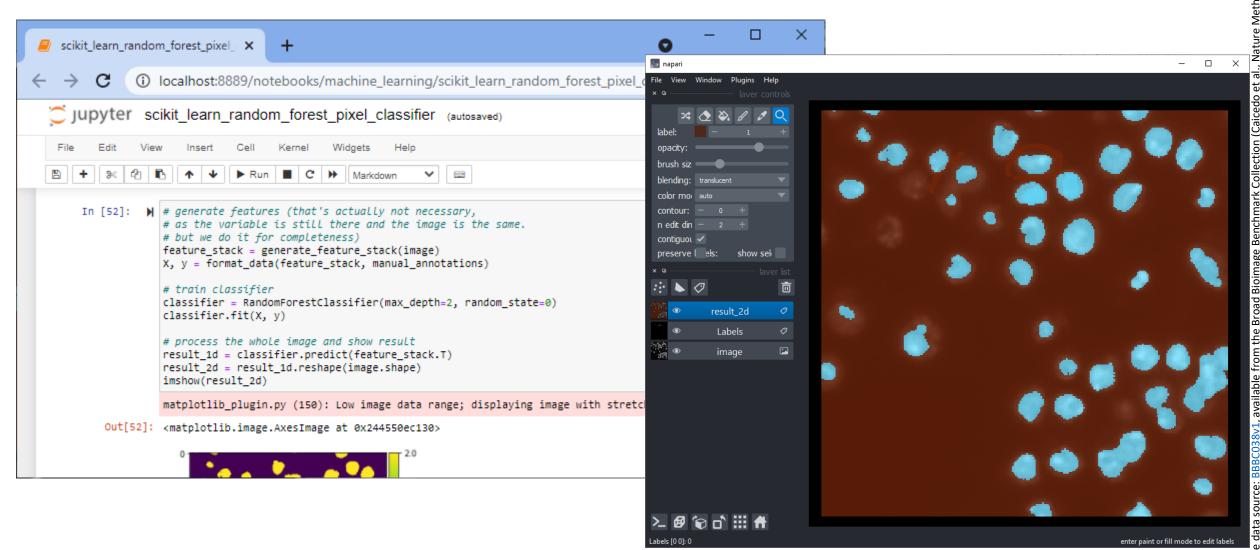
• Jupyter notebooks and napari side-by-side



Interactive pixel classification



Jupyter notebooks and napari side-by-side







Accelerated pixel and object classification (APOC)

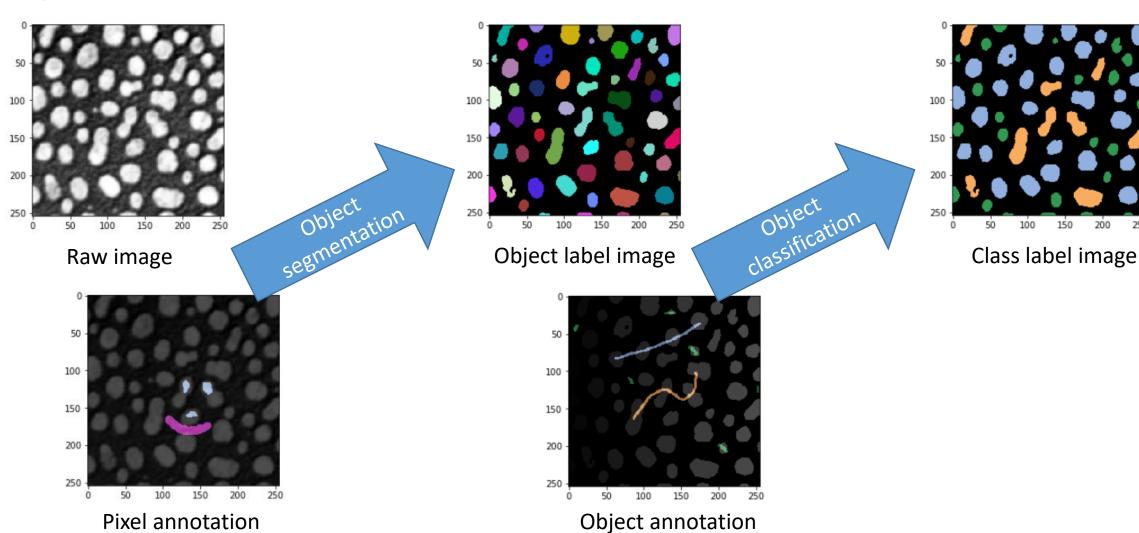
Robert Haase

Accelerated pixel and object classification

@haesleinhuepf



 APOC is a python library that makes use of OpenCL-compatible Graphics Cards to accelerate pixel and object classification

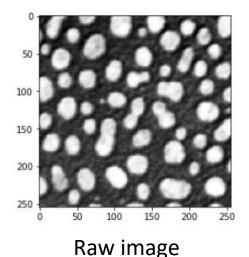


Bio-Image Analysis May 2022

Object segmentation



Pixel classification + connected component labeling



```
50 100 150 200 Pixel annotation
```

```
# define features
features = "gaussian_blur=1 gaussian_blur=5 sobel_of_gaussian_blur=1"

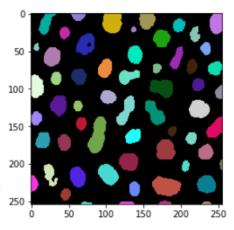
# this is where the model will be saved
cl_filename = 'my_object_segmenter.cl'

# delete classifier in case the file exists already
apoc.erase_classifier(cl_filename)

# train classifier
clf = apoc.ObjectSegmenter(opencl_filename=cl_filename, positive_class_identifier=2)
clf.train(features, manual_annotations, image)

segmentation_result = clf.predict(features=features, image=image)
cle.imshow(segmentation_result, labels=True)
```



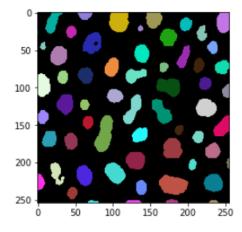


Object label image

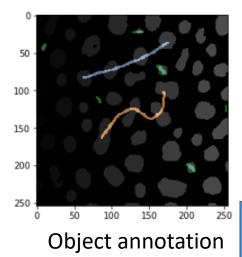
Object classification



Feature extraction + tabular classification



Object label image



```
# for the classification we define size and shape as criteria
features = 'area mean_max_distance_to_centroid_ratio'
```

```
# This is where the model will be saved
cl_filename_object_classifier = "my_object_classifier.cl"
```

```
# delete classifier in case the file exists already
apoc.erase_classifier(cl_filename_object_classifier)
```

```
# train the classifier
classifier = apoc.ObjectClassifier(cl_filename_object_classifier)
classifier.train(features, segmentation_result, annotation, image)
```

```
# determine object classification
classification_result = classifier.predict(segmentation_result, image)
cle.imshow(classification_result, labels=True)
```

50 - 100 - 150 - 200 - 250 - 2

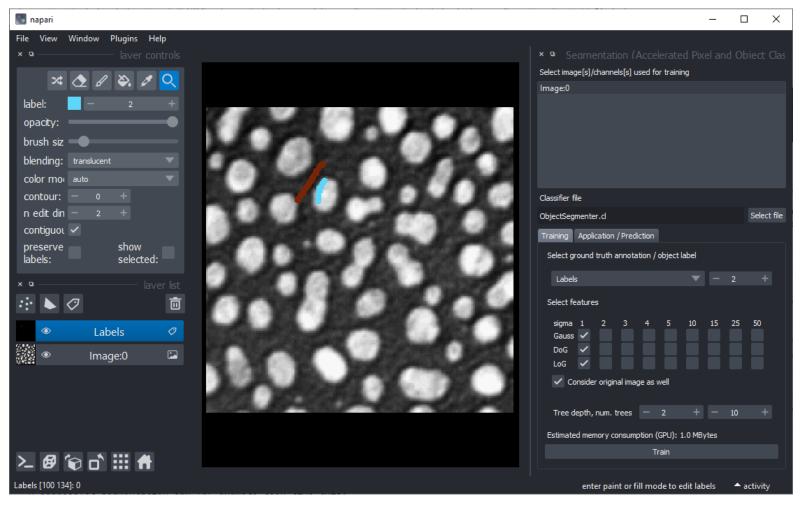
Class label image

Object classification

Graphical user interface



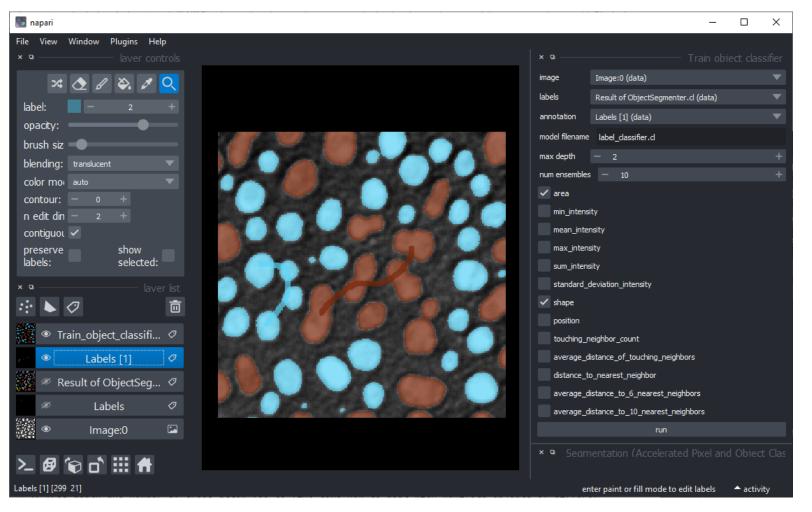
- Object segmentation
- https://github.com/haesleinhuepf/napari-accelerated-pixel-and-object-classification#object-and-semantic-segmentation



Graphical user interface



- Object classification
- https://github.com/haesleinhuepf/napari-accelerated-pixel-and-object-classification#object-classification





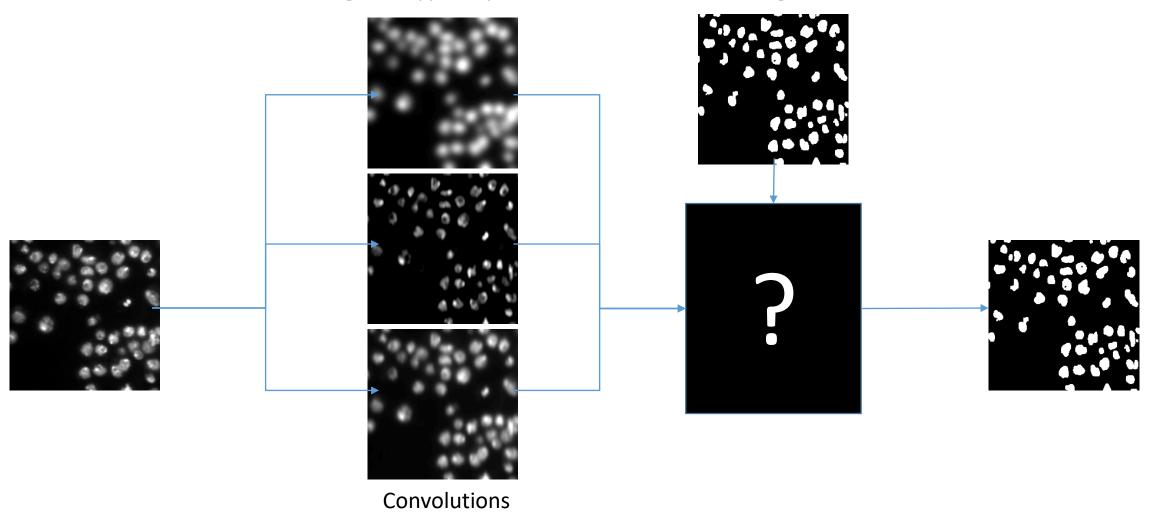


Summary & outlook

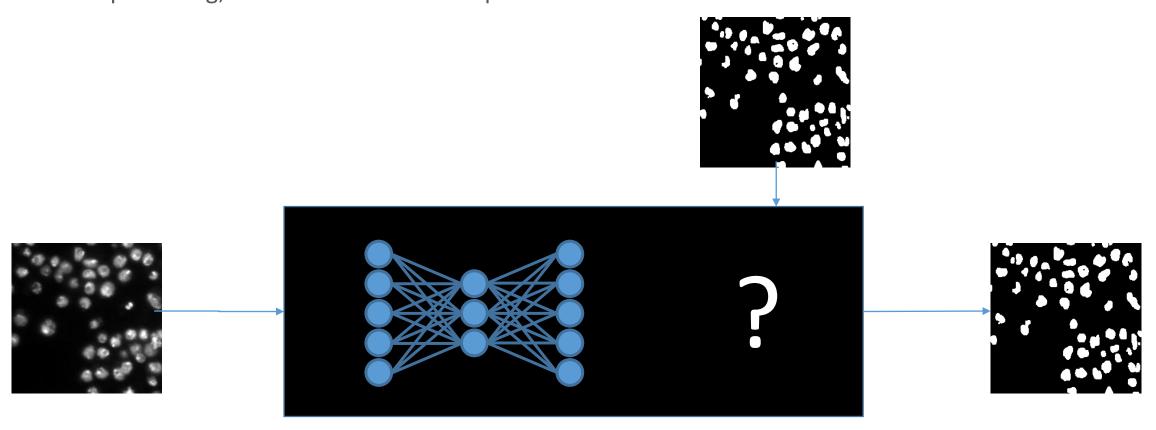




• In classifcal machine learning, we typically select features for training our classifier



In deep learning, this selection becomes part of the black box



Convolutional neural networks

Manual workflow design versus machine learning versus deep learning



Manual workflow design

Machine learning

Deep learning

Explicit definition of what's analysed

We specify features and how to combine them

Workload on human side is high

Semi-automatic feature selection

Fully automatic feature generation

We need to check carefully that the algorithm doesn't learn misleading features

Training can take long time

We can inspect code / models to understand how a decision is made

Allows us to analyze data we cannot manage manually

Uses neural networks

Uses artificial neural networks

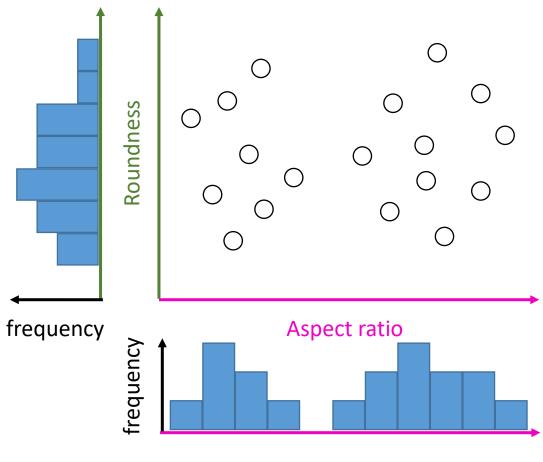
Needs training



Unsupervised machine learning



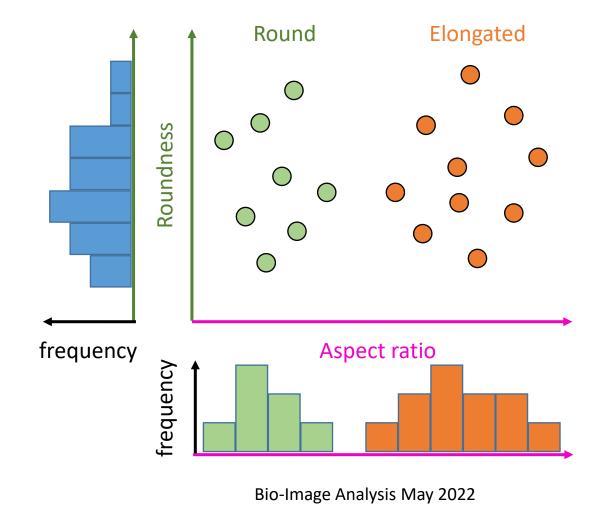
• If you don't provide ground truth, the algorithm is *unsupervised*.



Unsupervised machine learning



- If you don't provide ground truth, the algorithm is unsupervised.
- Nevertheless, algorithms can tell us something about the data



Further reading:

- Principal component analysis
- Cluster analysis





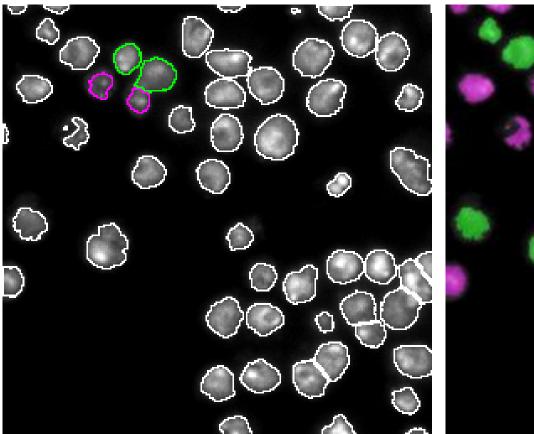
Exercises



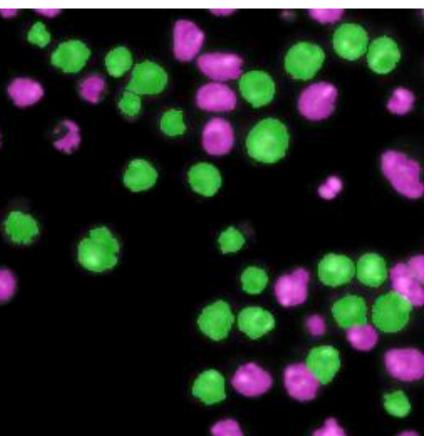


• Calculate precision and recall for the algorithm which created this prediction

Positive Negative



Ground truth annotation



Prediction

Summary



Today, you learned

- Machine learning for Pixel and Object classification
- Segmentation / prediction quality measurements
 - Precision
 - Recall
 - Jaccard Index
- Python
 - Scikit-learn / napari
 - Accelerated pixel and object classifiers (APOC)

Coming up next:

Hypothesis testing