

Machine Learning for Pixel and Object Segmentation



Lecture overview

Overview

- Machine learning for Pixel and Object Classification
 - Random Forest Classifiers
- Python
 - scikit-learn / napari
 - Accelerated pixel and object classification (APOC)

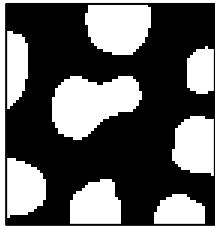
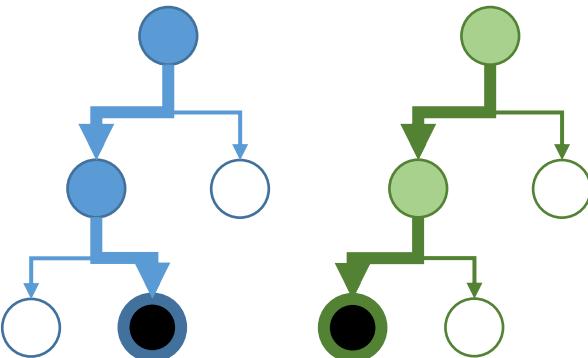
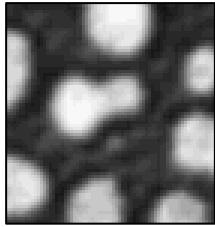
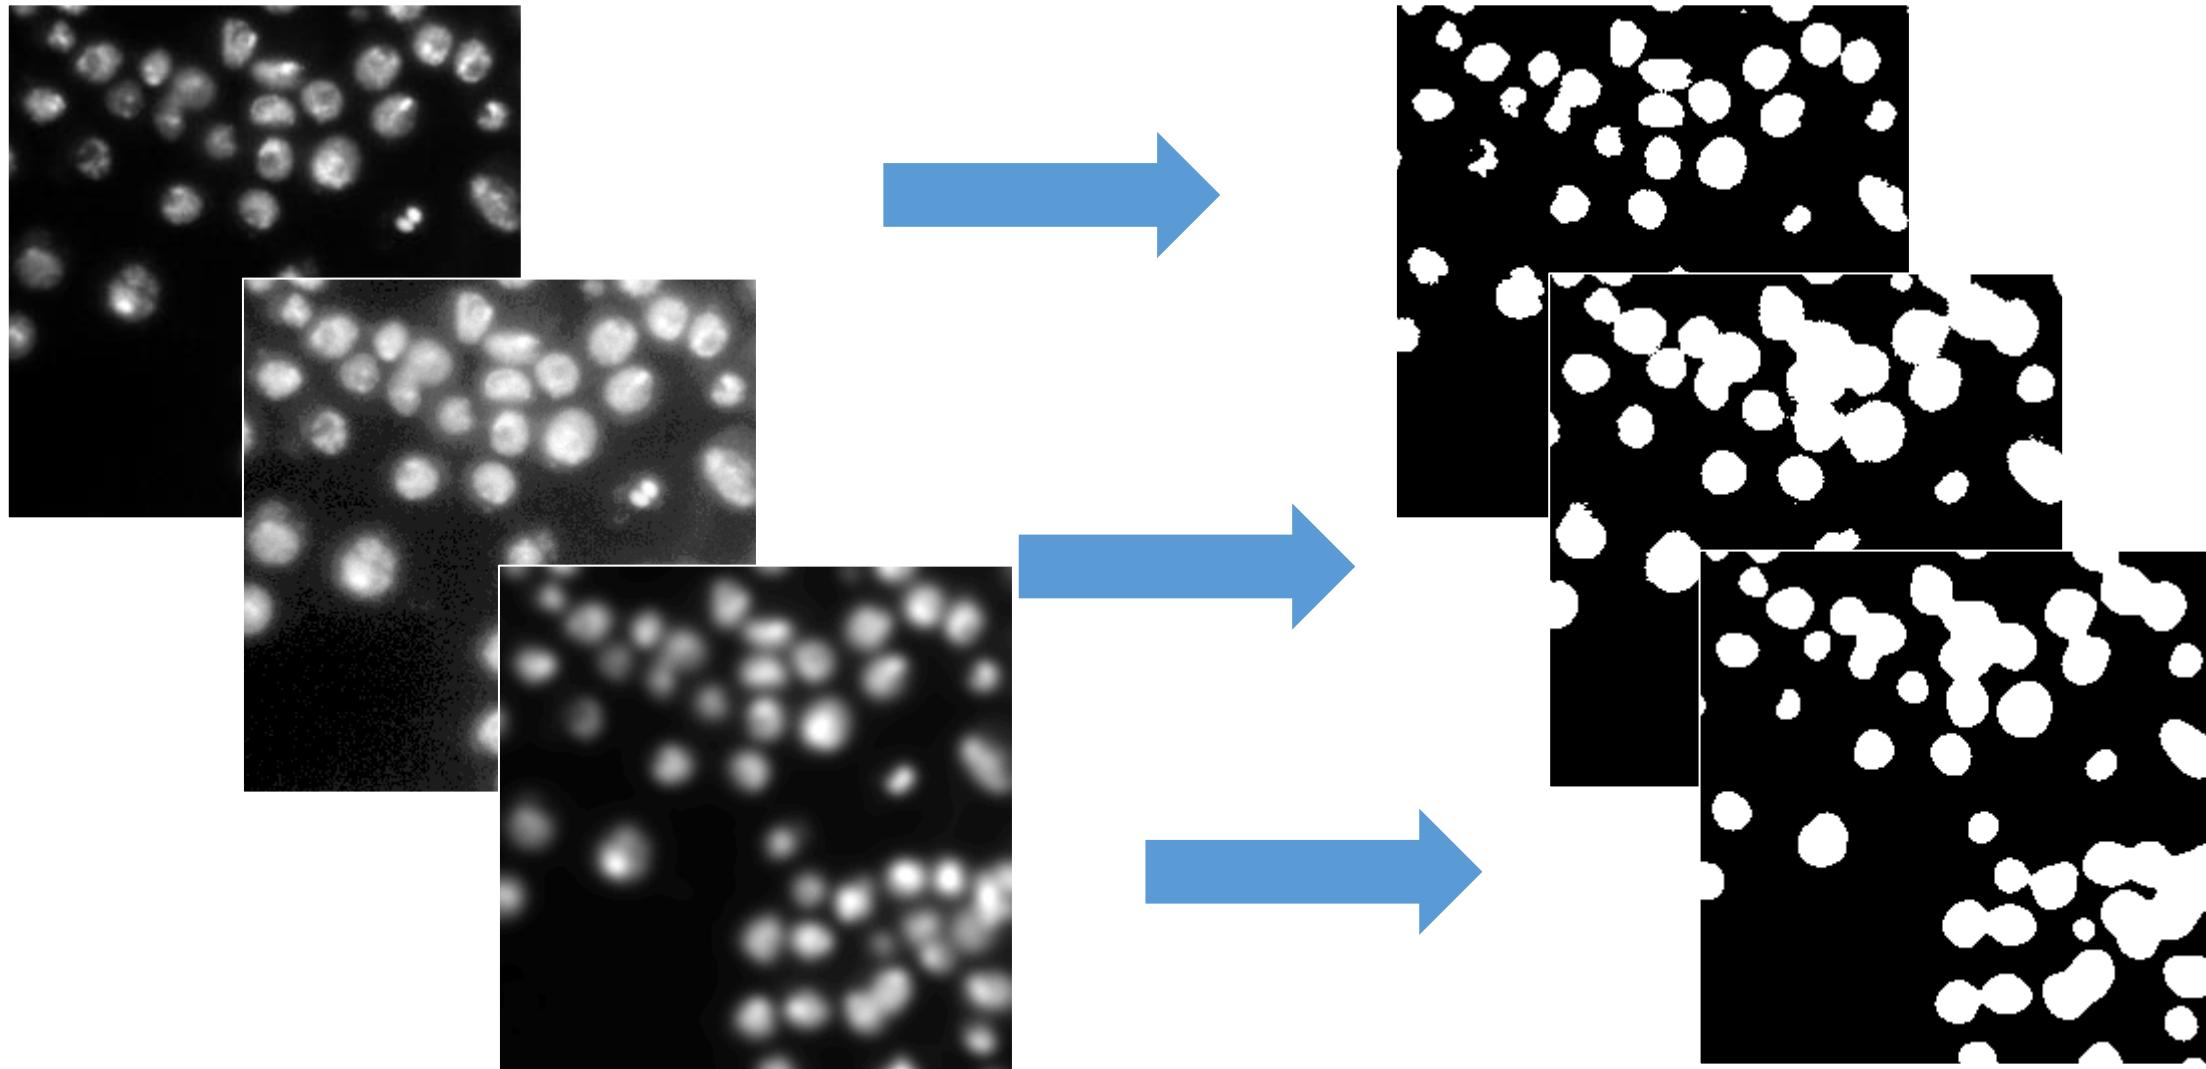


Image segmentation using thresholding

- Recap: Finding the right workflow towards a good segmentation takes time



June 2023

Image data source: [BBBC038v1](#), available from the Broad Bioimage Benchmark Collection (Caicedo et al., Nature Methods, 2019).

Image segmentation using thresholding

- Recap: Combining images, e.g. using Difference of Gaussian (DoG)

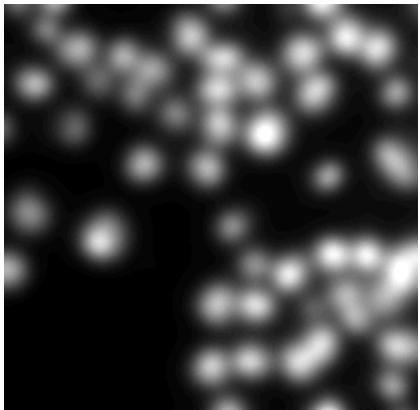
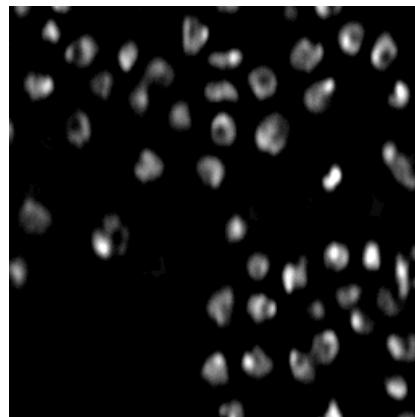
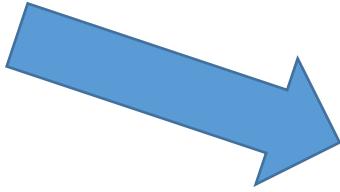
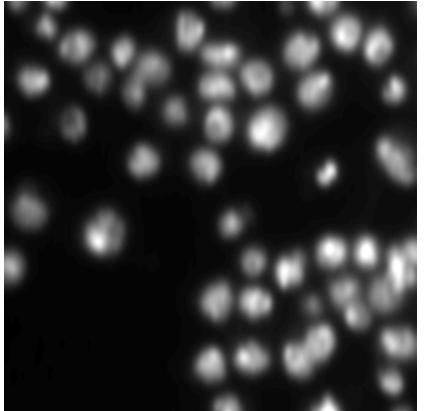
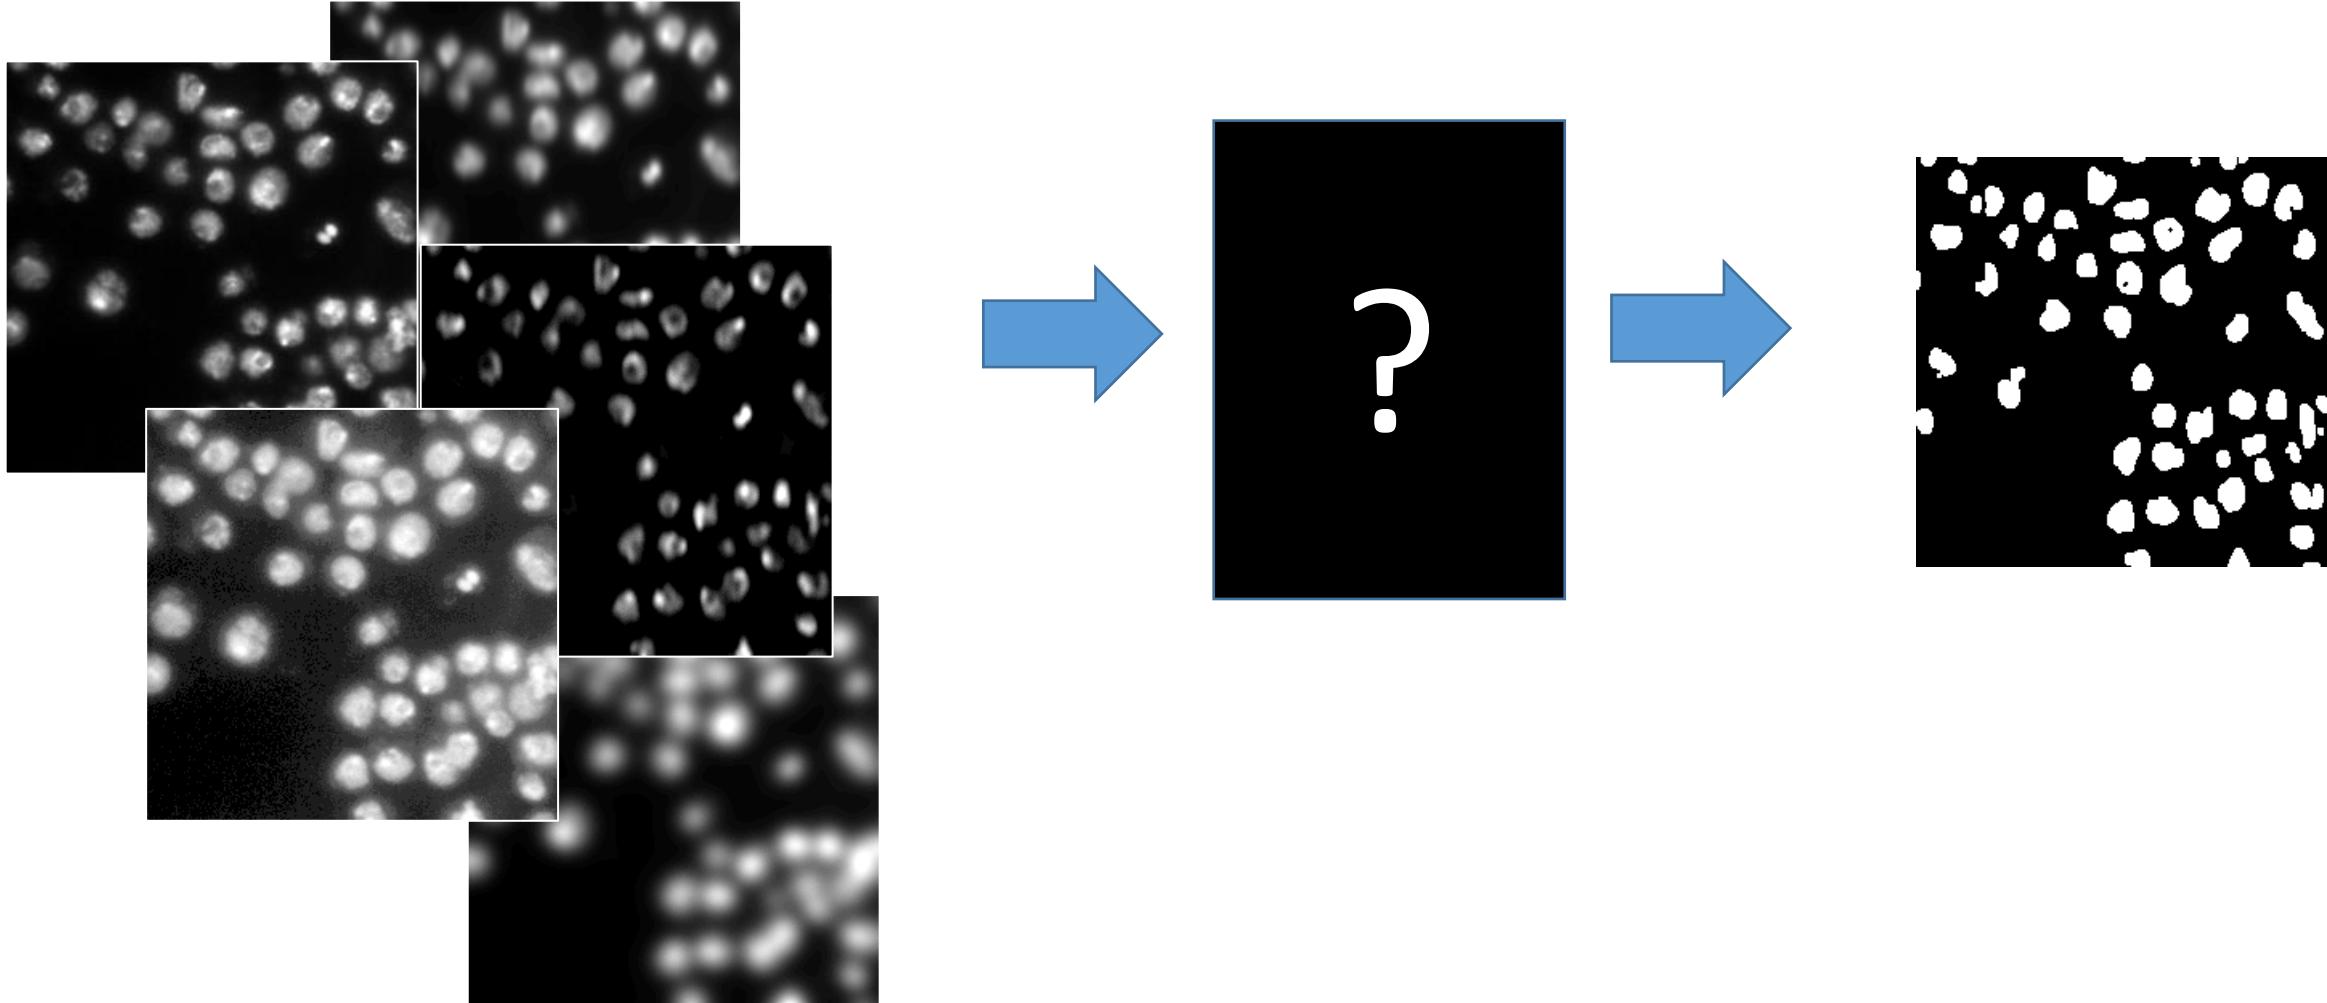


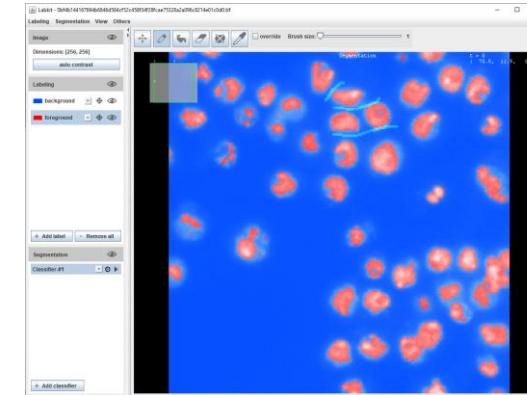
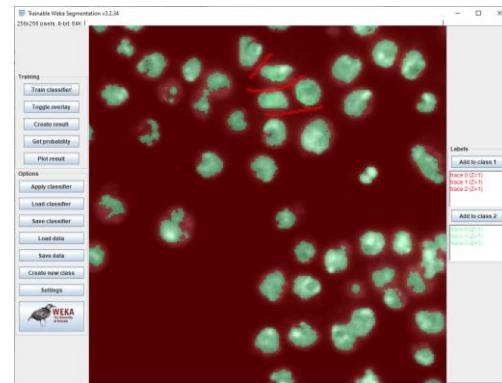
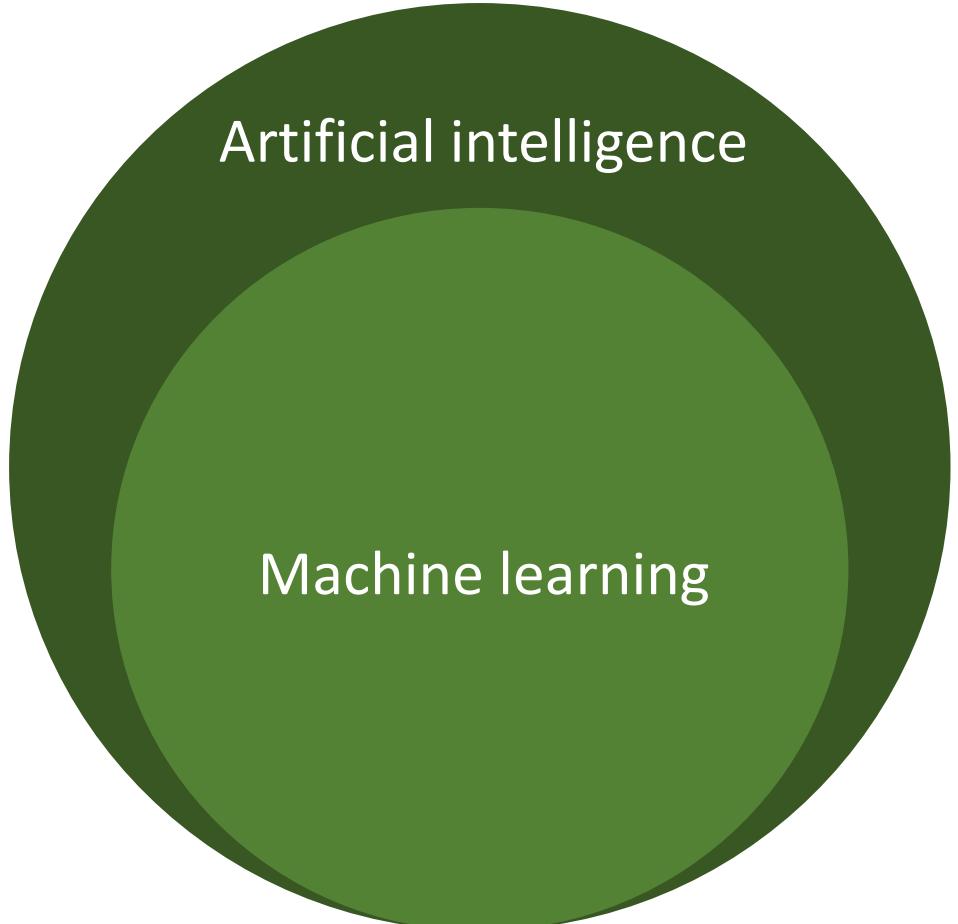
Image segmentation using thresholding

- Might there be a technology for optimization which combination of images can be used to get the best segmentation result?

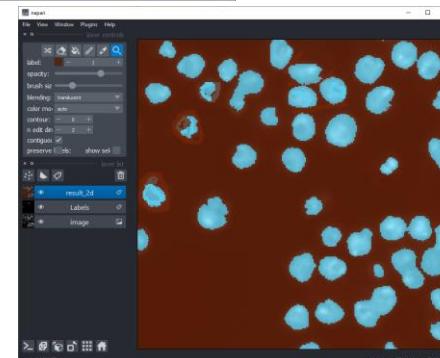


Machine learning

- A research field in computer science
- Finds more and more applications, also in life sciences.



Python /
scikit-learn /
napari /
apoc

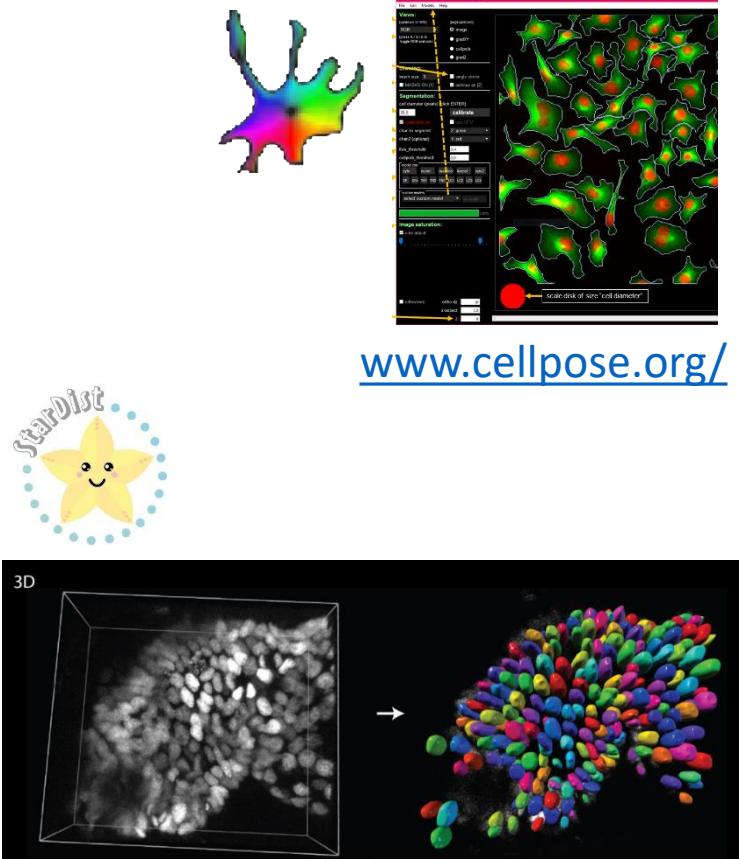
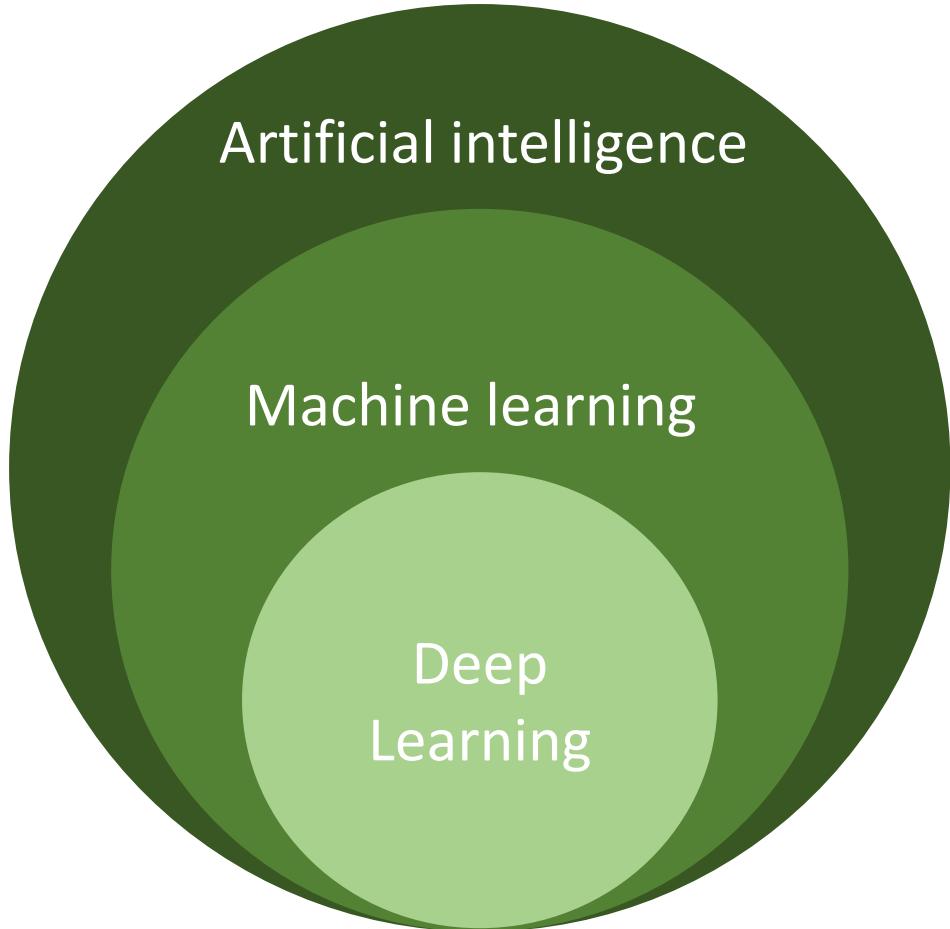


Trainable Weka Segmentation
<https://imagej.net/plugins/tws/>

LabKit
<https://imagej.net/plugins/labkit/>

Machine learning

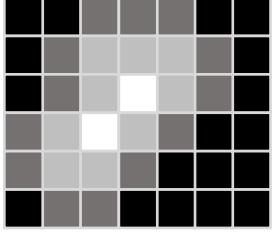
- A research field in computer science
- Finds more and more applications, also in life sciences.



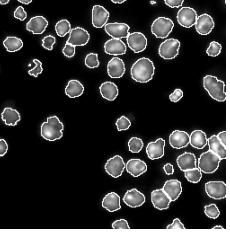
Machine learning

- Automatic construction of predictive models from given data

Pixels,



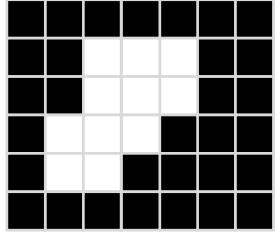
Objects,



Images, Audio, Text, Measurements, ...



Dense Segmentation / Binarization



Object classification

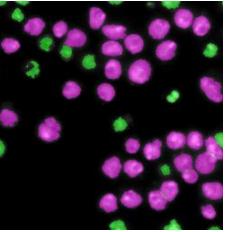
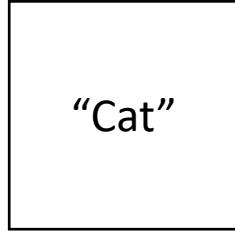
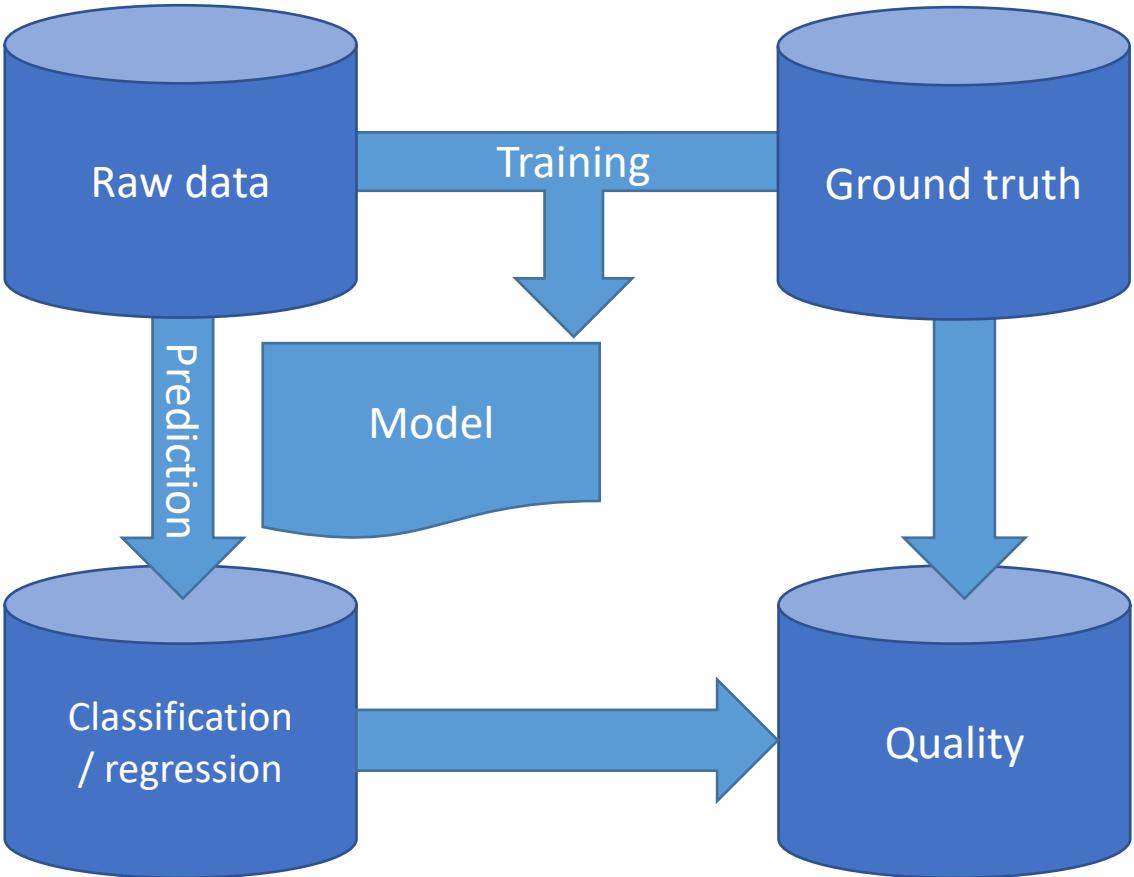
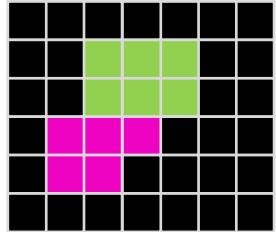


Image classification
“Cat”



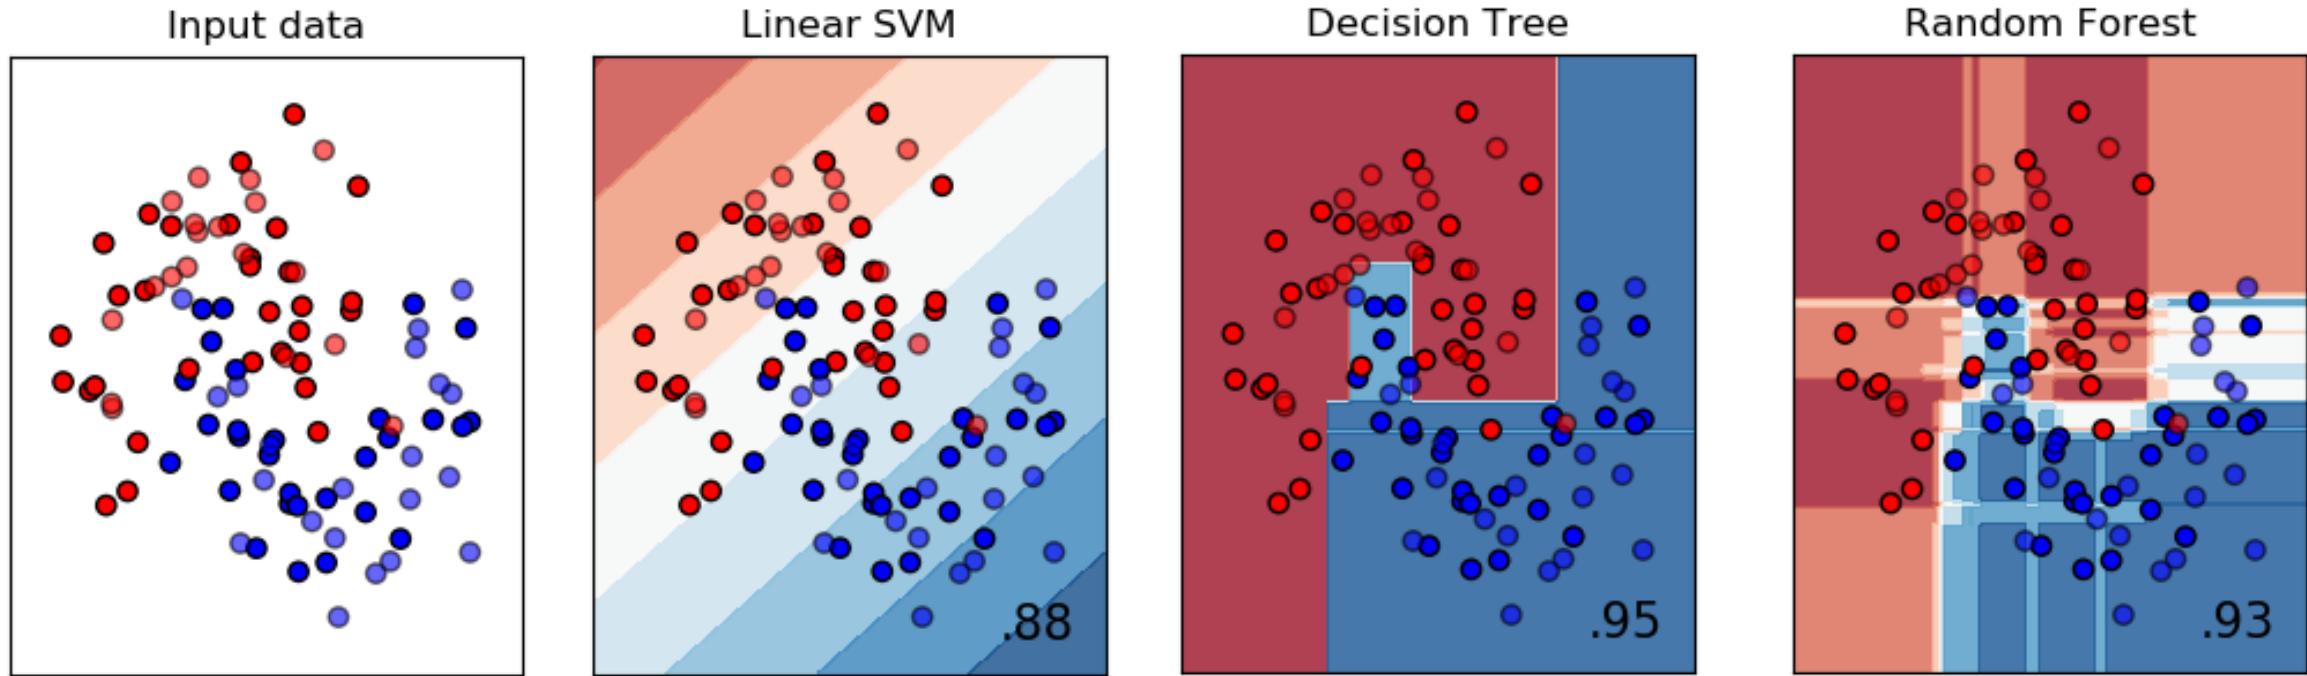
Instance segmentation



Annotated raw data, usually generated by humans

Precision,
Recall

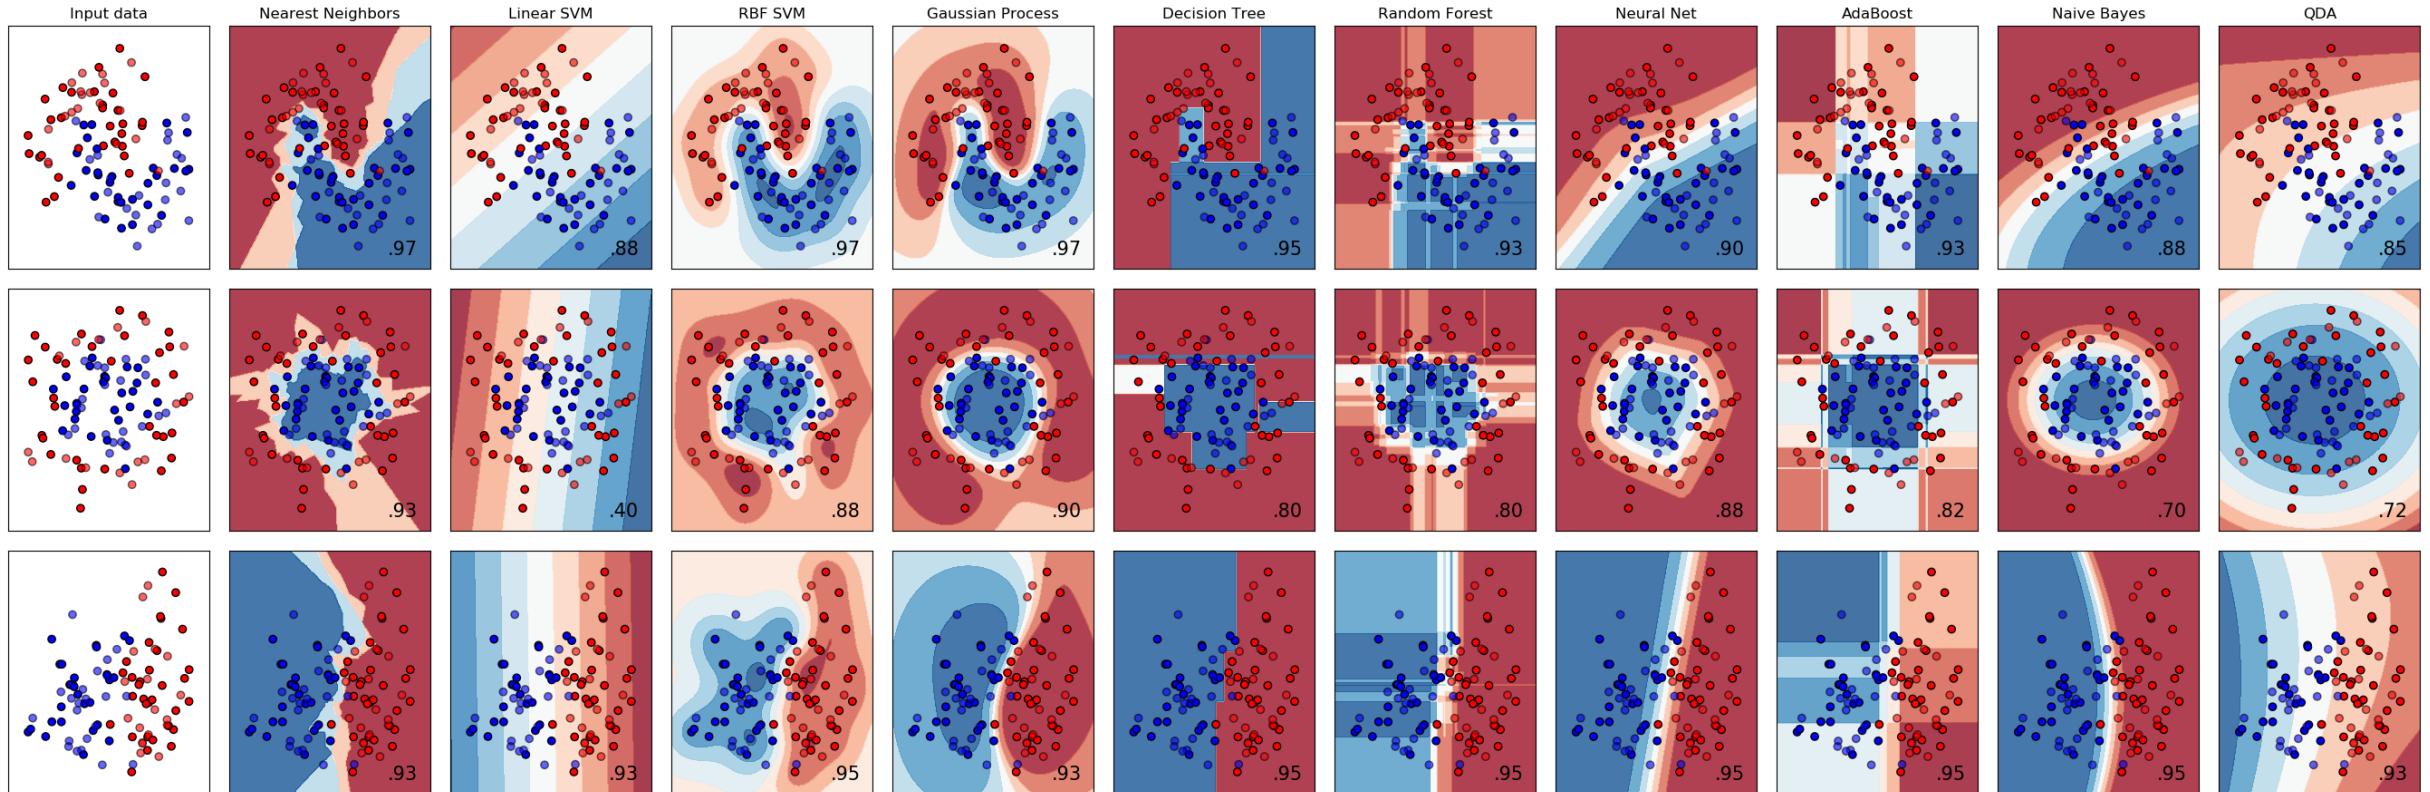
- Guess classification (**color**) from position of a sample in parameter space.



Adapted from https://scikit-learn.org/stable/auto_examples/classification/plot_classifier_comparison.html

Approaches

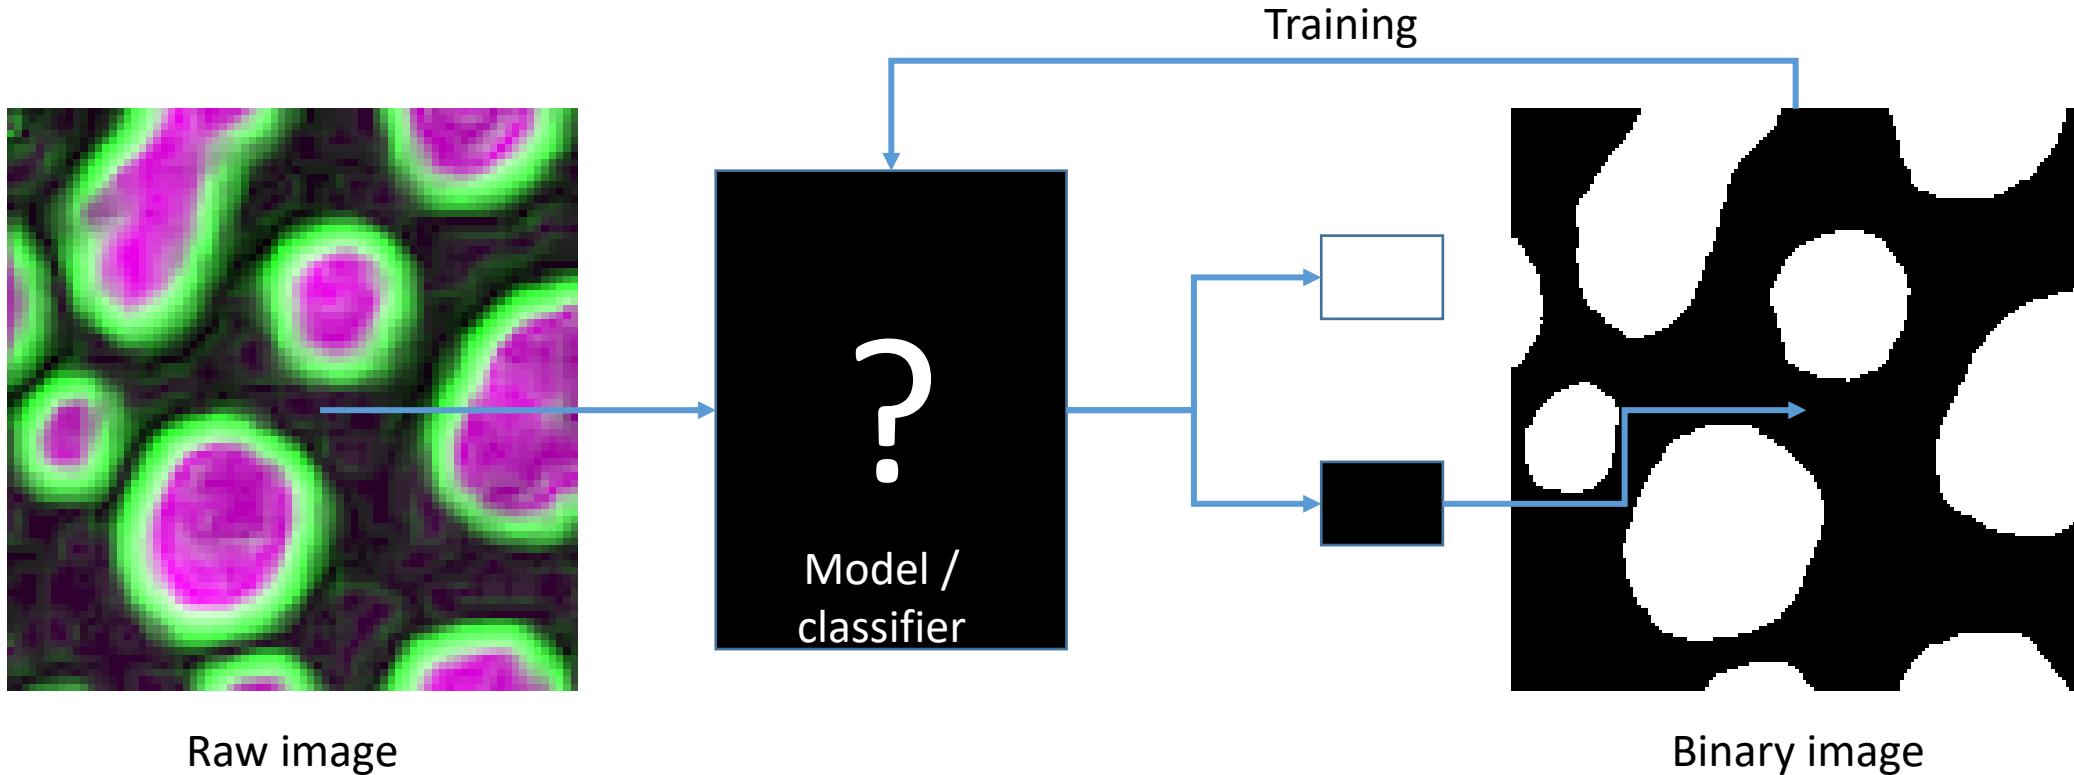
- The right approach depends on data, computational resources and desired quality



Adapted from https://scikit-learn.org/stable/auto_examples/classification/plot_classifier_comparison.html

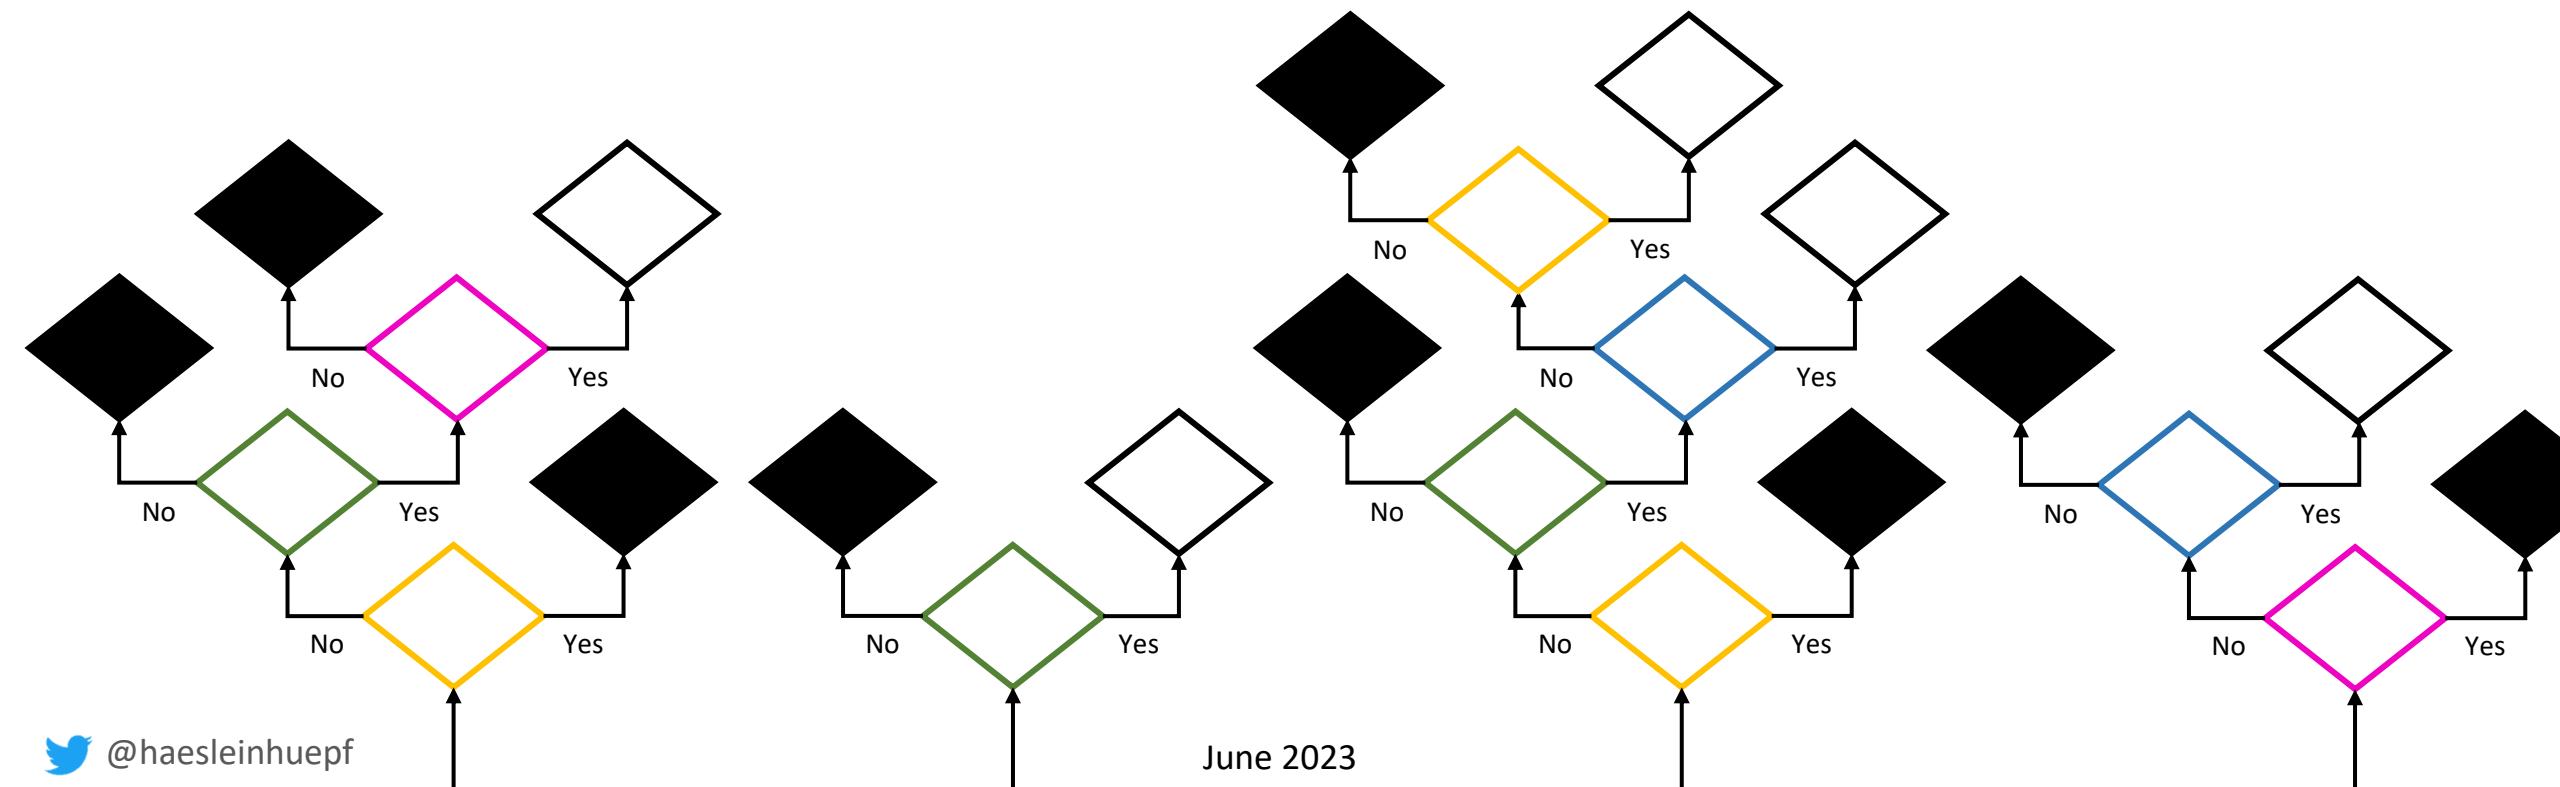
Machine learning for image segmentation

- *Supervised* machine learning: We give the computer some ground truth to learn from
- The computer derives a *model* or a *classifier* which can judge if a pixel should be foreground (white) or background (black)
- Example: Binary classifier



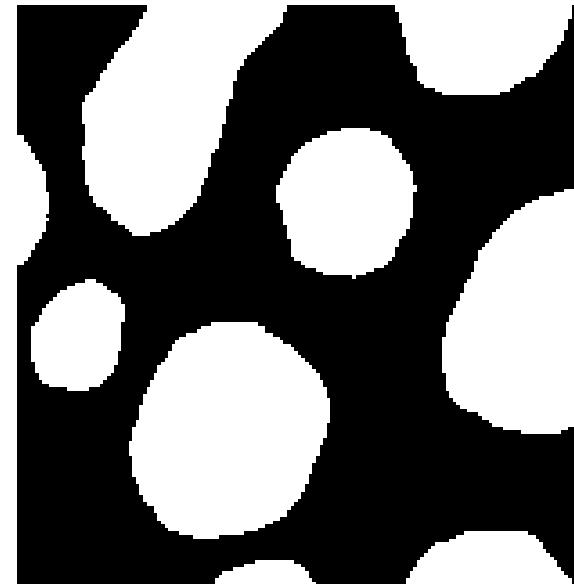
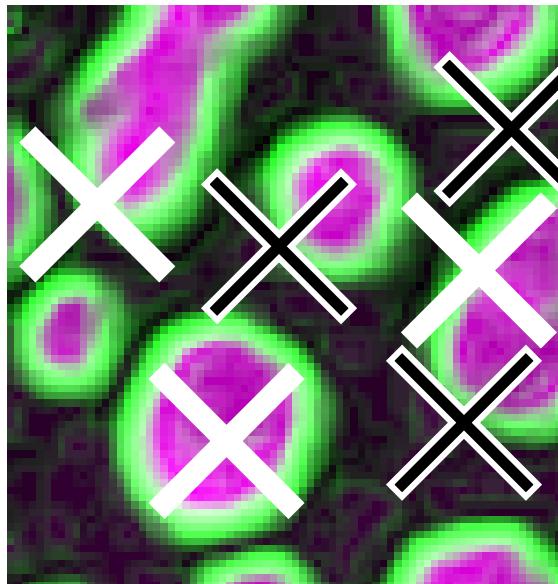
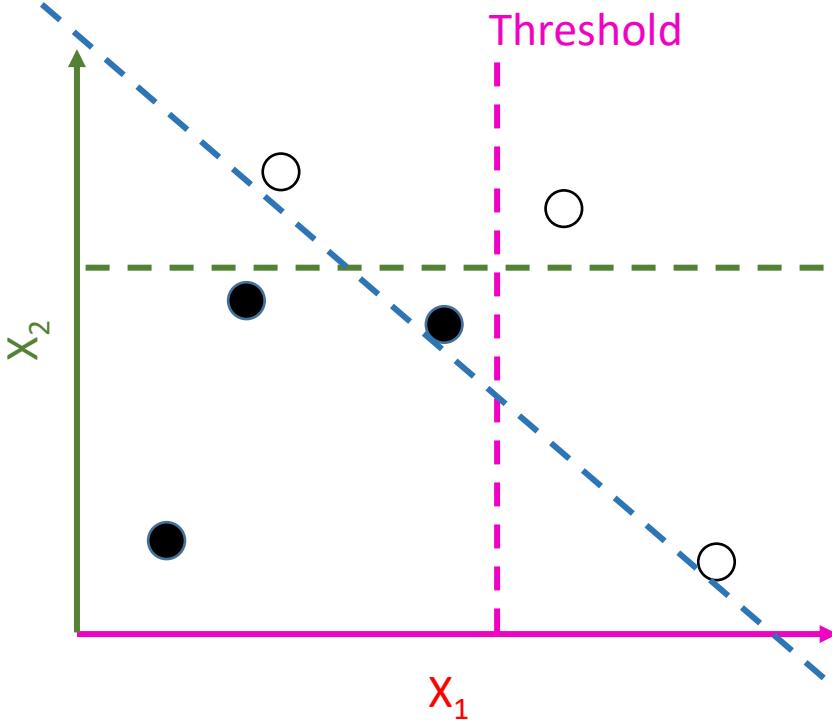
Random forest based image segmentation

- Decision trees are classifiers, they decide if a pixel should be white or black
 - Random decision trees are randomly initialized, afterwards evaluated and selected
 - Random forests consist of many random decision trees
-
- Example: Random forest of binary decision trees



Deriving random decision trees

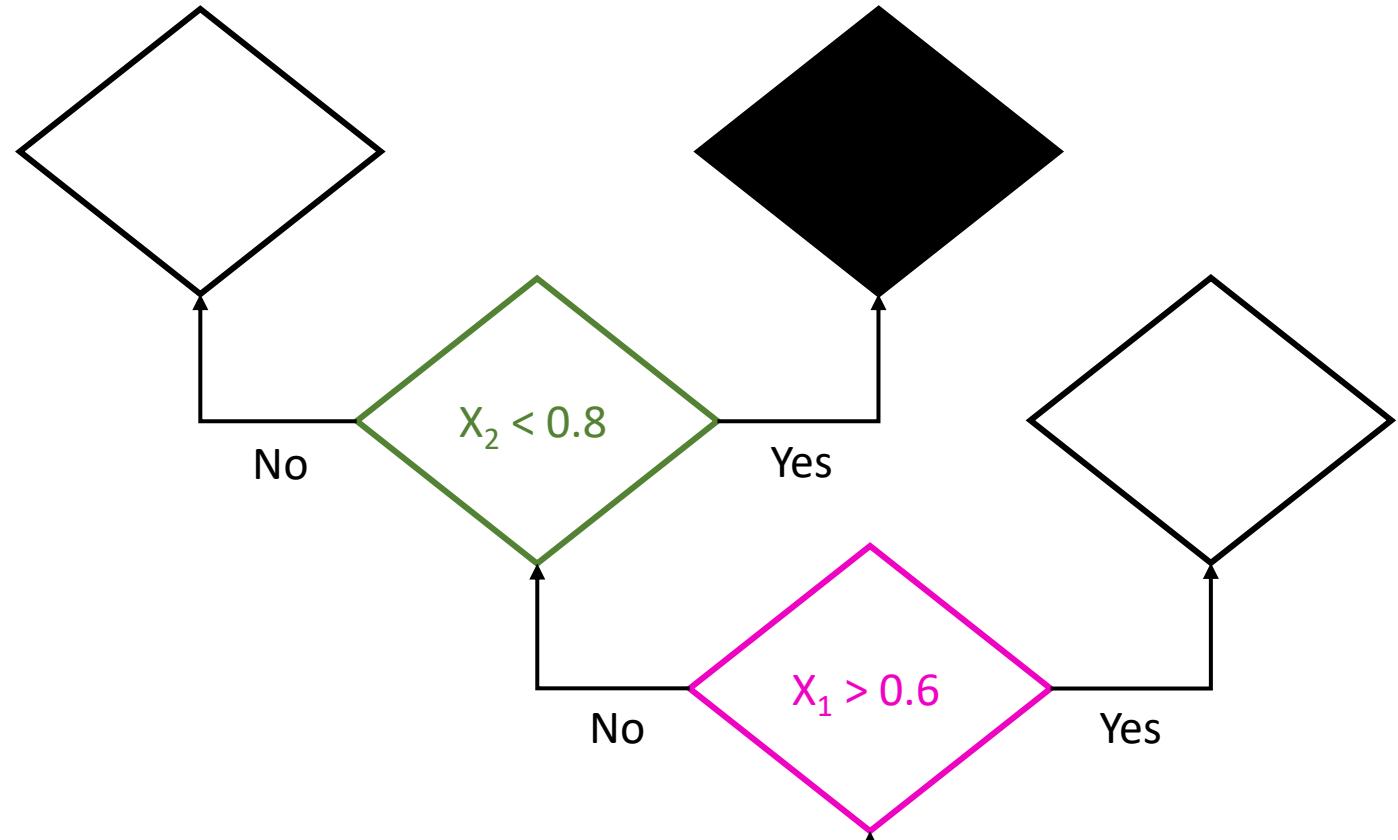
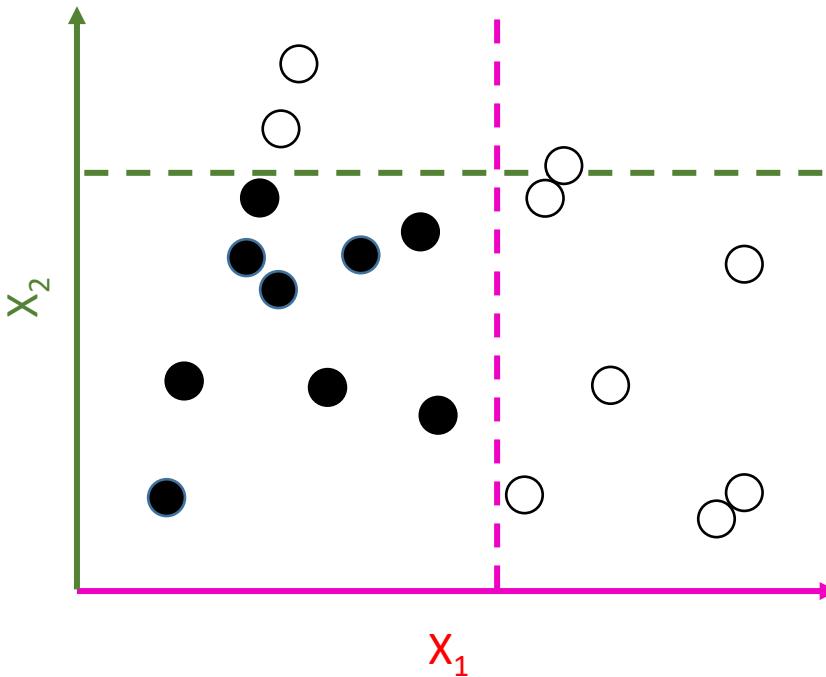
- For efficient processing, we randomly *sample* our data set
 - Individual pixels, their intensity and their classification



Note: You cannot use a single threshold to make the decision correctly

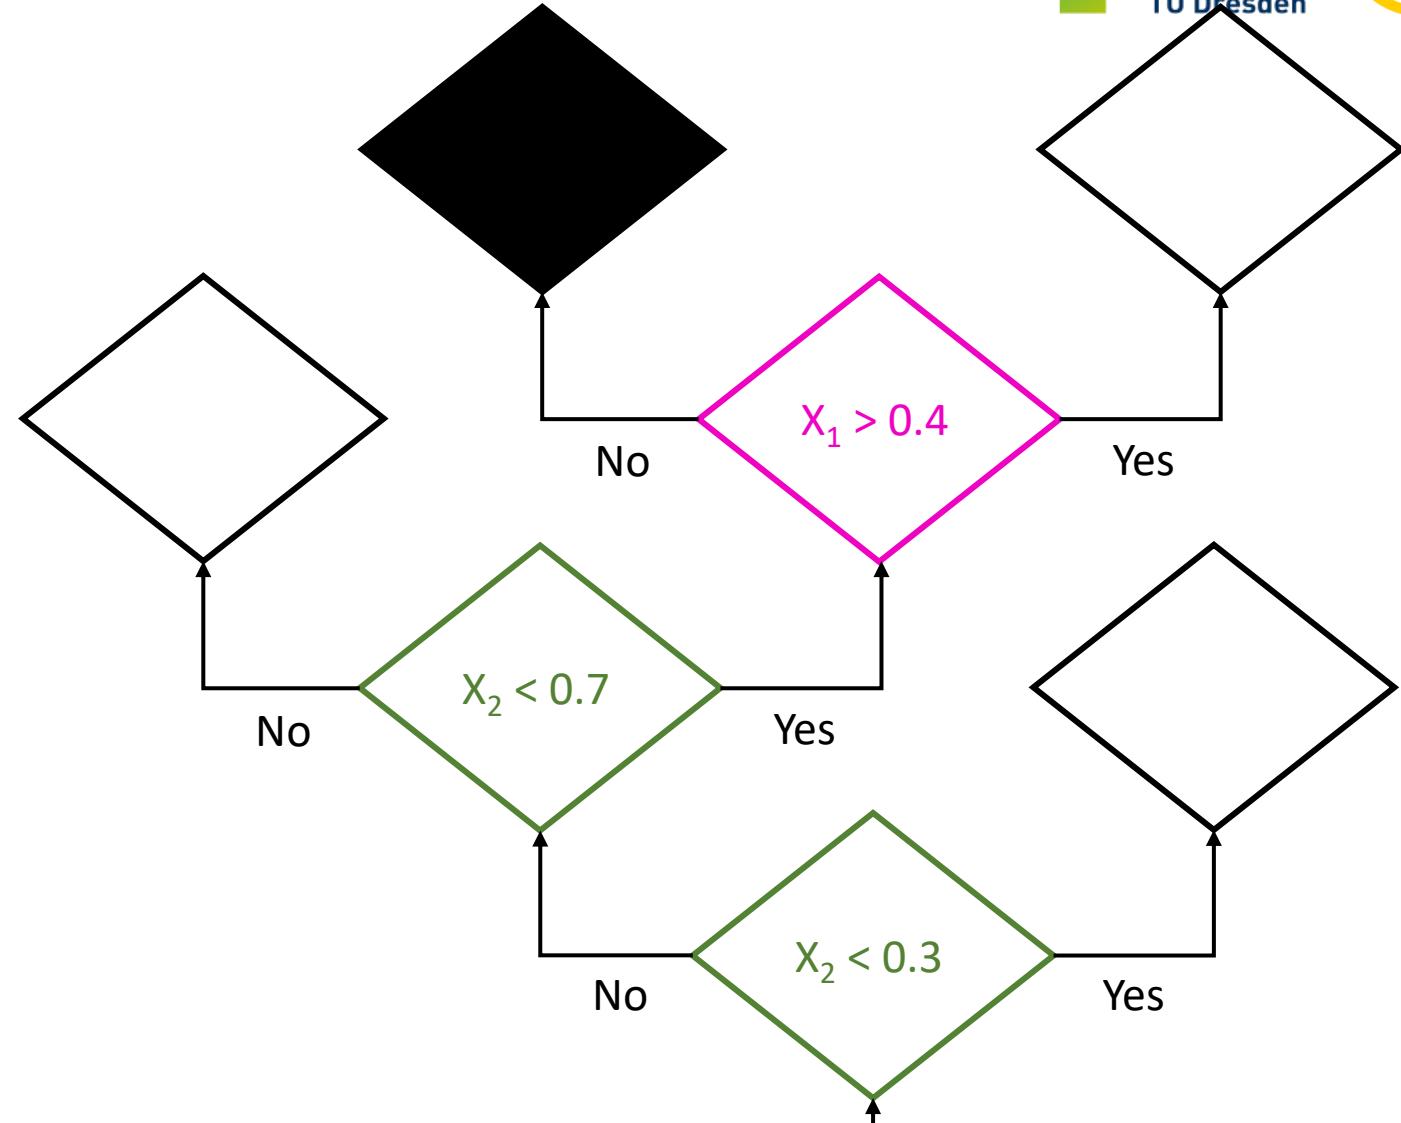
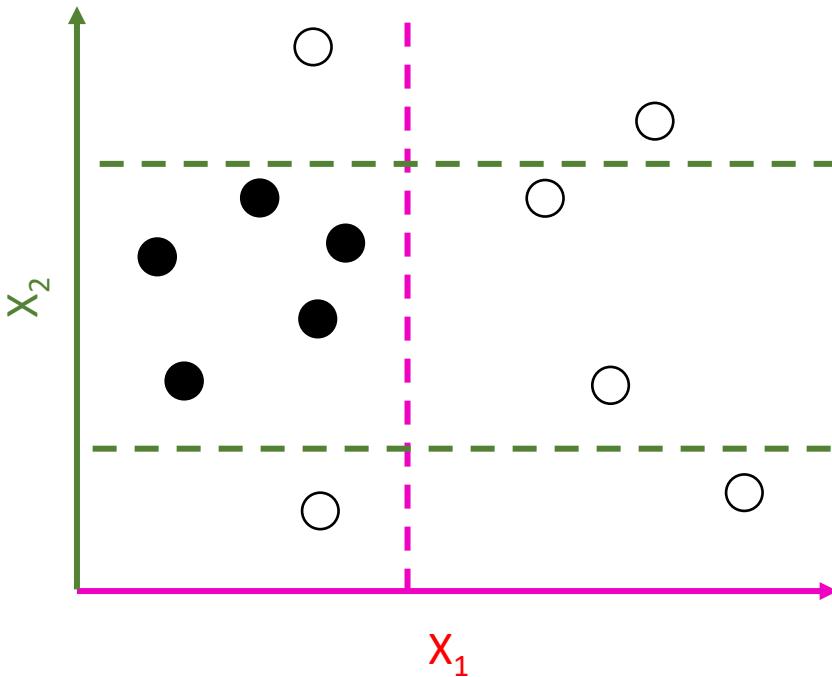
Deriving random decision trees

- Decision trees combine several thresholds on several parameters



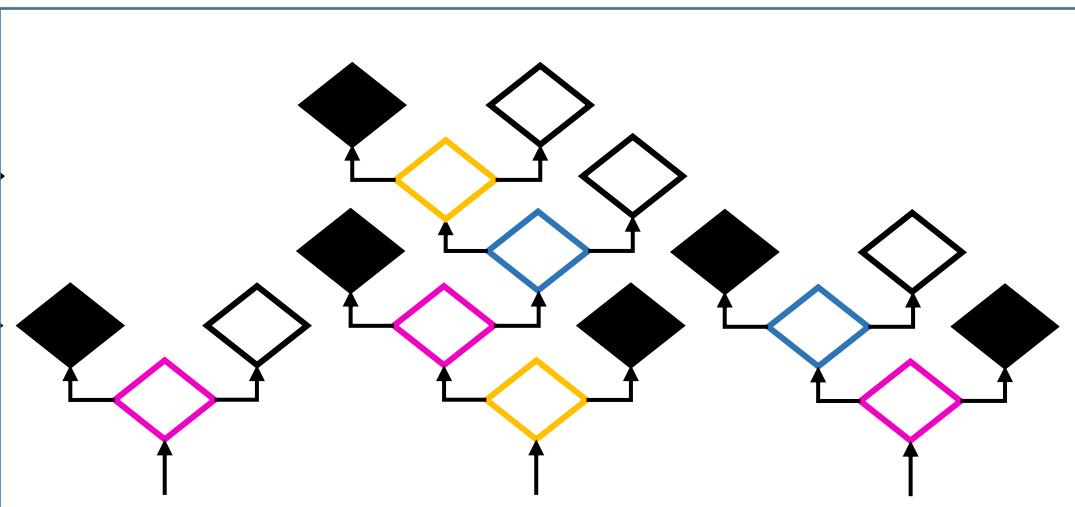
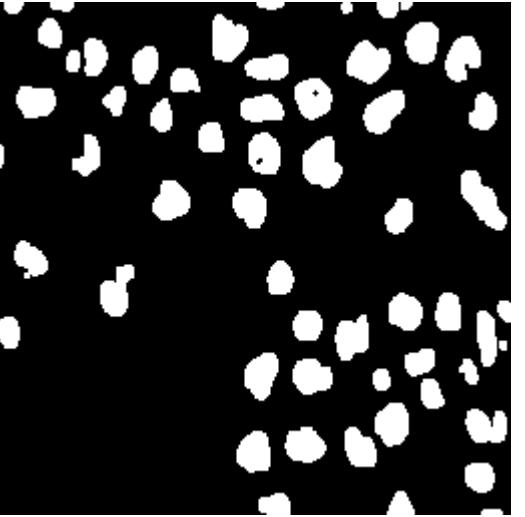
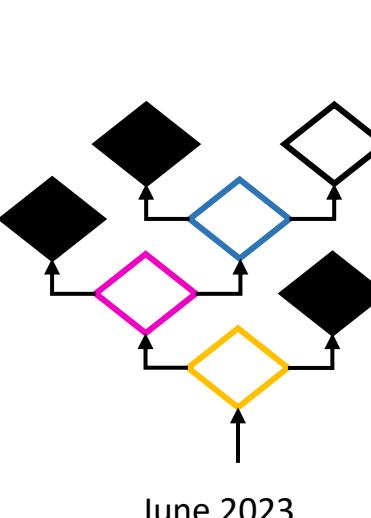
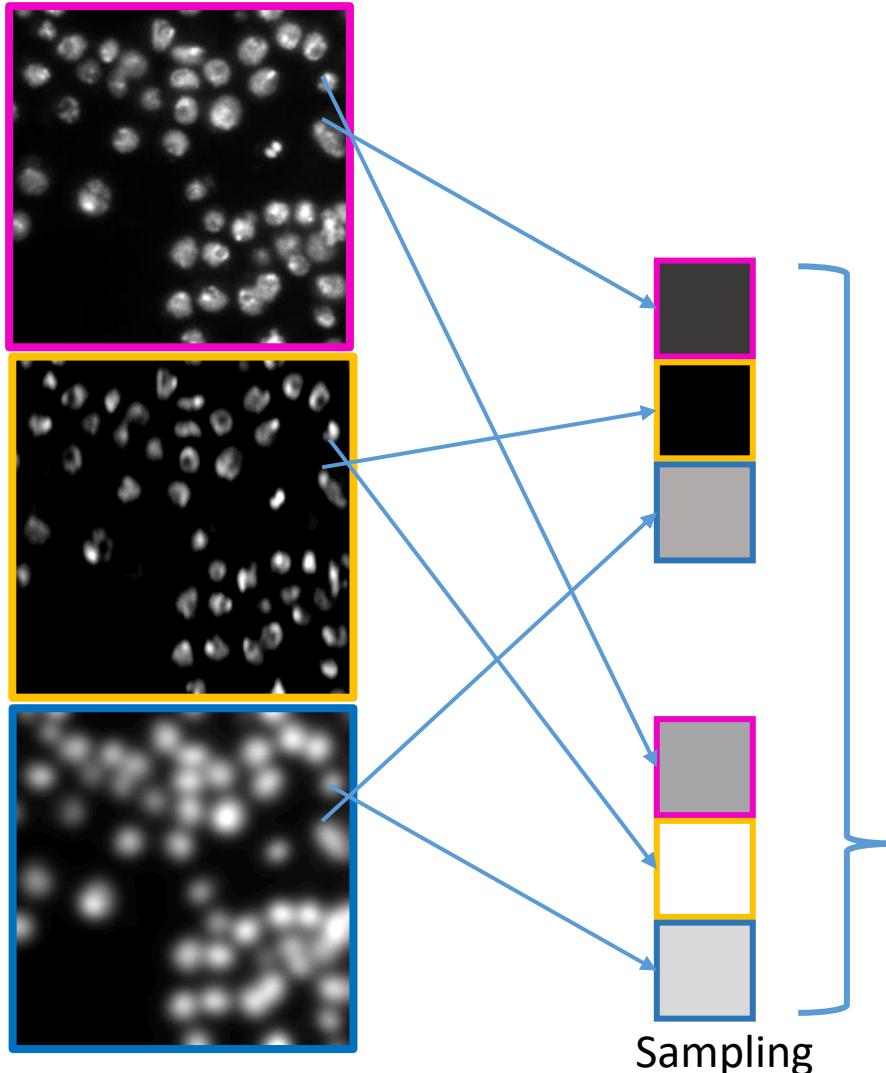
Deriving random decision trees

- Depending on sampling, the decision trees are different



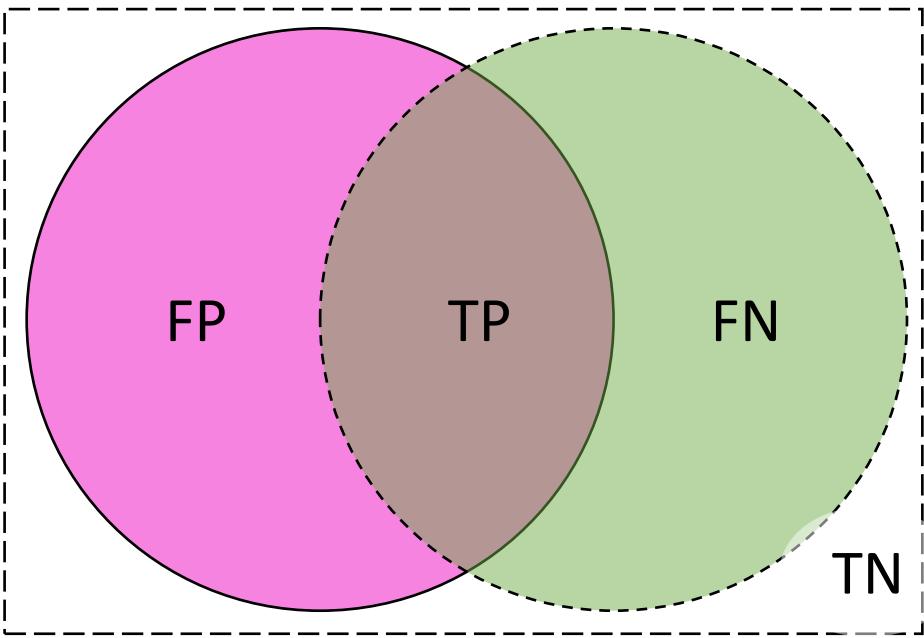
Random Forest Pixel Classifiers

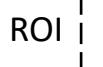
- By comparing performance of individual decision trees, good ones can be selected, bad ones excluded.



Segmentation quality estimation

- In general
 - Define what's positive and what's negative.
 - Compare with a reference to figure out what was true and false
- Welcome to the Theory of Sets



 A	Prediction A
 B	Reference B (ground truth)
 ROI	Region of interest
 TP	True-positive
 FN	False-negative
 FP	False-positive
TN	True-negative

Overlap
(a.k.a. Jaccard index) $\frac{TP}{TP + FN + FP}$

How much do A and B overlap?

Precision $\frac{TP}{TP + FP}$

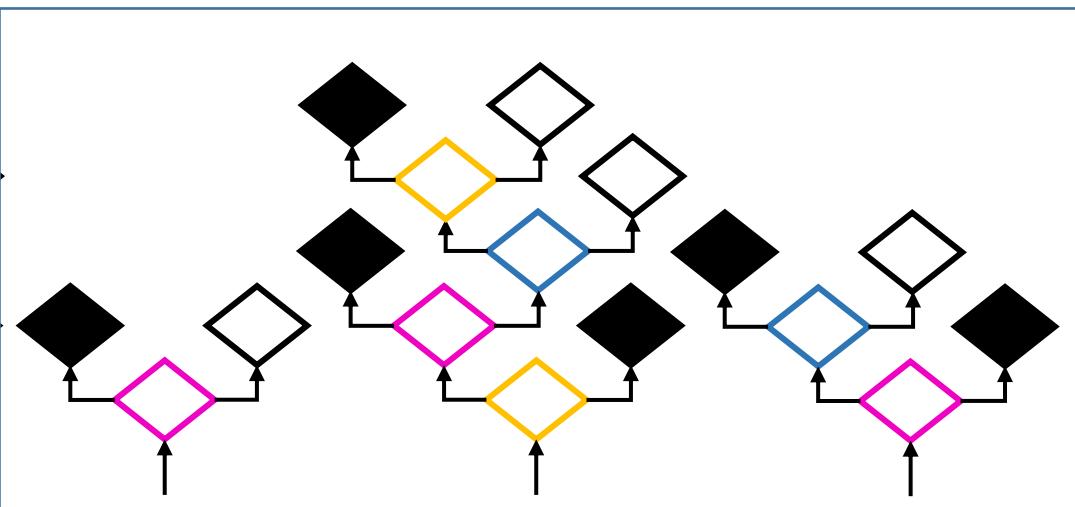
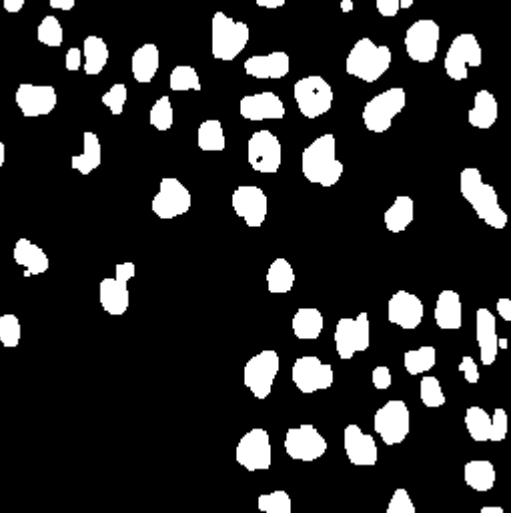
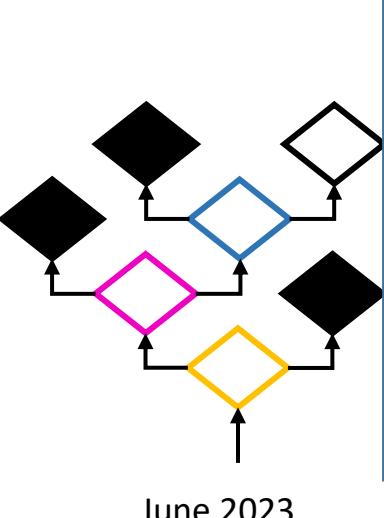
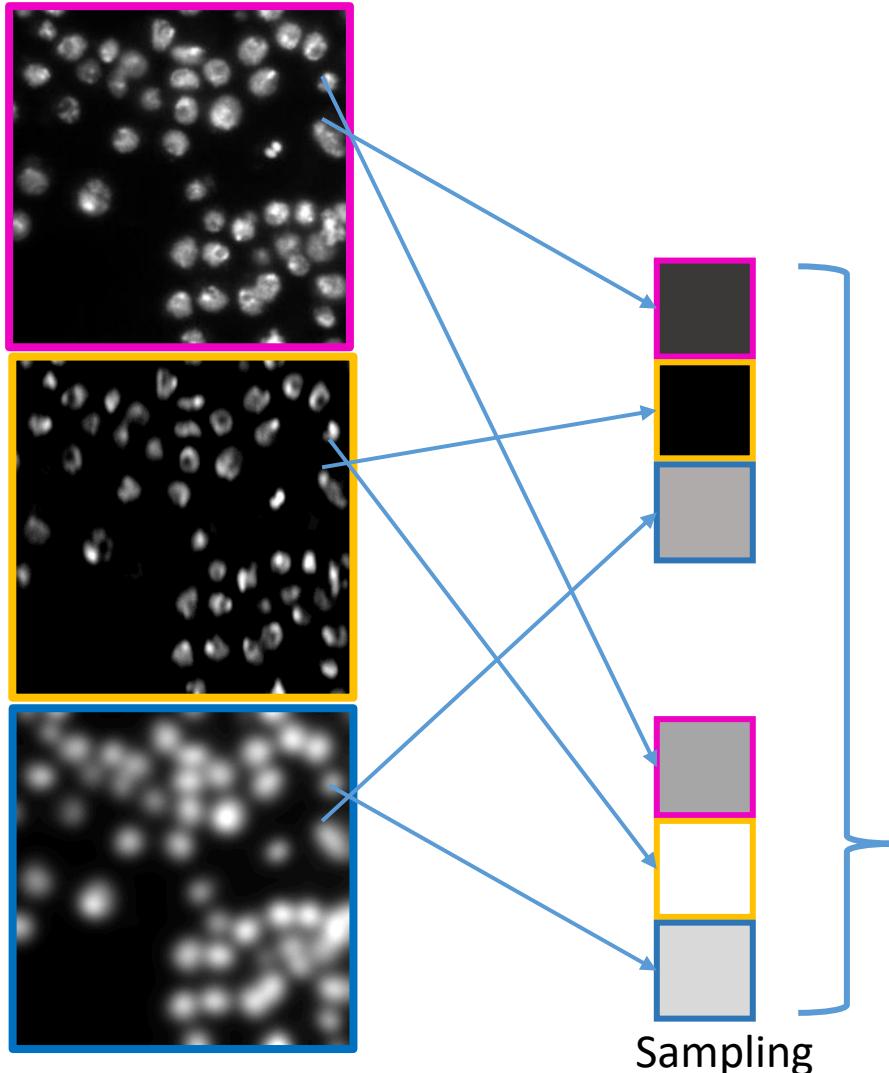
What fraction of points that were predicted as positives were really positive?

Recall
(a.k.a. sensitivity) $\frac{TP}{TP + FN}$

What fraction of positives points were predicted as positives?

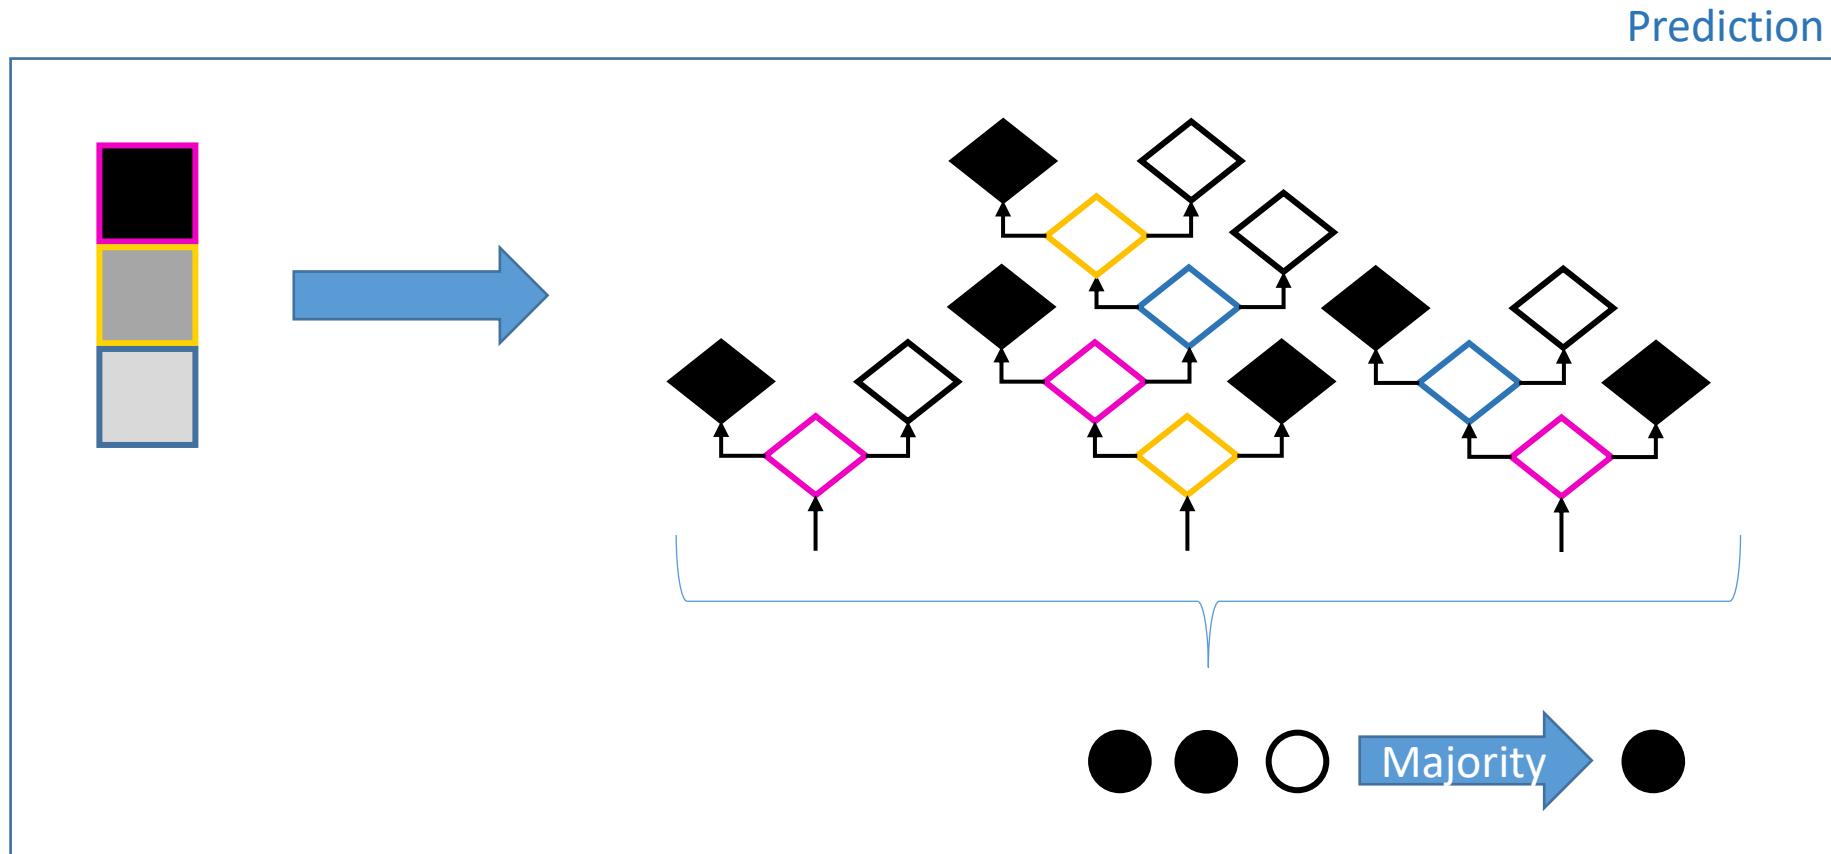
Random Forest Pixel Classifiers

- By comparing performance of individual decision trees, good ones can be selected, bad ones excluded.



Random Forest Pixel Classifiers

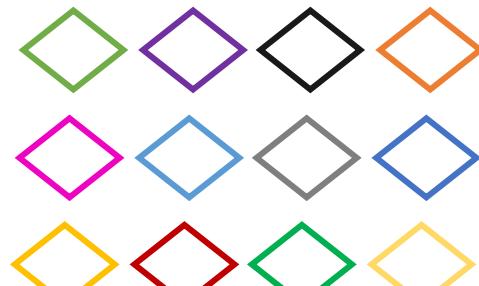
- Combination of individual tree decisions by voting or max / mean



Random Forest Pixel Classifiers

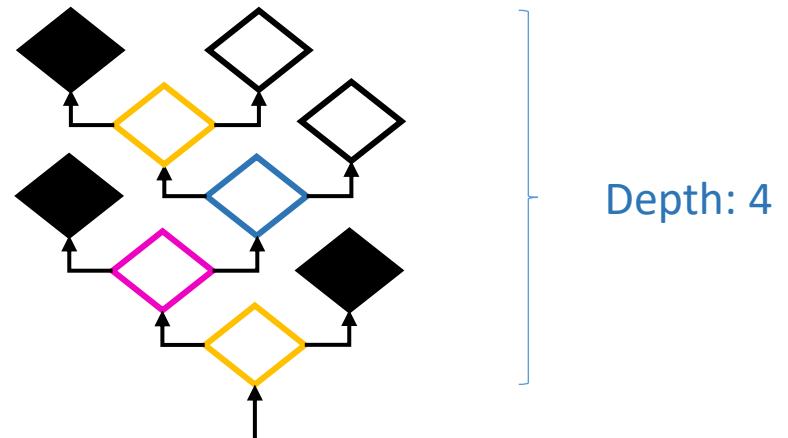
- Typical numbers for pixel classifiers in microscopy

Available features: > 20

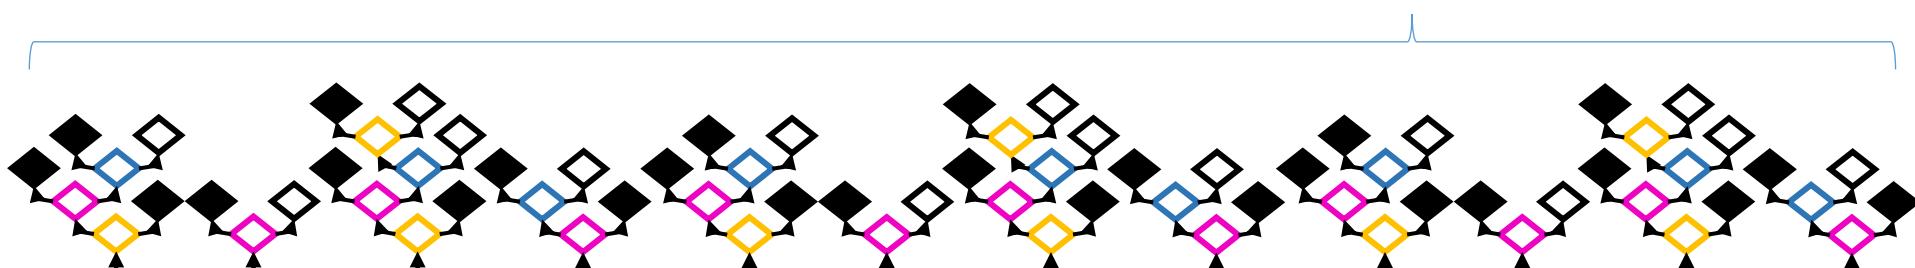


- Gaussian blur image
- DoG image
- LoG image
- Hessian
-

Selected features: <= depth



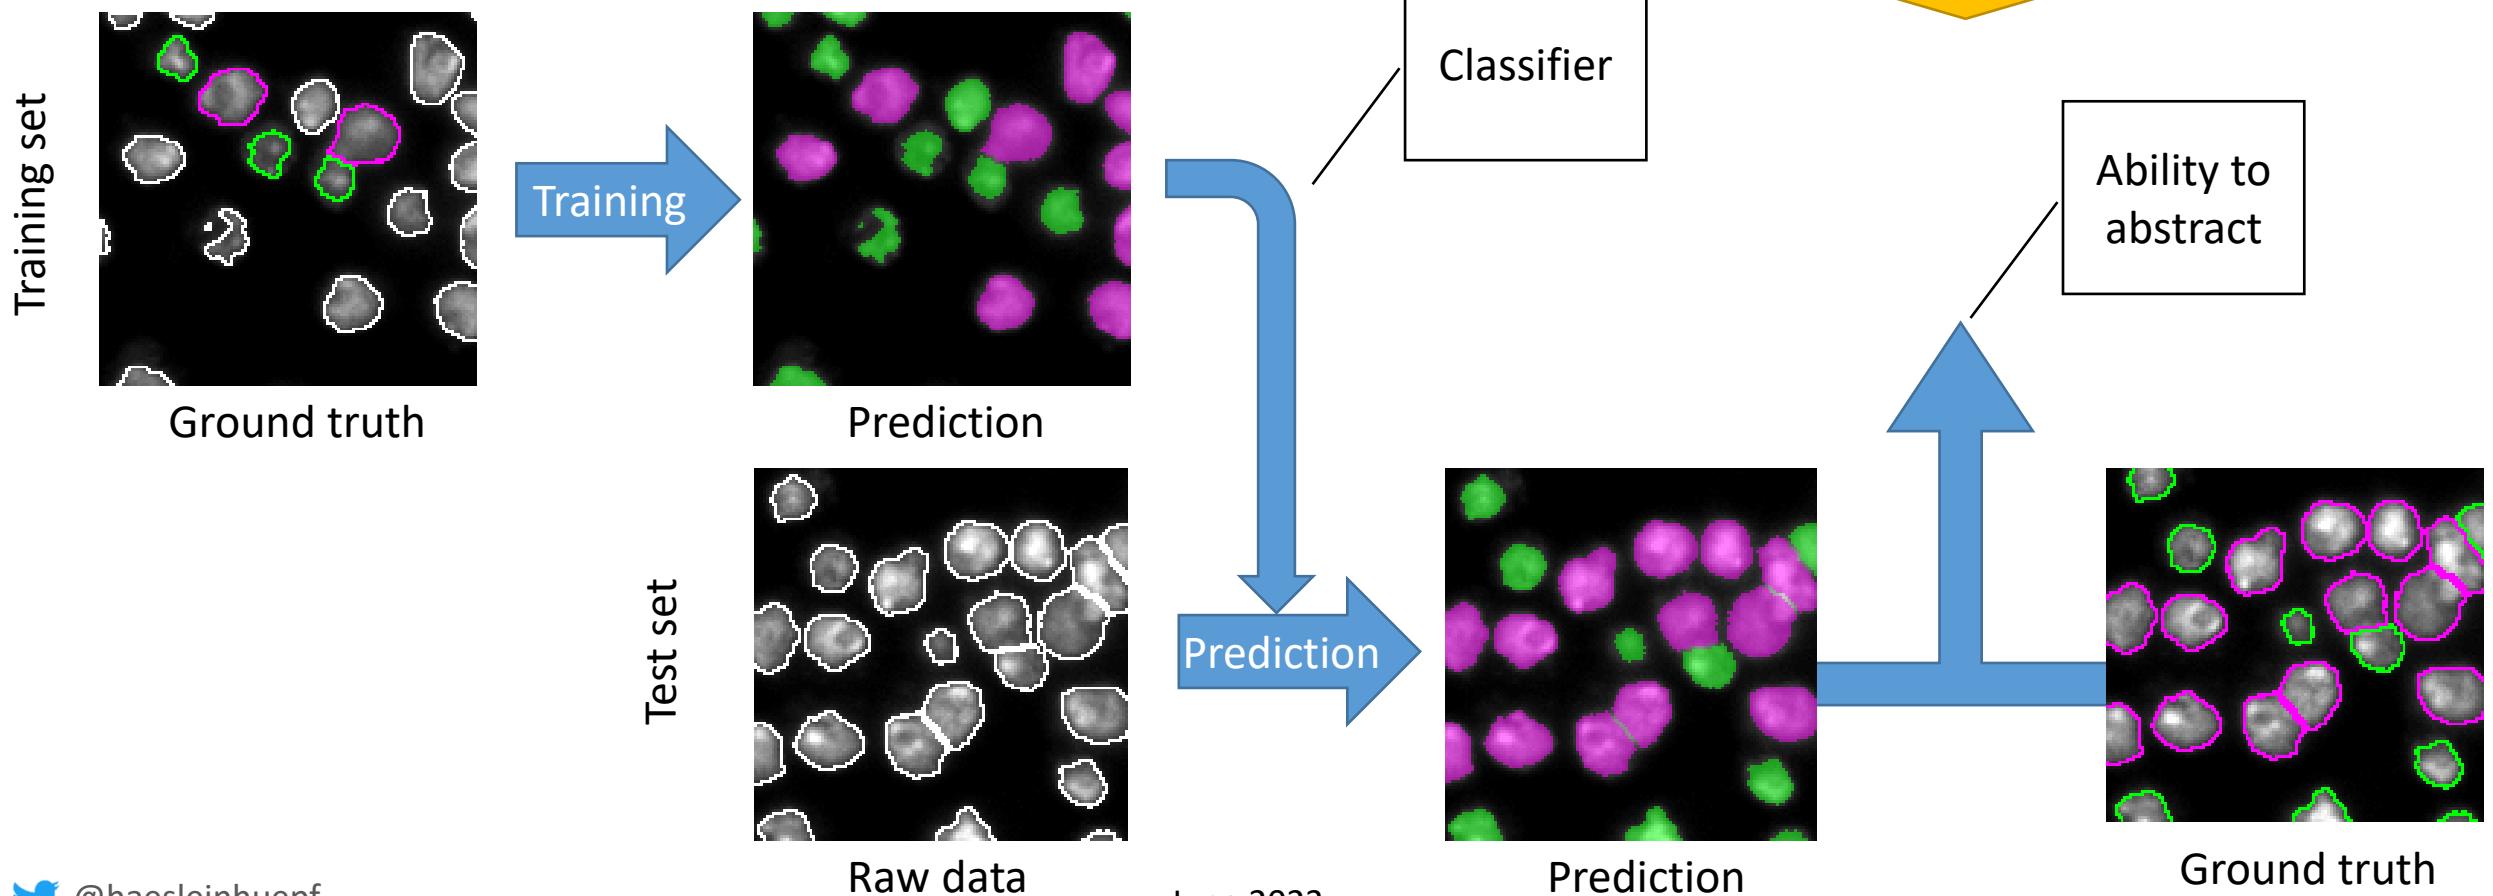
Number of trees: > 100



- Underfitting
 - A trained model that is not even able to properly process the data it was trained on
- Overfitting
 - A model that is able to process data it was trained on well
 - It processes other data poorly

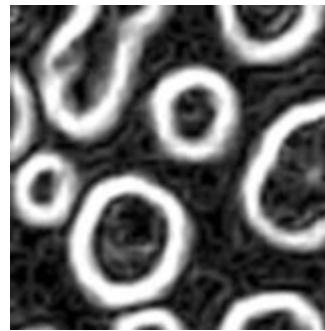
Model validation

- A good classifier is trained on a hand full of datasets and works on thousands similarly well.
- In order to assess that, we split the ground truth into two set
 - Training set (50%-90% of the available data)
 - Test set (10%-50% of the available data)

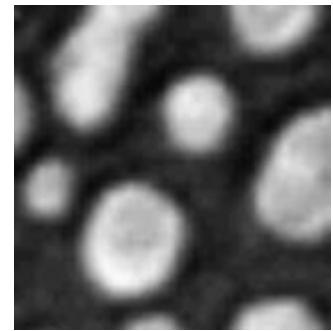
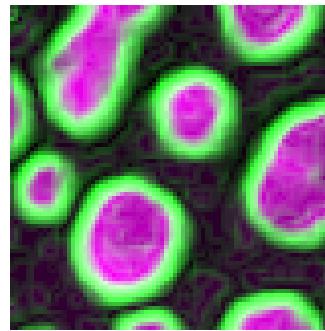
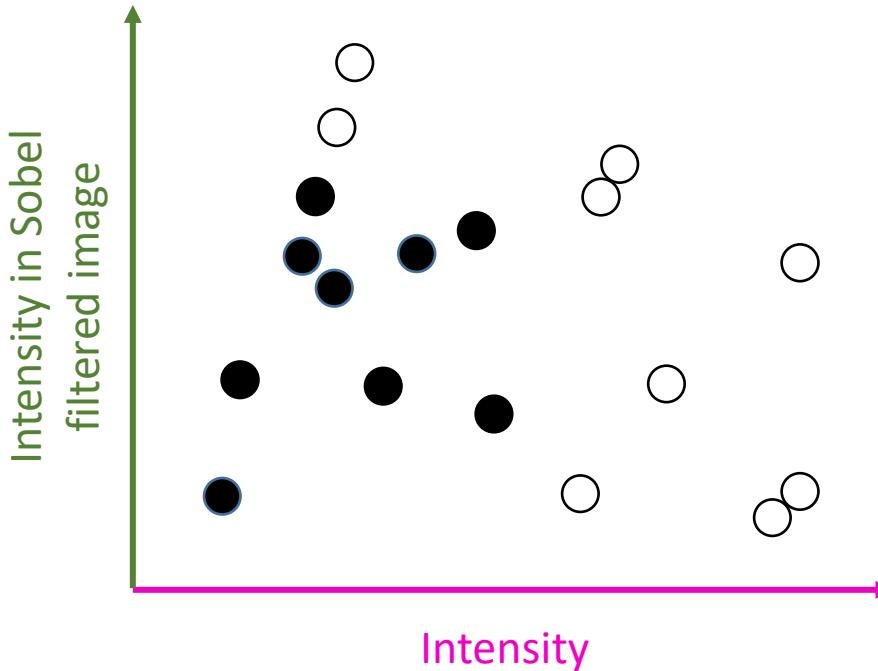


Object classification

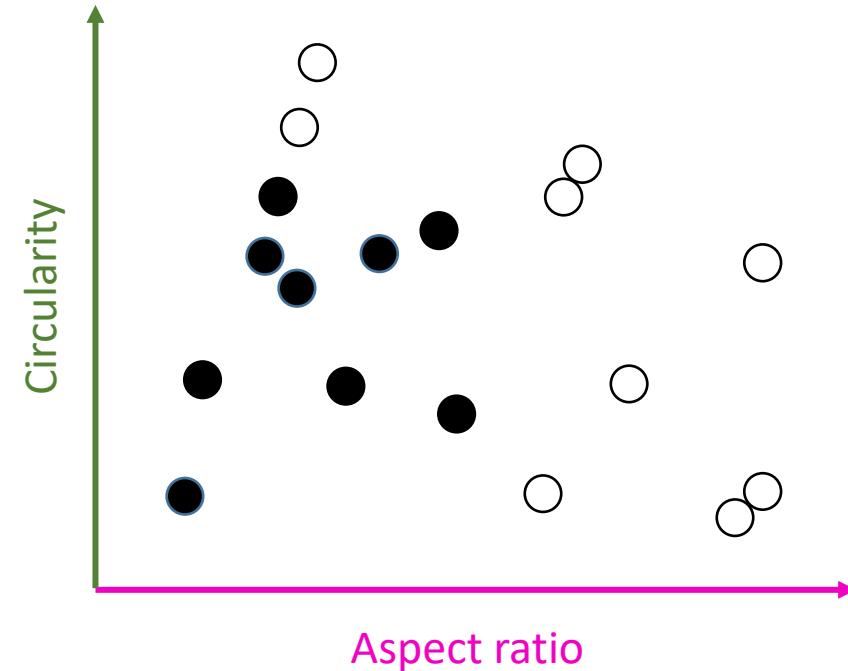
- What if we exchange pixel features with object features?



Pixel classification



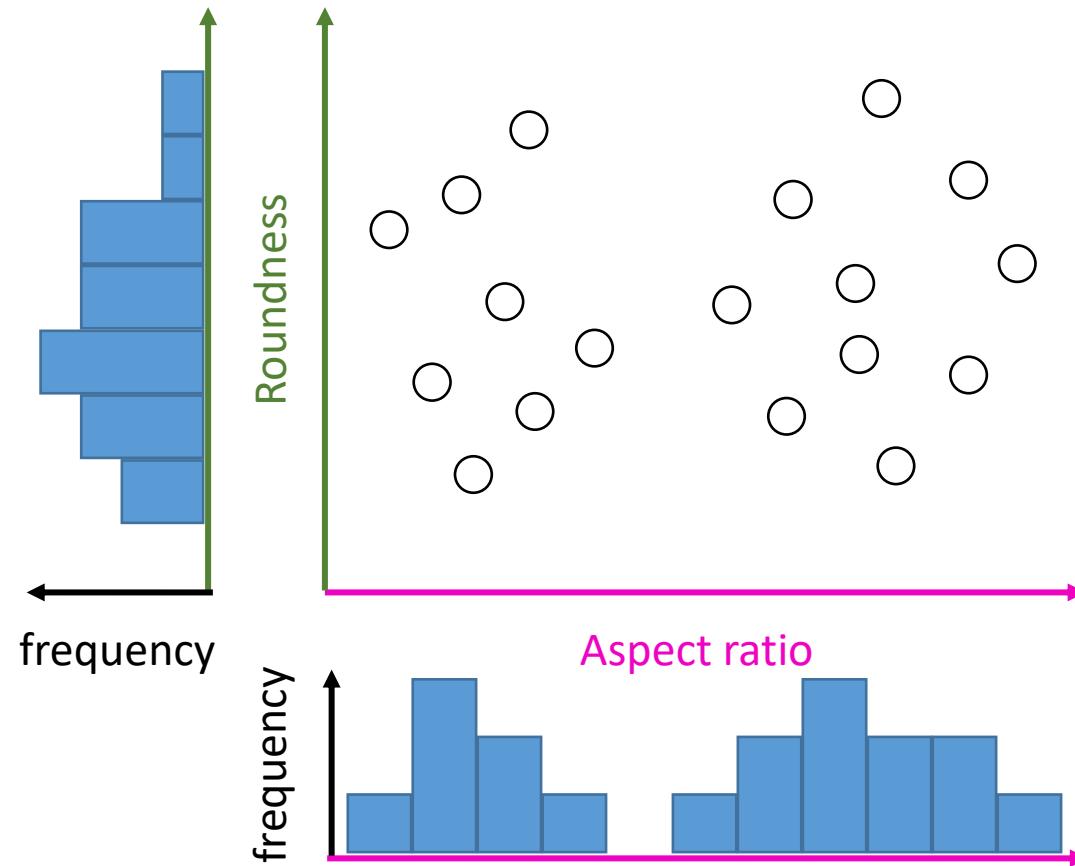
Object classification



- The algorithms work the same but with different
 - Features
 - Number of features
 - Tree / forest parameters
 - Selection criteria

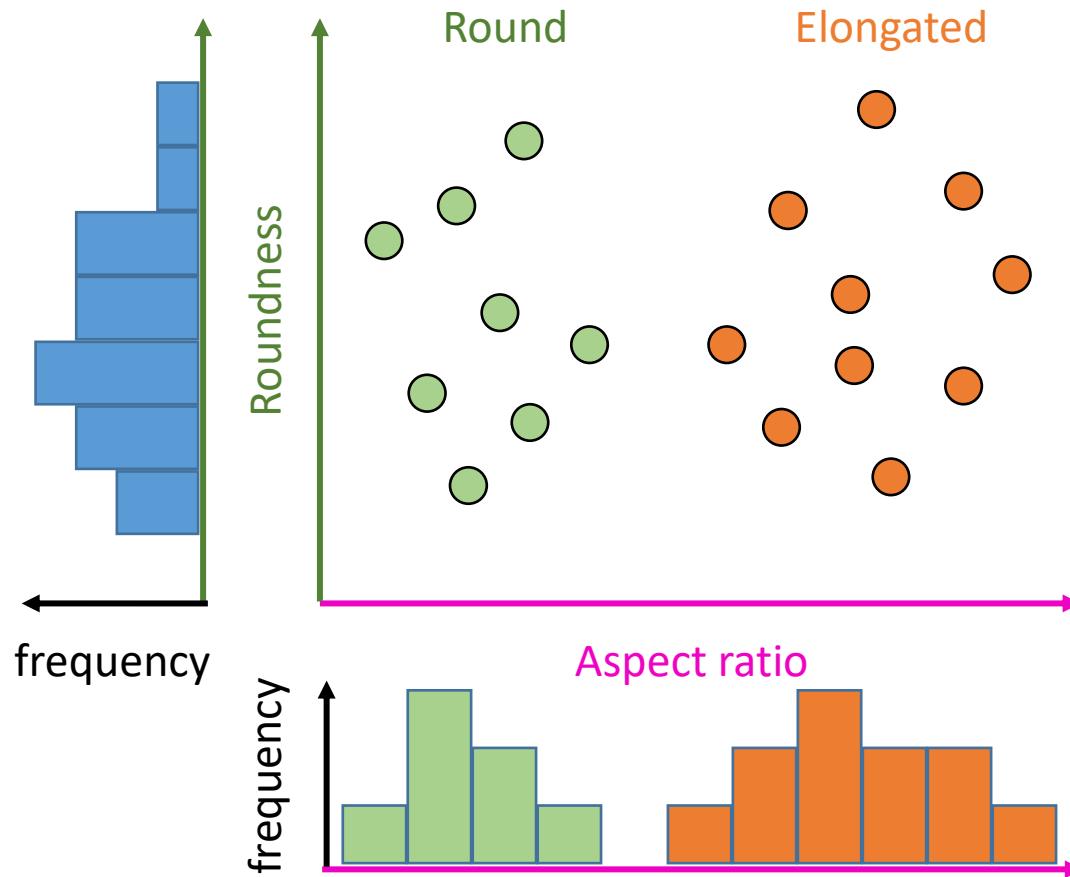
Unsupervised machine learning

- If you don't provide ground truth, the algorithm is *unsupervised*.



Unsupervised machine learning

- If you don't provide ground truth, the algorithm is unsupervised.
- Nevertheless, algorithms can tell us something about the data



Feature extraction

Robert Haase

With material from

Johannes Soltwedel, BiaPoL, PoL TU Dresden

Marcelo Zoccoler, BiaPol, PoL, TU Dresden

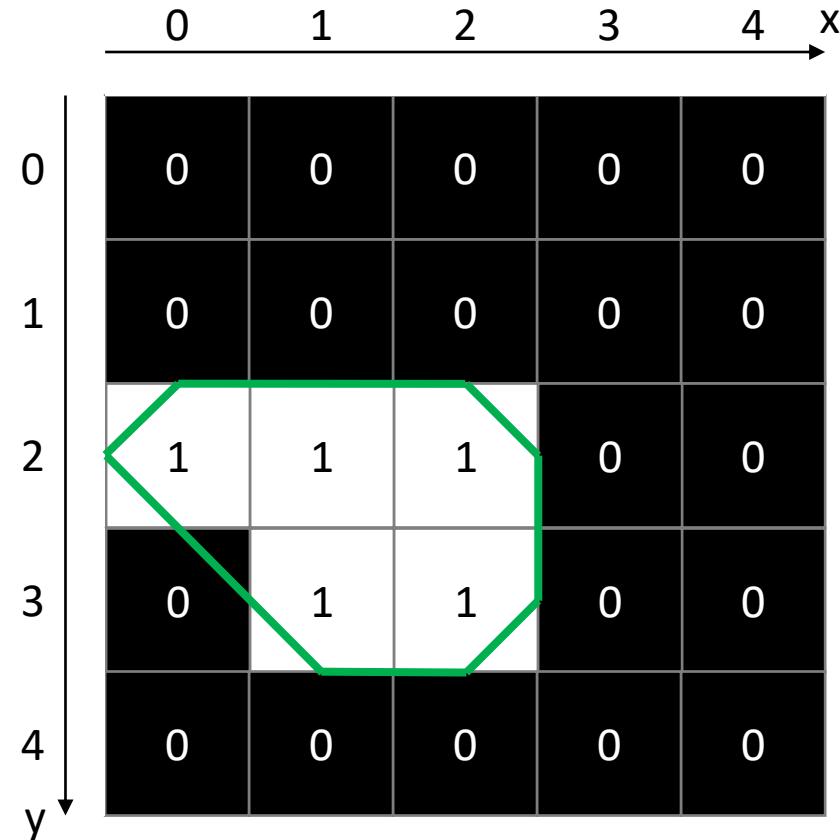
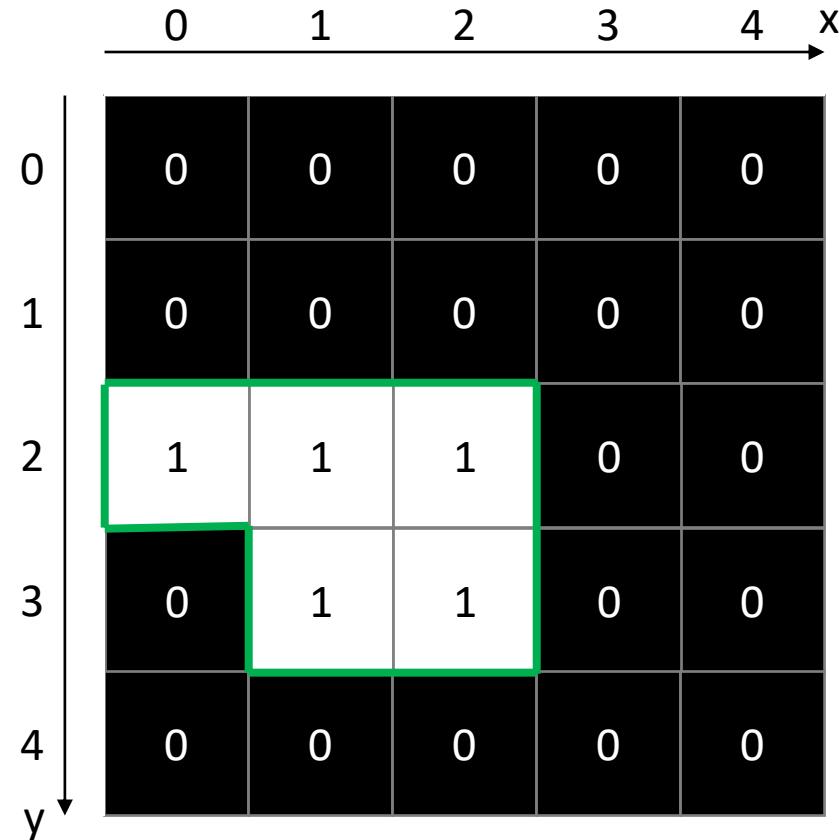
June 2023

These slides can be reused under the
[CC-BY 4.0](#) license unless mentioned otherwise.

- A *feature* is a countable or measurable property of an image or object.
- Goal of feature extraction is finding a minimal set of features to describe an object well enough to differentiate it from other objects.
- **Intensity based**
 - Mean intensity
 - Standard deviation
 - Total intensity
 - Textures
- **Shape based /spatial**
 - Area / Volume
 - Roundness
 - Solidity
 - Circularity / Sphericity
 - Elongation
 - Centroid
 - Bounding box
- **Spatio-temporal**
 - Displacement,
 - Speed,
 - Acceleration
- **Others**
 - Overlap
 - Colocalization
 - Neighborhood
- **Mixed features**
 - Center of mass
 - Local minima / maxima

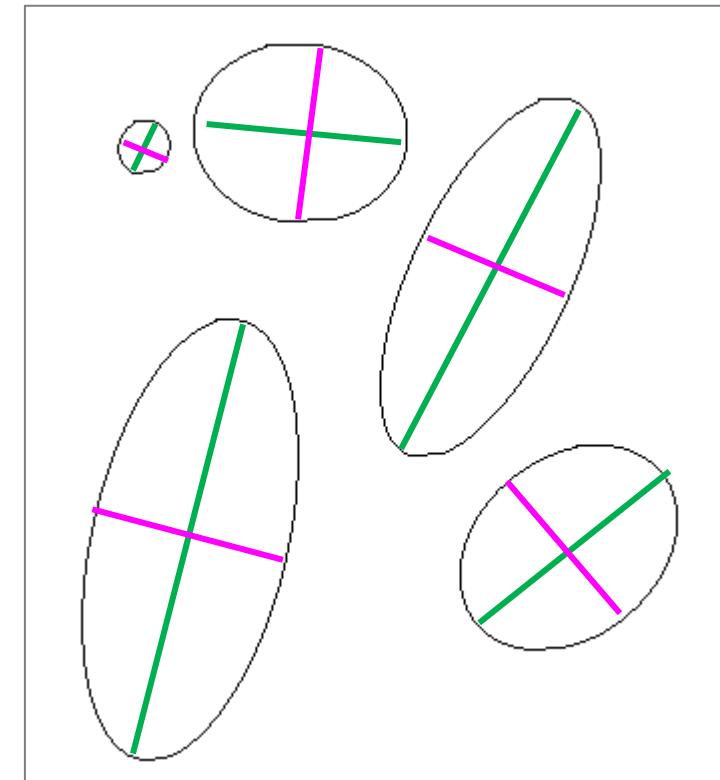
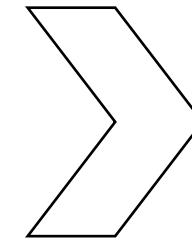
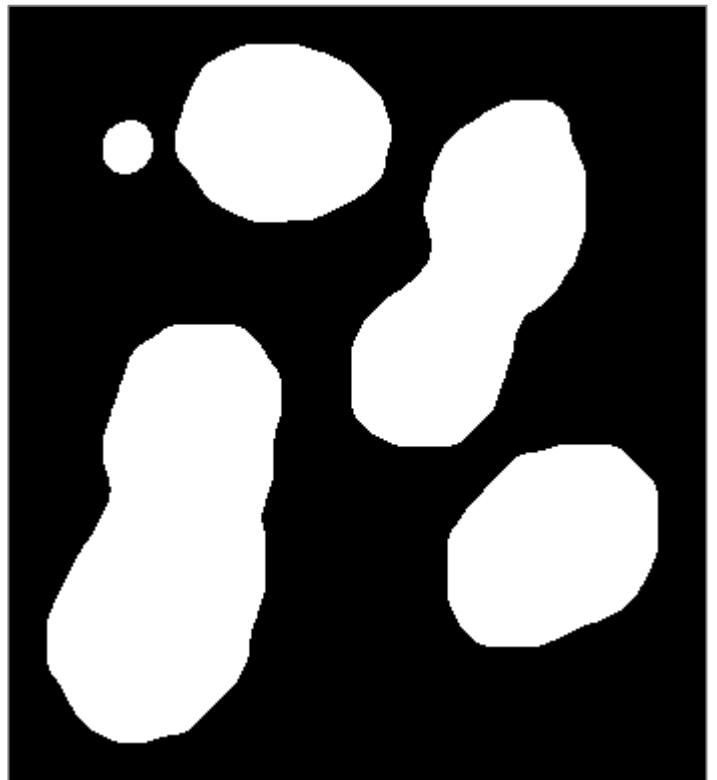


- Length of the outline around an object
- Depends on the actual implementation



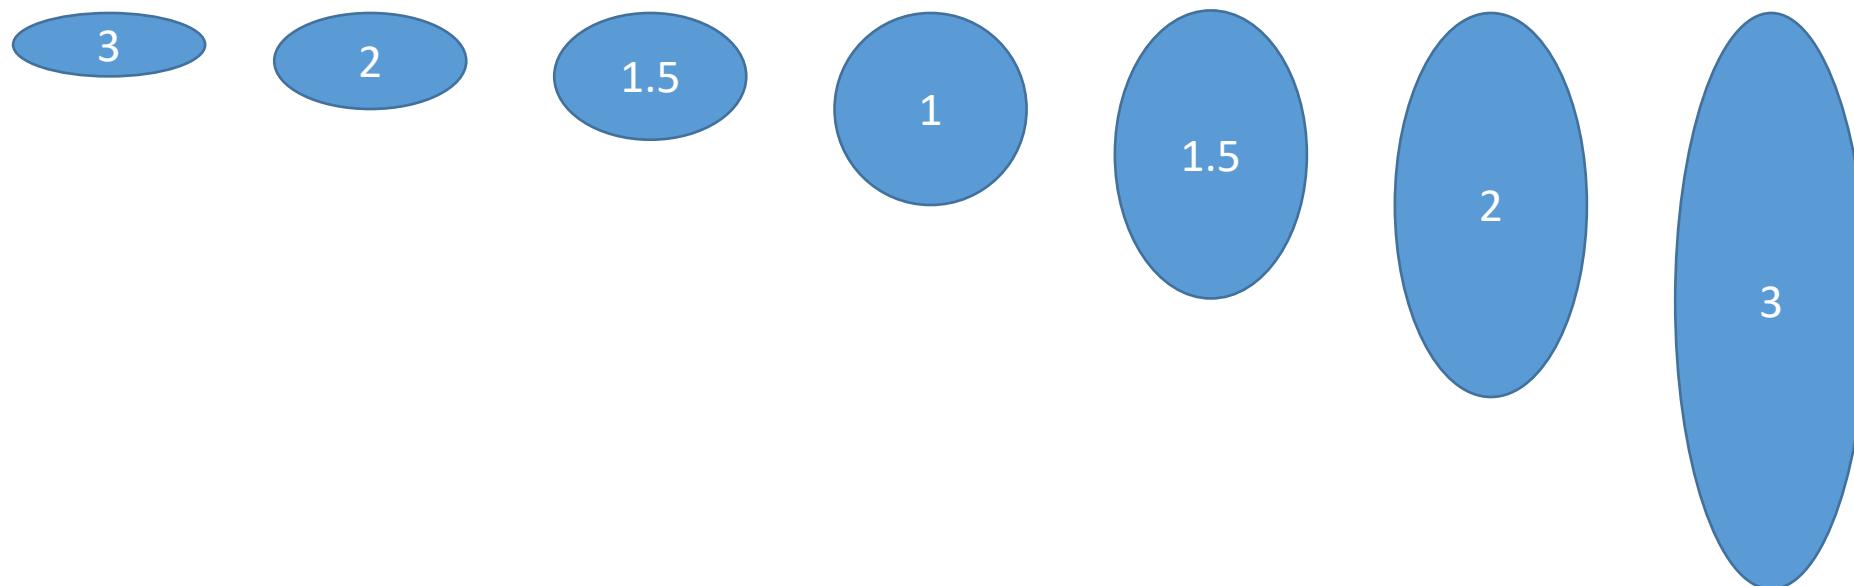
Fit ellipse

- For every object, find the optimal ellipse simplifying the object.
- Major axis ... long diameter
- Minor axis ... short diameter
- Major and minor axis are perpendicular to each other



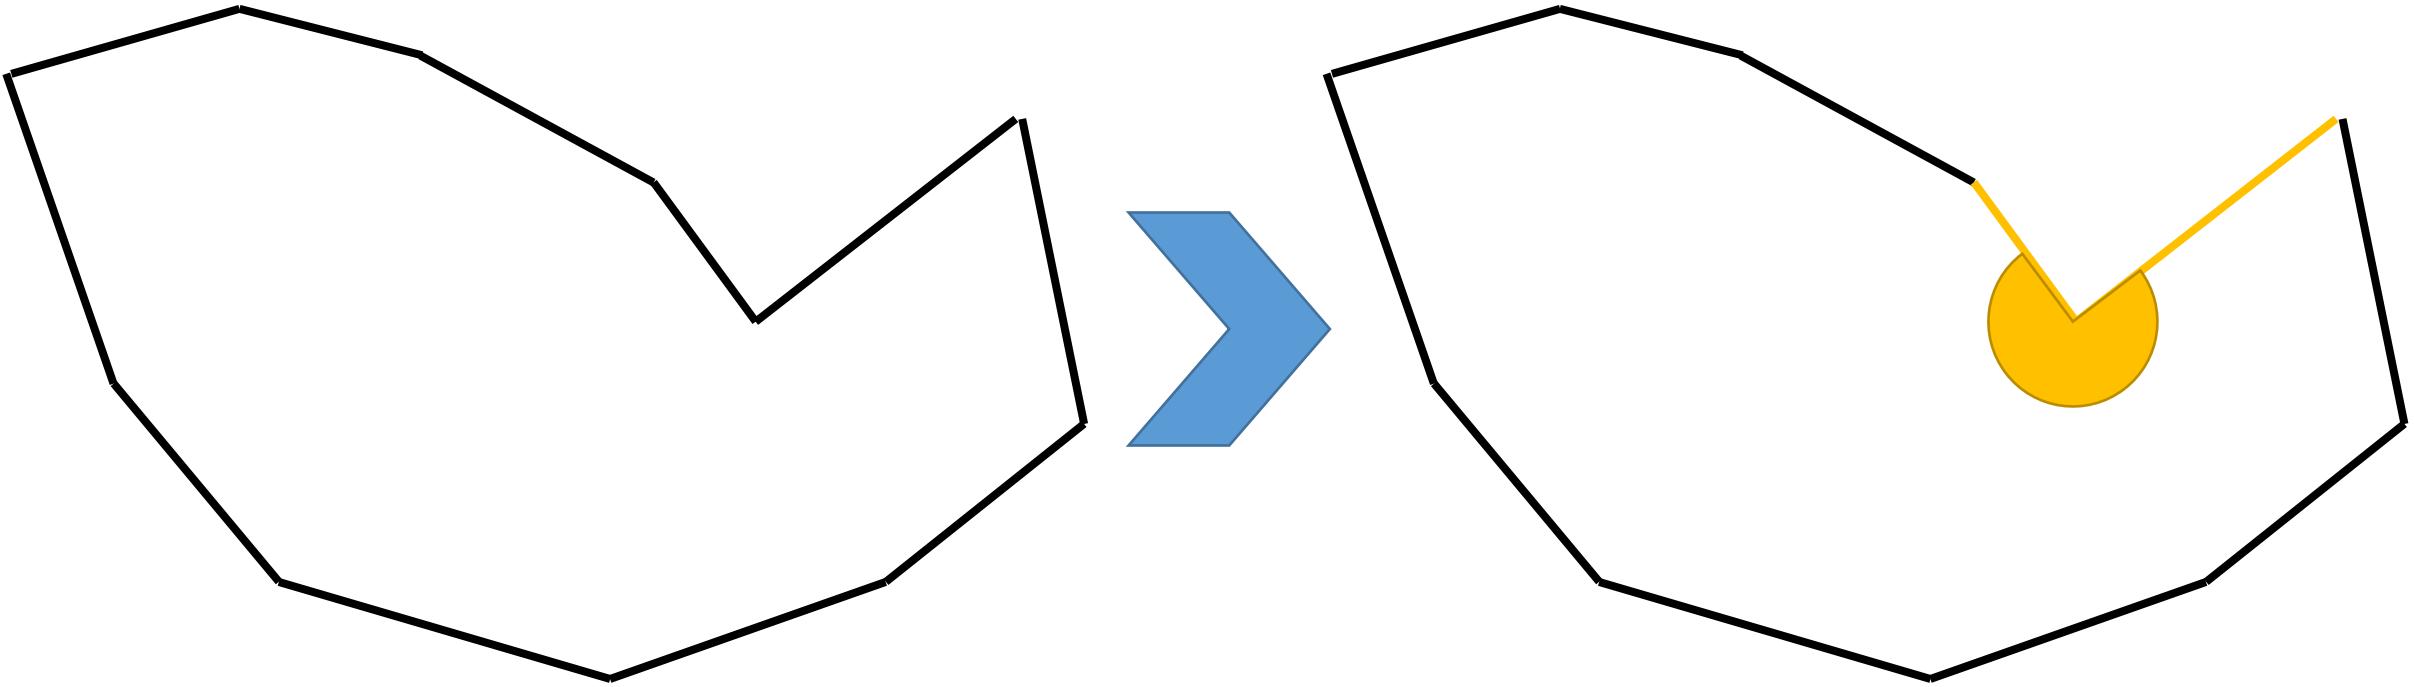
- The aspect ratio describes the elongation of an object.

$$AR = \text{major} / \text{minor}$$



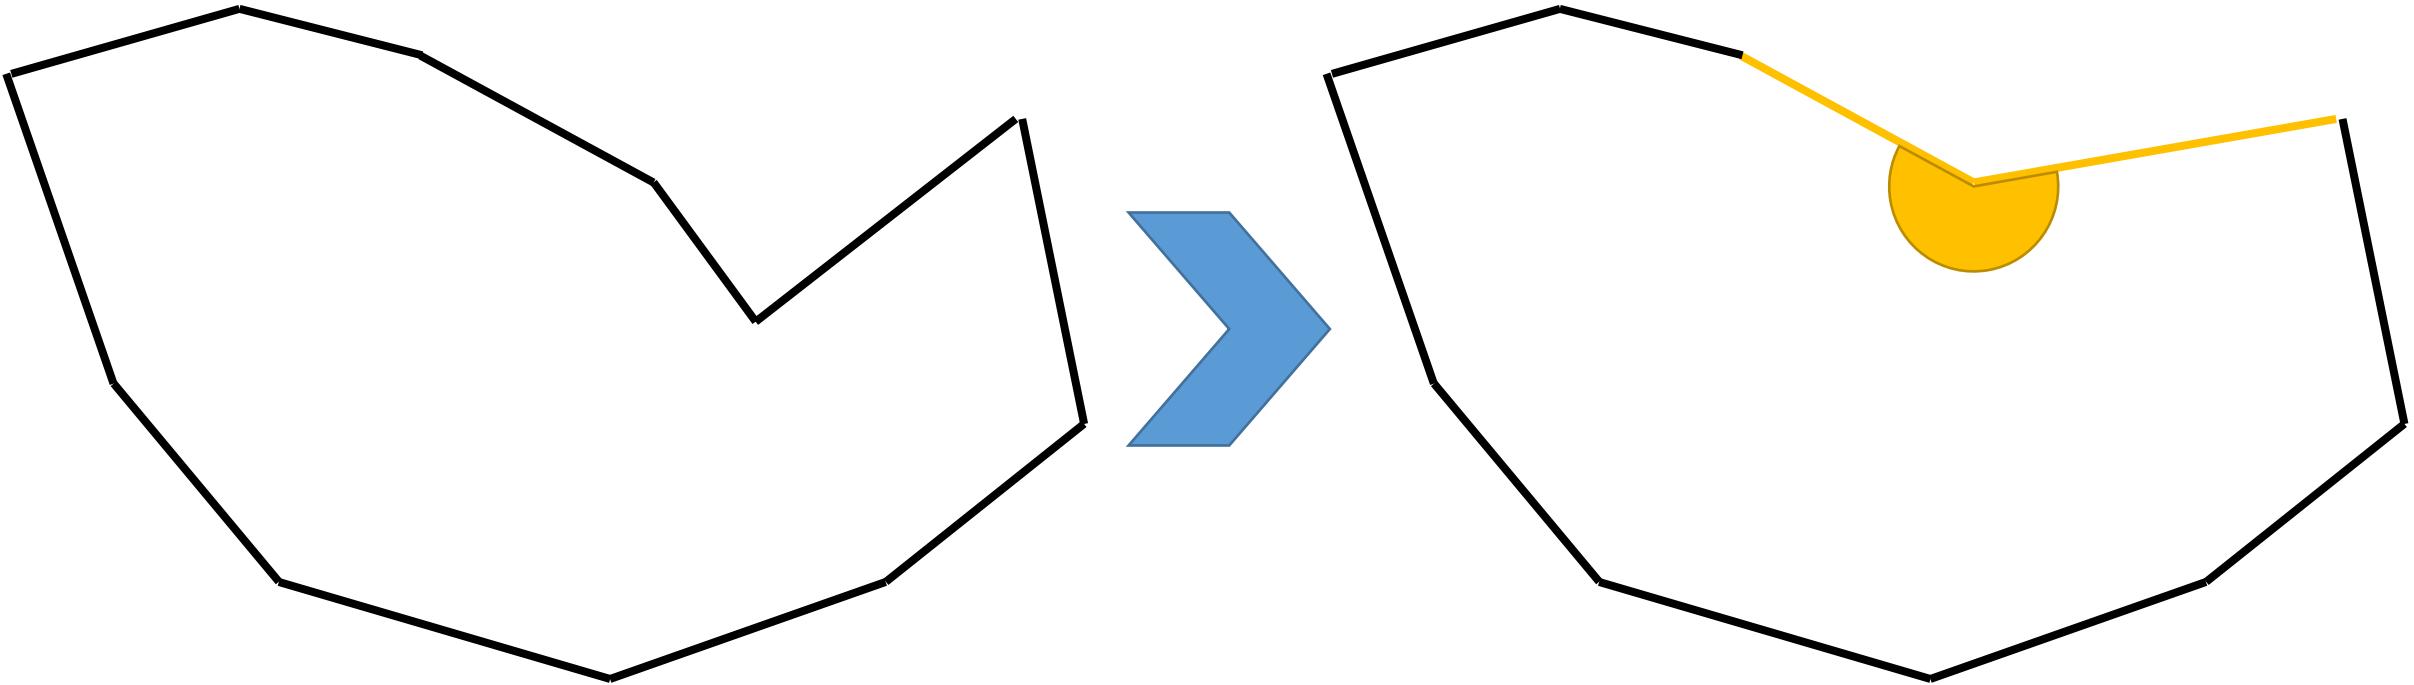
Convex hull

- By removing all concave corners of an object, we retrieve its **convex hull**.



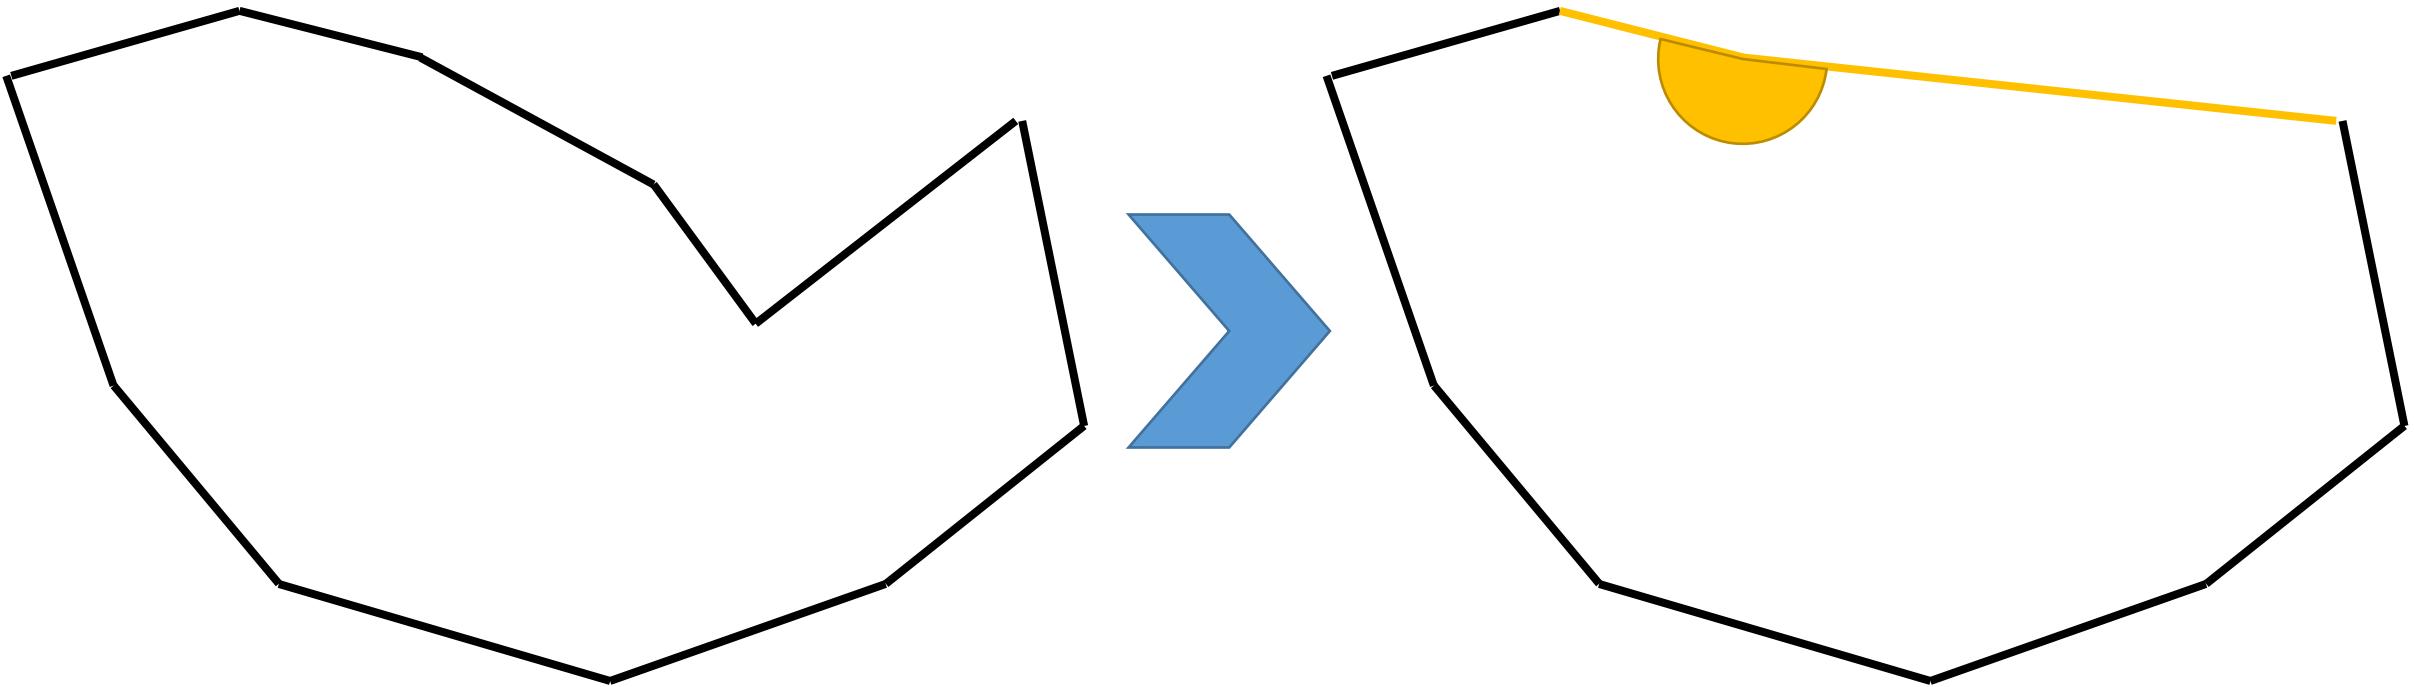
Convex hull

- By removing all concave corners of an object, we retrieve its **convex hull**.

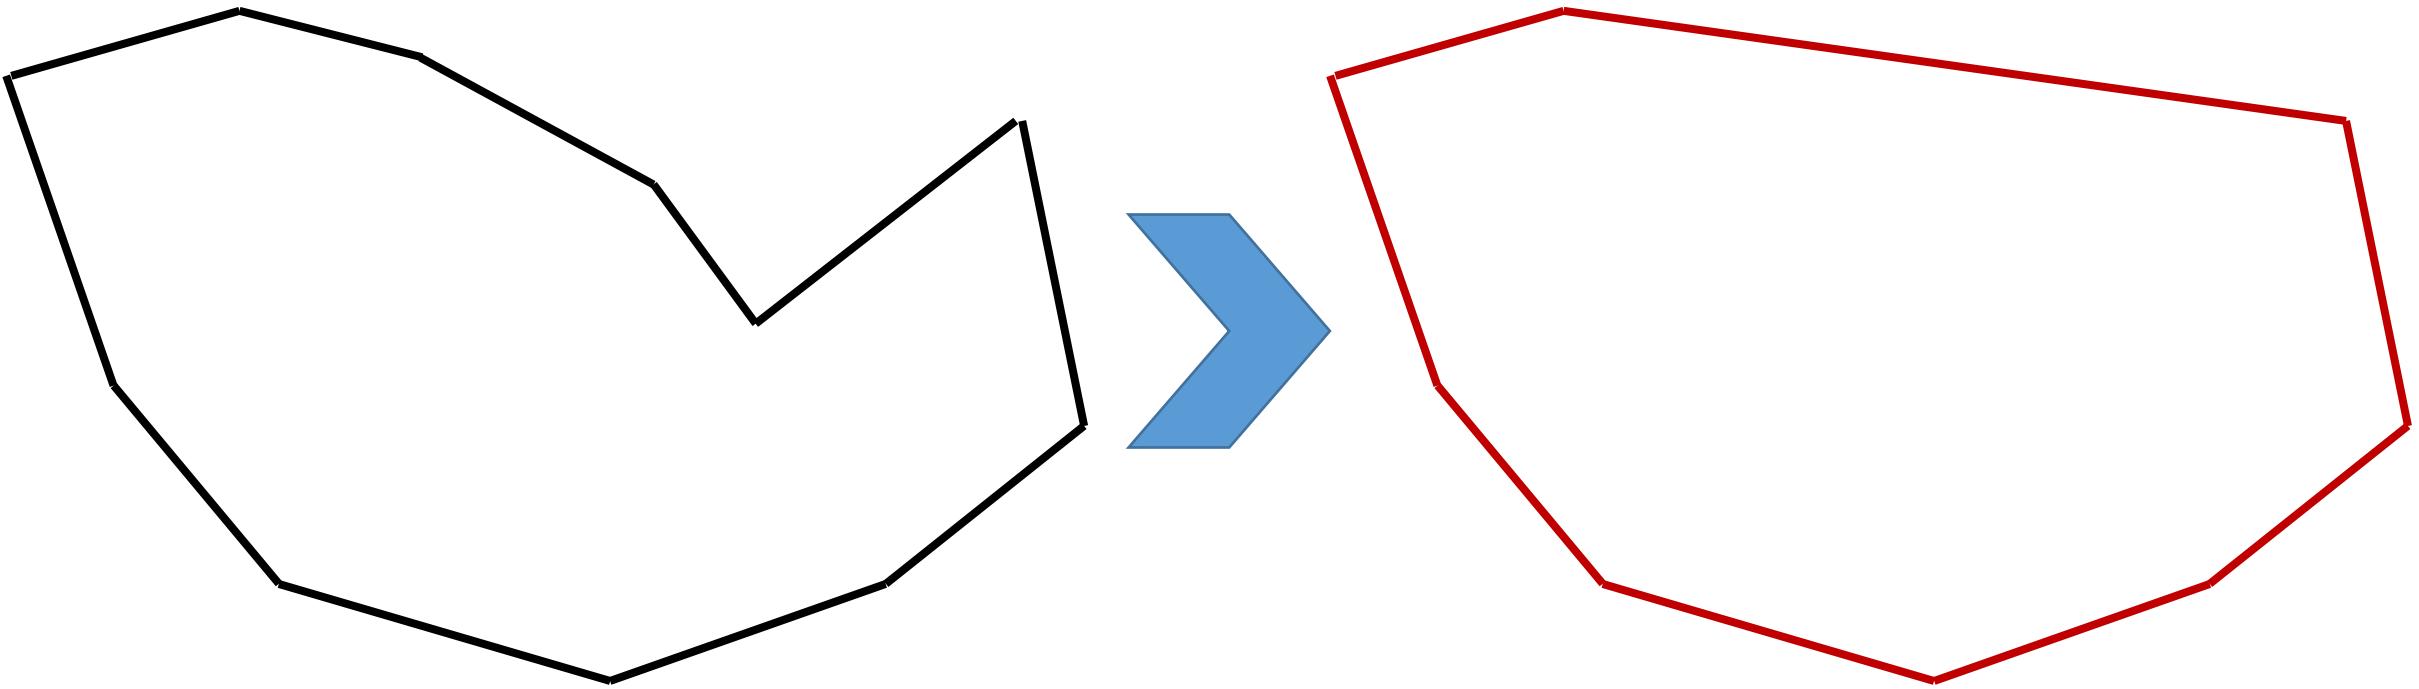


Convex hull

- By removing all concave corners of an object, we retrieve its **convex hull**.



- By removing all concave corners of an object, we retrieve its **convex hull**.



$$solidity = \frac{A}{A_{convexHull}}$$

Roundness and circularity

- The definition of a circle leads us to measurements of circularity and roundness.
- In case you use these measures, define them correctly. They are not standardized!

Diameter

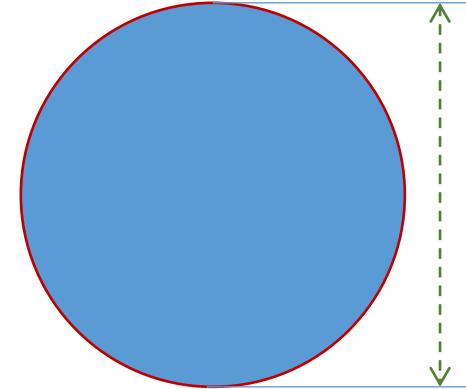
Circumference

Area

d

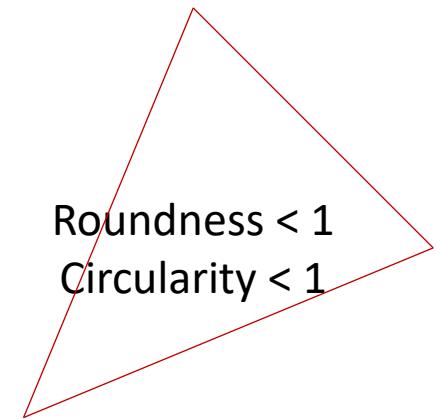
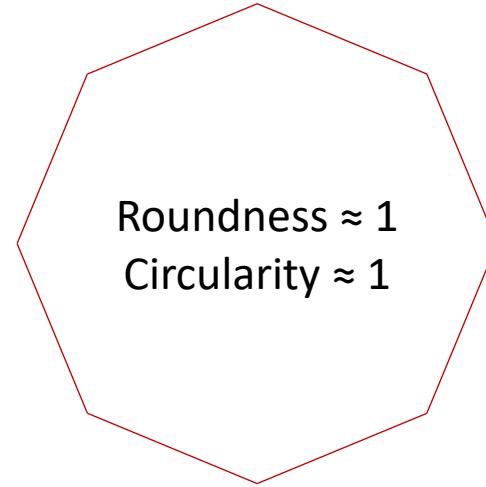
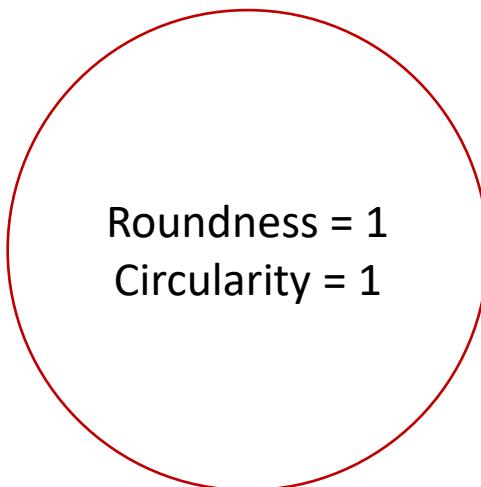
$C = \pi d$

$$A = \frac{\pi d^2}{4}$$



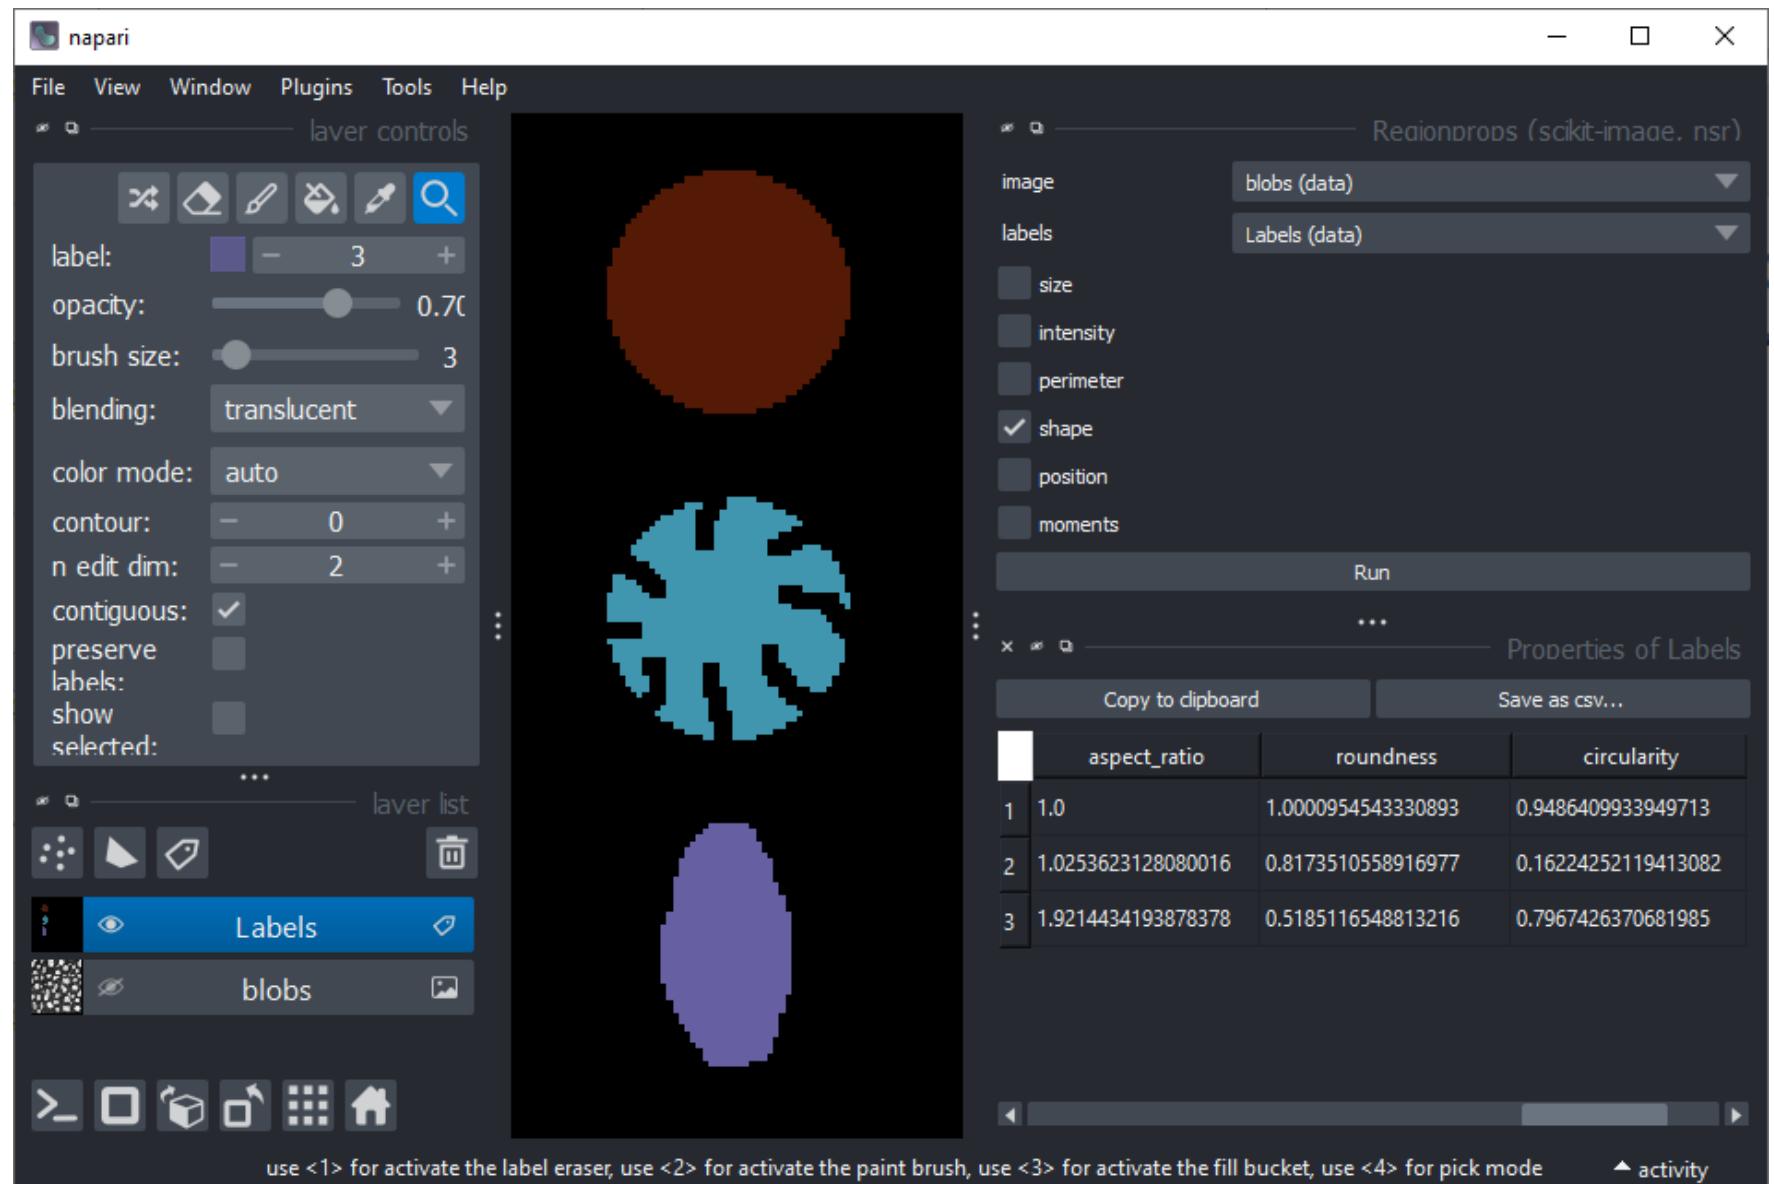
$$\text{roundness} = \frac{4 * A}{\pi \text{ major}^2}$$

$$\text{circularity} = \frac{4\pi * A}{\text{perimeter}^2}$$

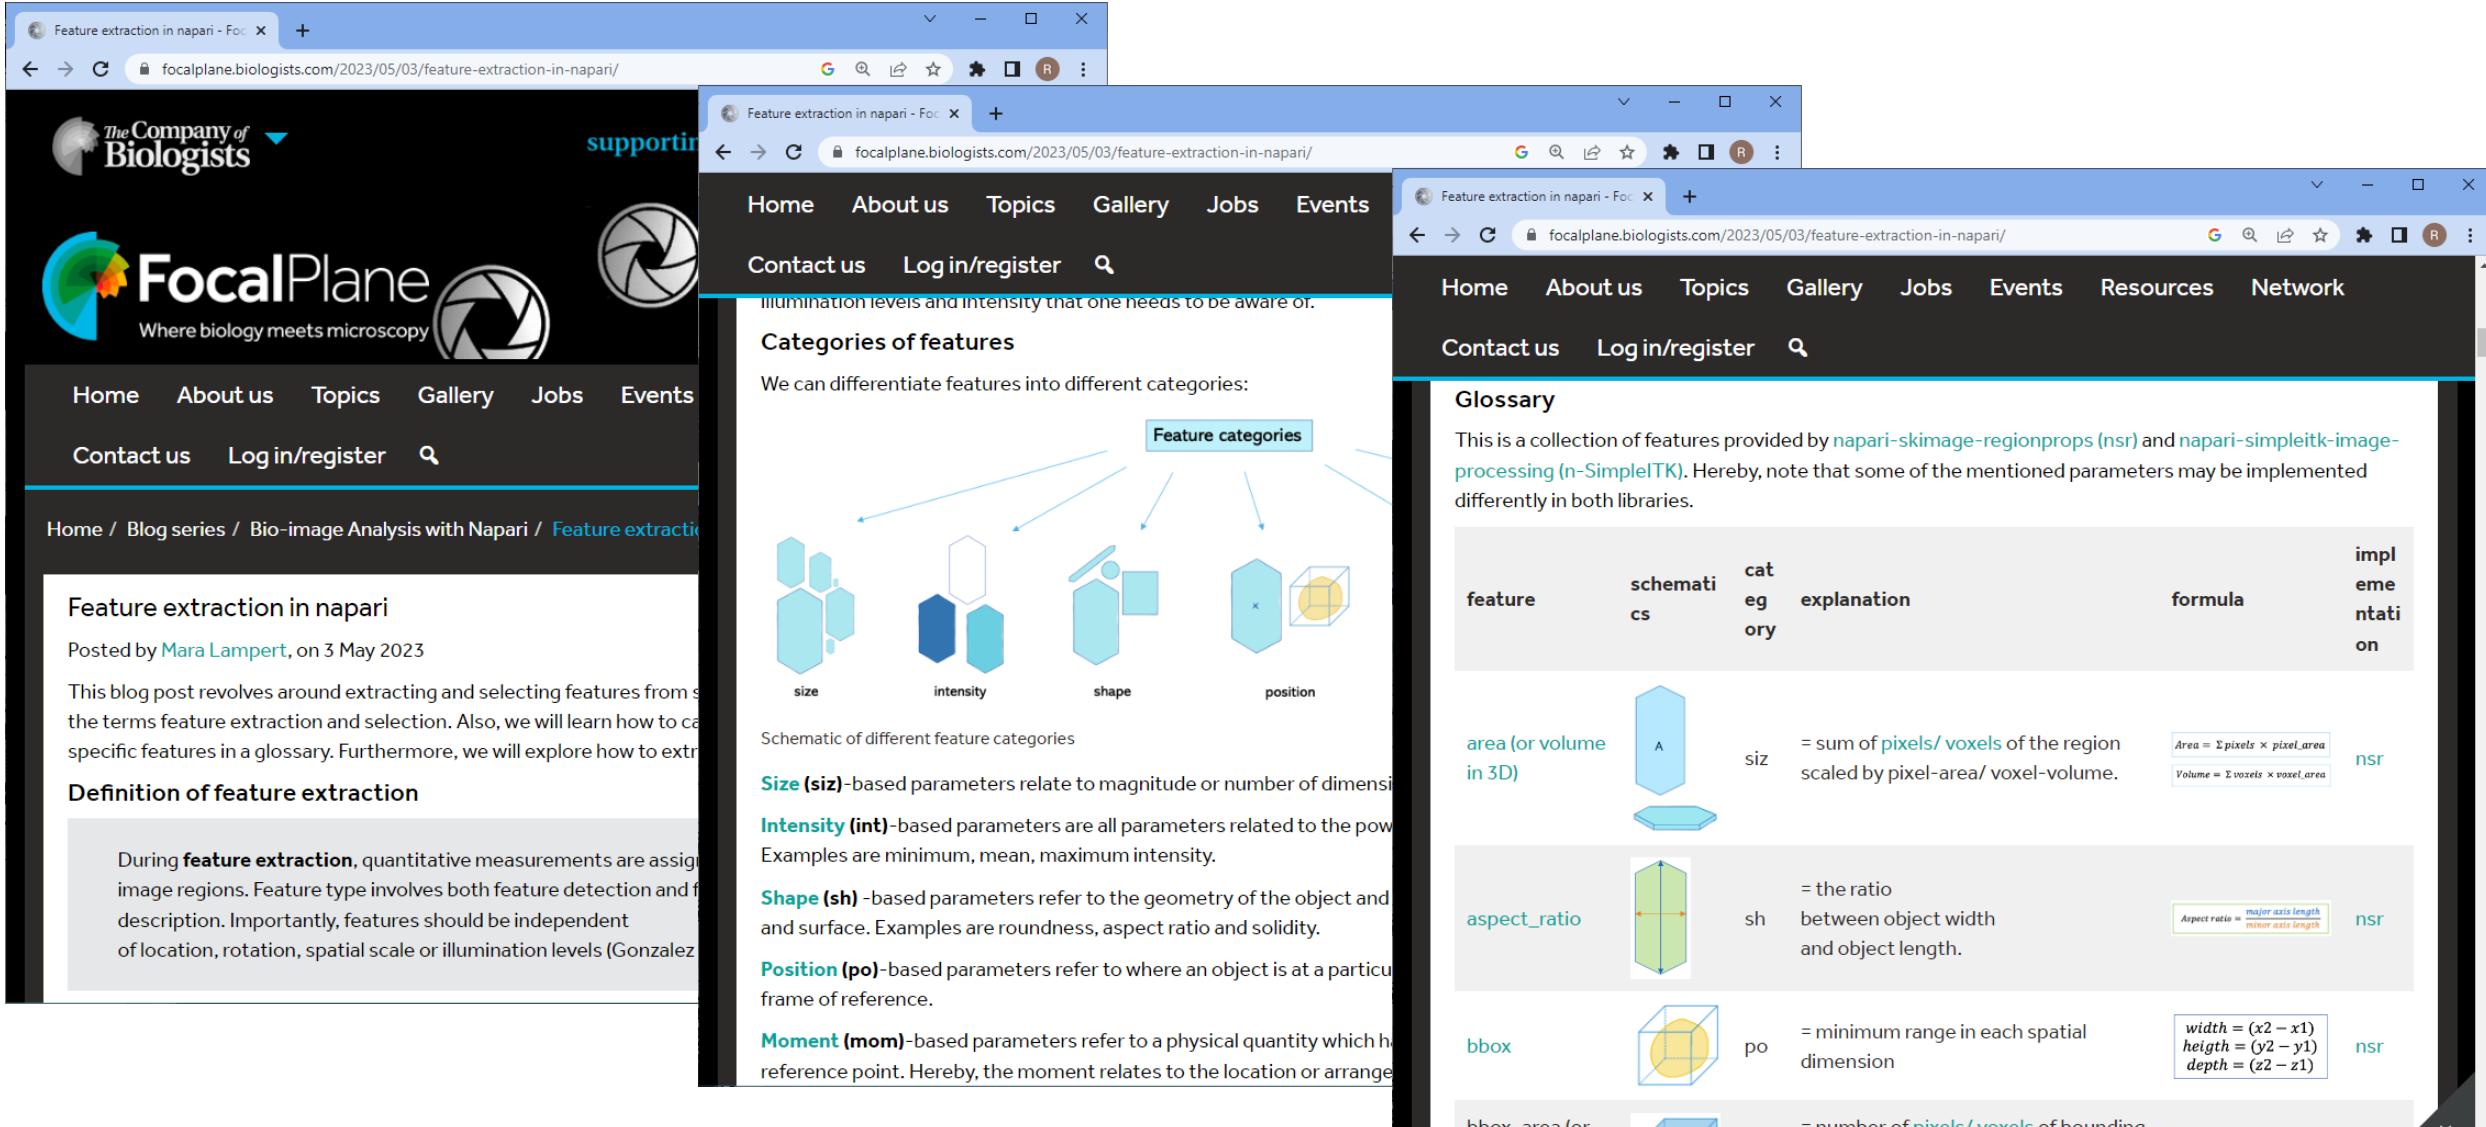


What's the difference between circularity and roundness?

- Draw representative shapes and study measurements



Further reading



The figure consists of three side-by-side screenshots of a web browser displaying a blog post. The left screenshot shows the main navigation bar of the FocalPlane website. The middle screenshot shows the content of the blog post, which includes a diagram of feature categories and text about size, intensity, shape, and position. The right screenshot shows a glossary table comparing parameters from napari-skimage-regionprops (nsr) and napari-simpleitk-image-processing (n-SimpleITK).

Categories of features

We can differentiate features into different categories:

- size**
- intensity**
- shape**
- position**

Schematic of different feature categories

Size (siz)-based parameters relate to magnitude or number of dimensions in 3D.

Intensity (int)-based parameters are all parameters related to the power of light.

Shape (sh)-based parameters refer to the geometry of the object and its surface.

Position (po)-based parameters refer to where an object is at a particular frame of reference.

Moment (mom)-based parameters refer to a physical quantity which has a reference point. Hereby, the moment relates to the location or arrangement of mass around a point.

feature	schematics	category	explanation	formula	implementation
area (or volume in 3D)		siz	= sum of pixels/ voxels of the region scaled by pixel-area/ voxel-volume.	$Area = \sum pixels \times pixel_area$ $Volume = \sum voxels \times voxel_area$	nsr
aspect_ratio		sh	= the ratio between object width and object length.	$Aspect\ ratio = \frac{major\ axis\ length}{minor\ axis\ length}$	nsr
bbox		po	= minimum range in each spatial dimension	$width = (x2 - x1)$ $height = (y2 - y1)$ $depth = (z2 - z1)$	nsr
bbox_area (or bbox_volume)		siz	= number of pixels/ voxels of bounding box scaled by pixel-area/ voxel-volume	$Bbox\ area = width \times height \times depth$	nsr

<https://focalplane.biologists.com/2023/05/03/feature-extraction-in-napari/>

Pixel classification using scikit-learn

Robert Haase



June 2023

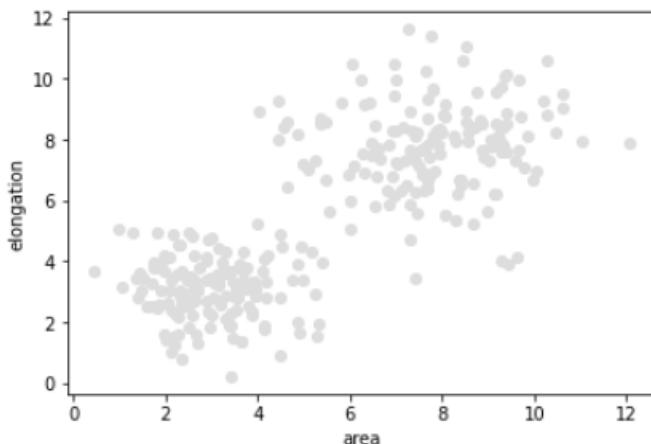
The scikit-learn logo is BSD3 licensed by the scikit-learn developers
https://commons.wikimedia.org/wiki/File:Scikit_learn_logo_small.svg

Tabular object classification

- Classify objects starting from feature vectors (table columns)

Raw data

	area	elongation
0	3.950088	2.848643
1	4.955912	3.390093
2	7.469852	5.575289
3	2.544467	3.017479
4	3.465662	1.463756
5	3.156507	3.232181
6	9.978705	6.676372
7	6.001683	5.047063
8	2.457139	3.416050
9	3.672295	3.407462
10	9.413702	7.598608



“Ground truth” annotation

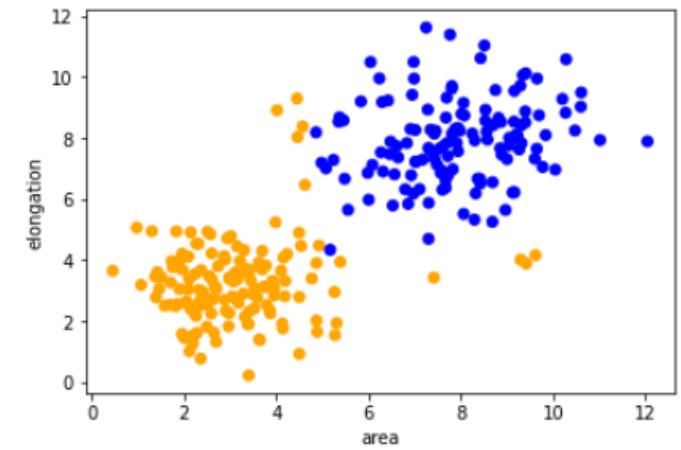
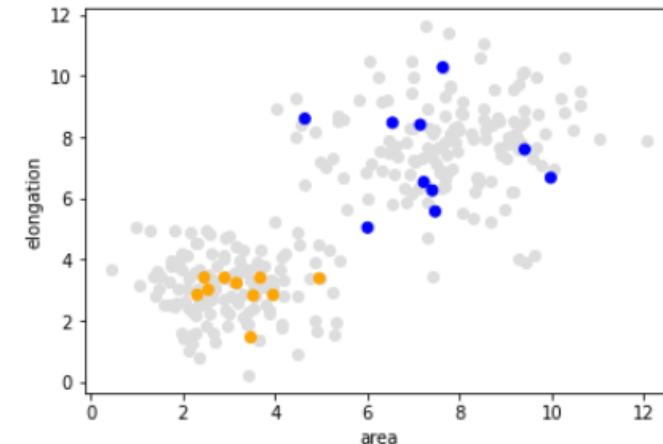
```
annotation = [1, 1, 2, 1, 1, 1, 2, 2,
```

```
classifier = RandomForestClassifier()  
classifier.fit(train_data, train_annotation)
```

Classifier training

Classifier prediction

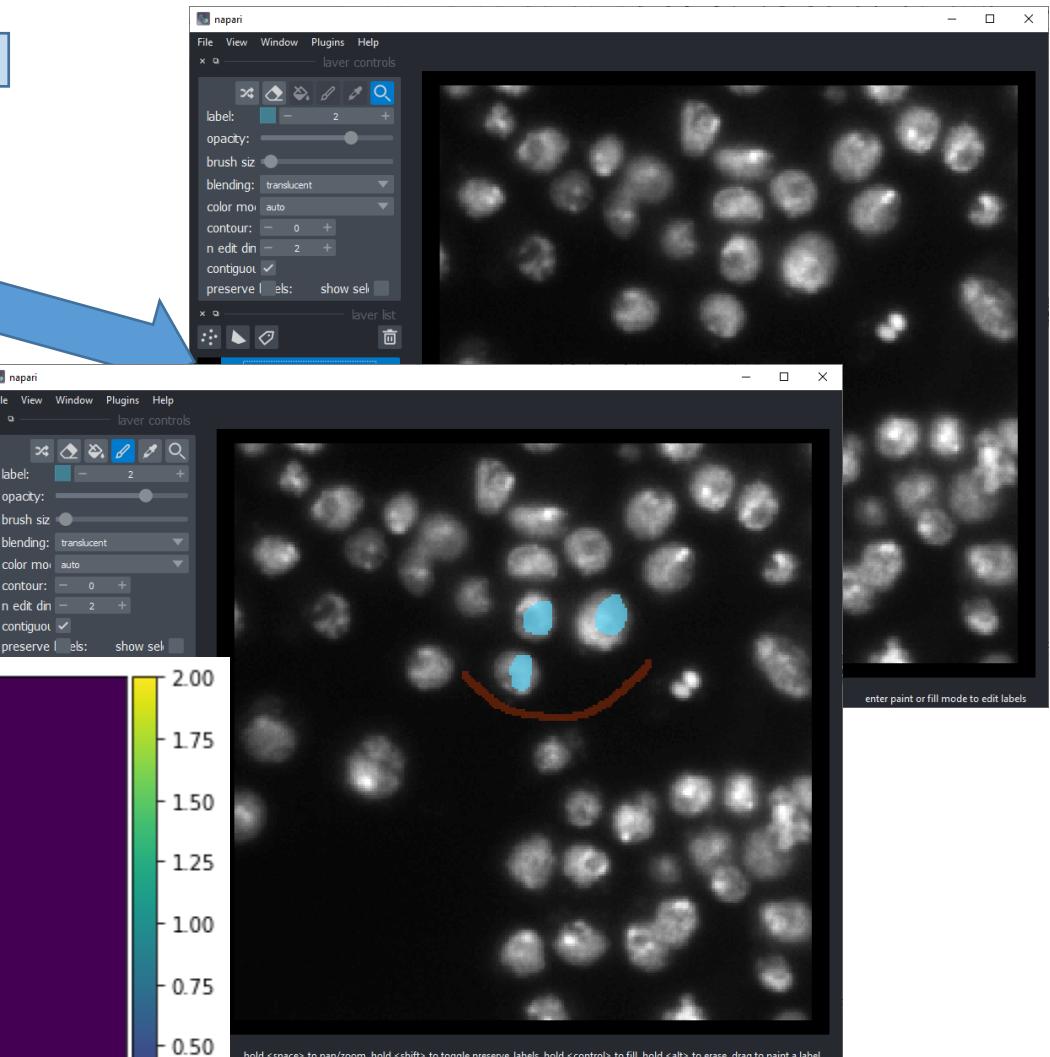
```
result = classifier.predict(validation_data)
```



Interactive pixel classification

- Prepare an empty layer for annotations and keep a reference

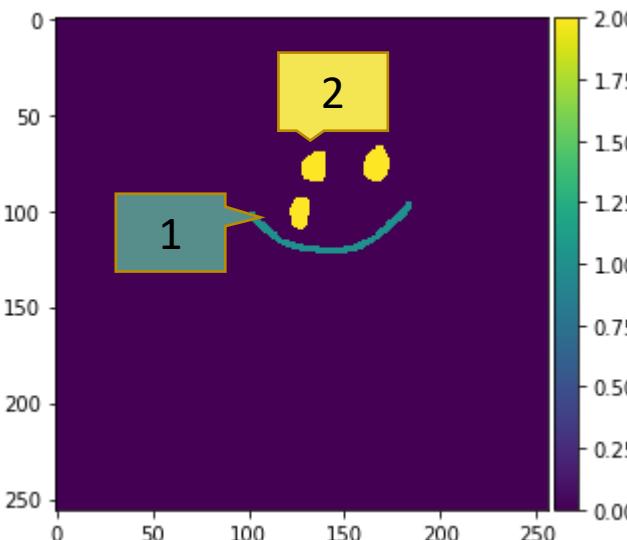
```
labels = viewer.add_labels(  
    np.zeros(image.shape).astype(int))
```



- Read annotations

```
manual_annotations = labels.data
```

```
from skimage.io import imshow  
  
imshow(manual_annotations,  
       vmin=0, vmax=2)
```



- Pixel classification using scikit-learn
 - Expects one-dimensional arrays for
 - every feature individually
 - ground truth

```
# train classifier
from sklearn.ensemble import RandomForestClassifier
classifier = RandomForestClassifier(max_depth=2, random_state=0)
classifier.fit(X, y)
```

Image data

Ground truth /
annotation

Image data

y_ = classifier.predict(X)

prediction

Interactive pixel classification

- Pixel classification using scikit-learn

- Expects one-dimensional arrays for
 - every feature individually
 - ground truth

```
# for training, we need to generate features
```

```
feature_stack = generate_feature_stack(image)
```

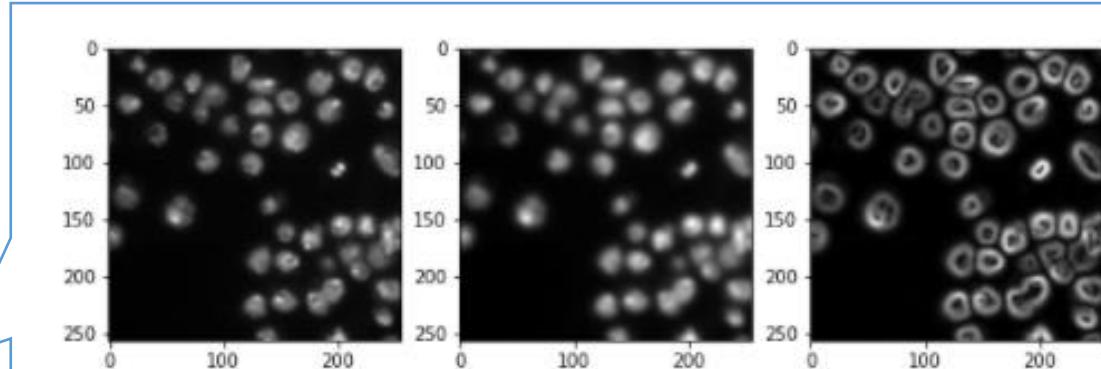
```
X, y = format_data(feature_stack, manual_annotations)
```

```
# train classifier
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
classifier = RandomForestClassifier(max_depth=2, random_state=0)
```

```
classifier.fit(X, y)
```

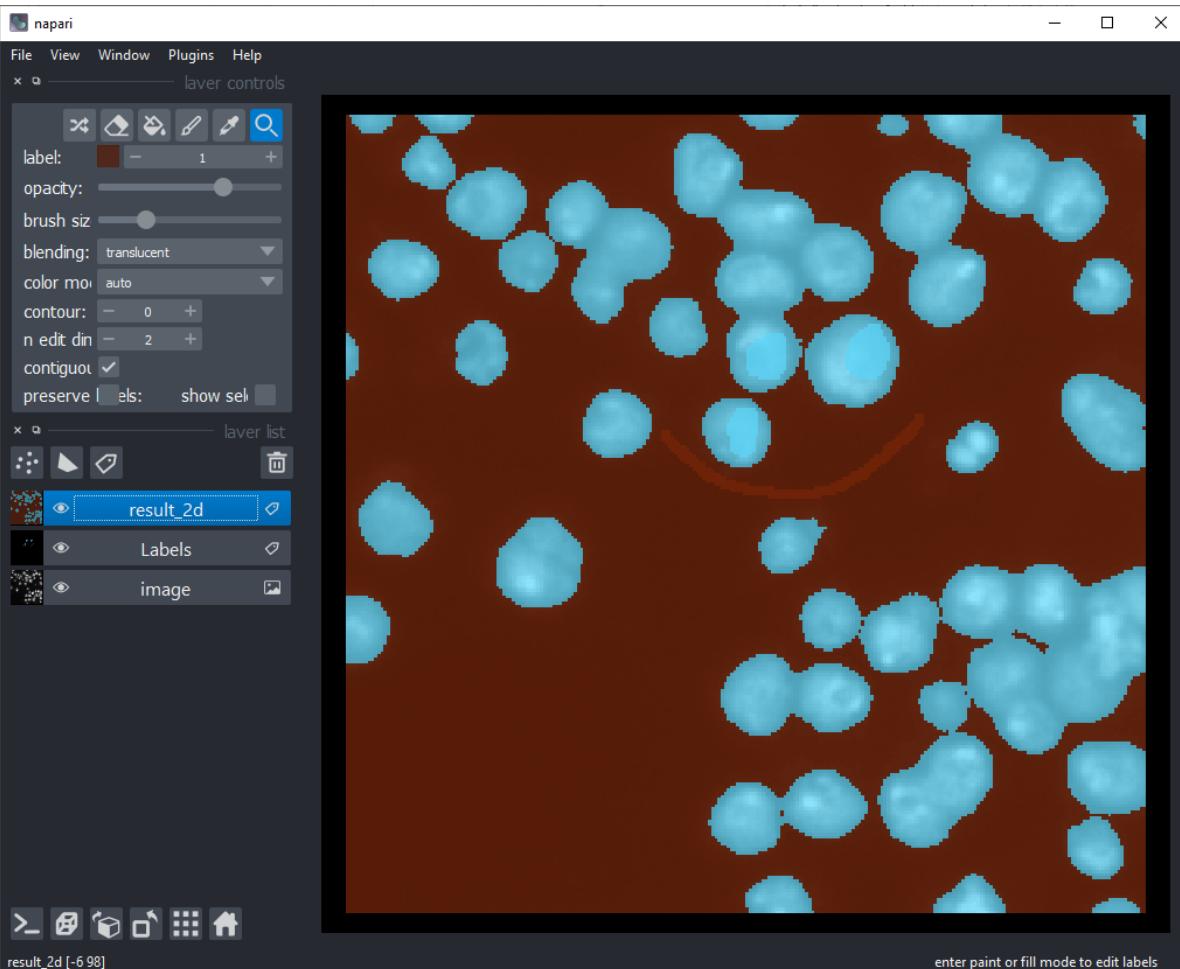


Interactive pixel classification

- Pixel classification using scikit-learn

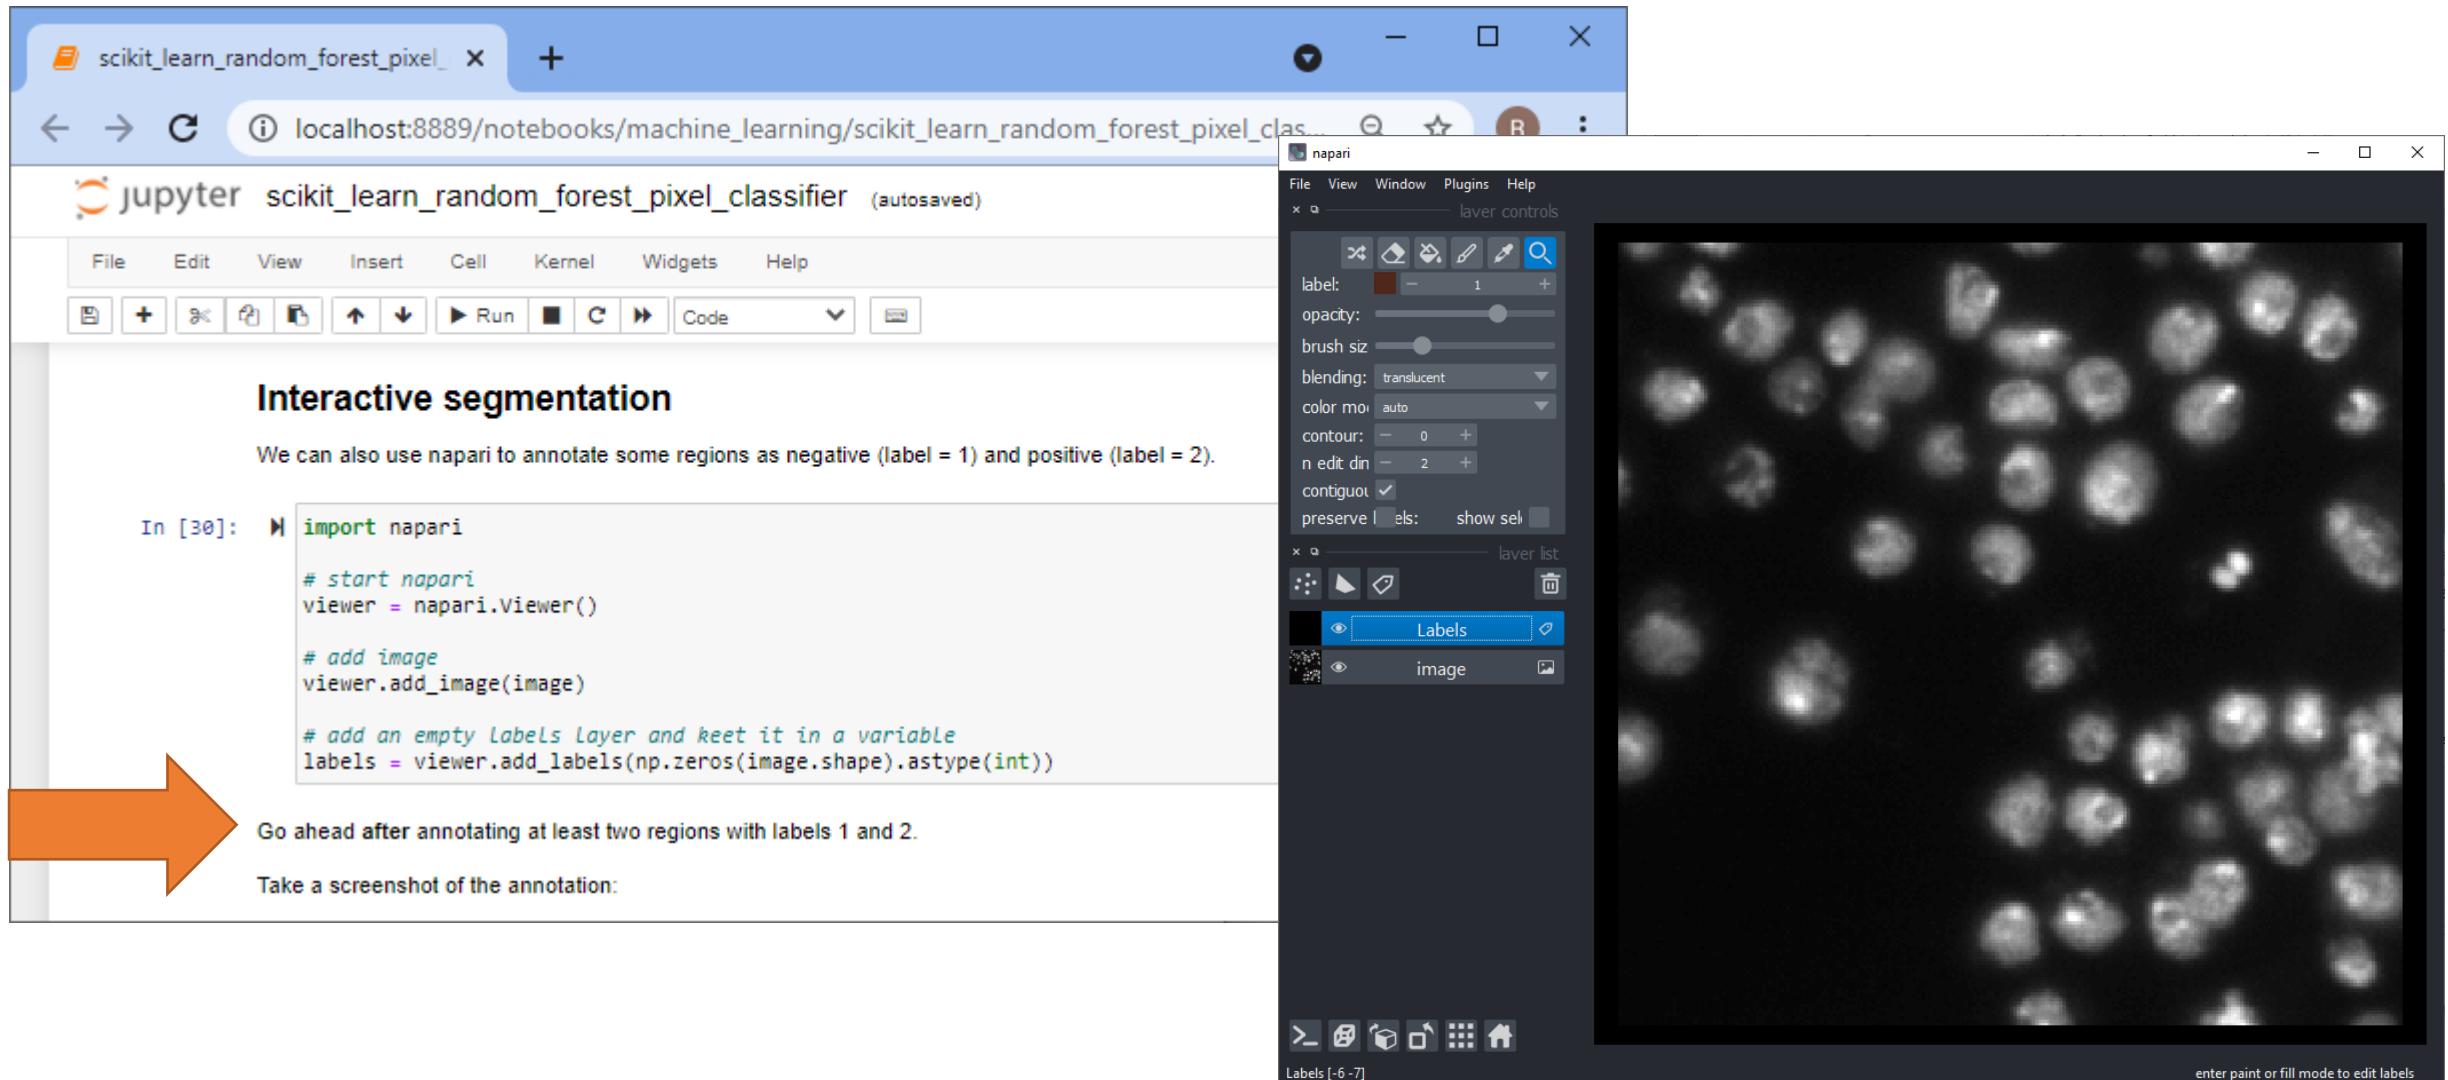
```
# process the whole image and show result
result_1d = classifier.predict(feature_stack.T)
result_2d = result_1d.reshape(image.shape)

viewer.add_labels(result_2d)
```



Interactive pixel classification

- Jupyter notebooks and napari side-by-side



The screenshot shows a Jupyter notebook interface and the napari image viewer running side-by-side.

Jupyter Notebook:

- Tab title: scikit_learn_random_forest_pixel_
- URL: localhost:8889/notebooks/machine_learning/scikit_learn_random_forest_pixel_clas...
- Code cell (In [30]):

```
import napari

# start napari
viewer = napari.Viewer()

# add image
viewer.add_image(image)

# add an empty Labels Layer and keep it in a variable
labels = viewer.add_labels(np.zeros(image.shape).astype(int))
```
- Text below code:

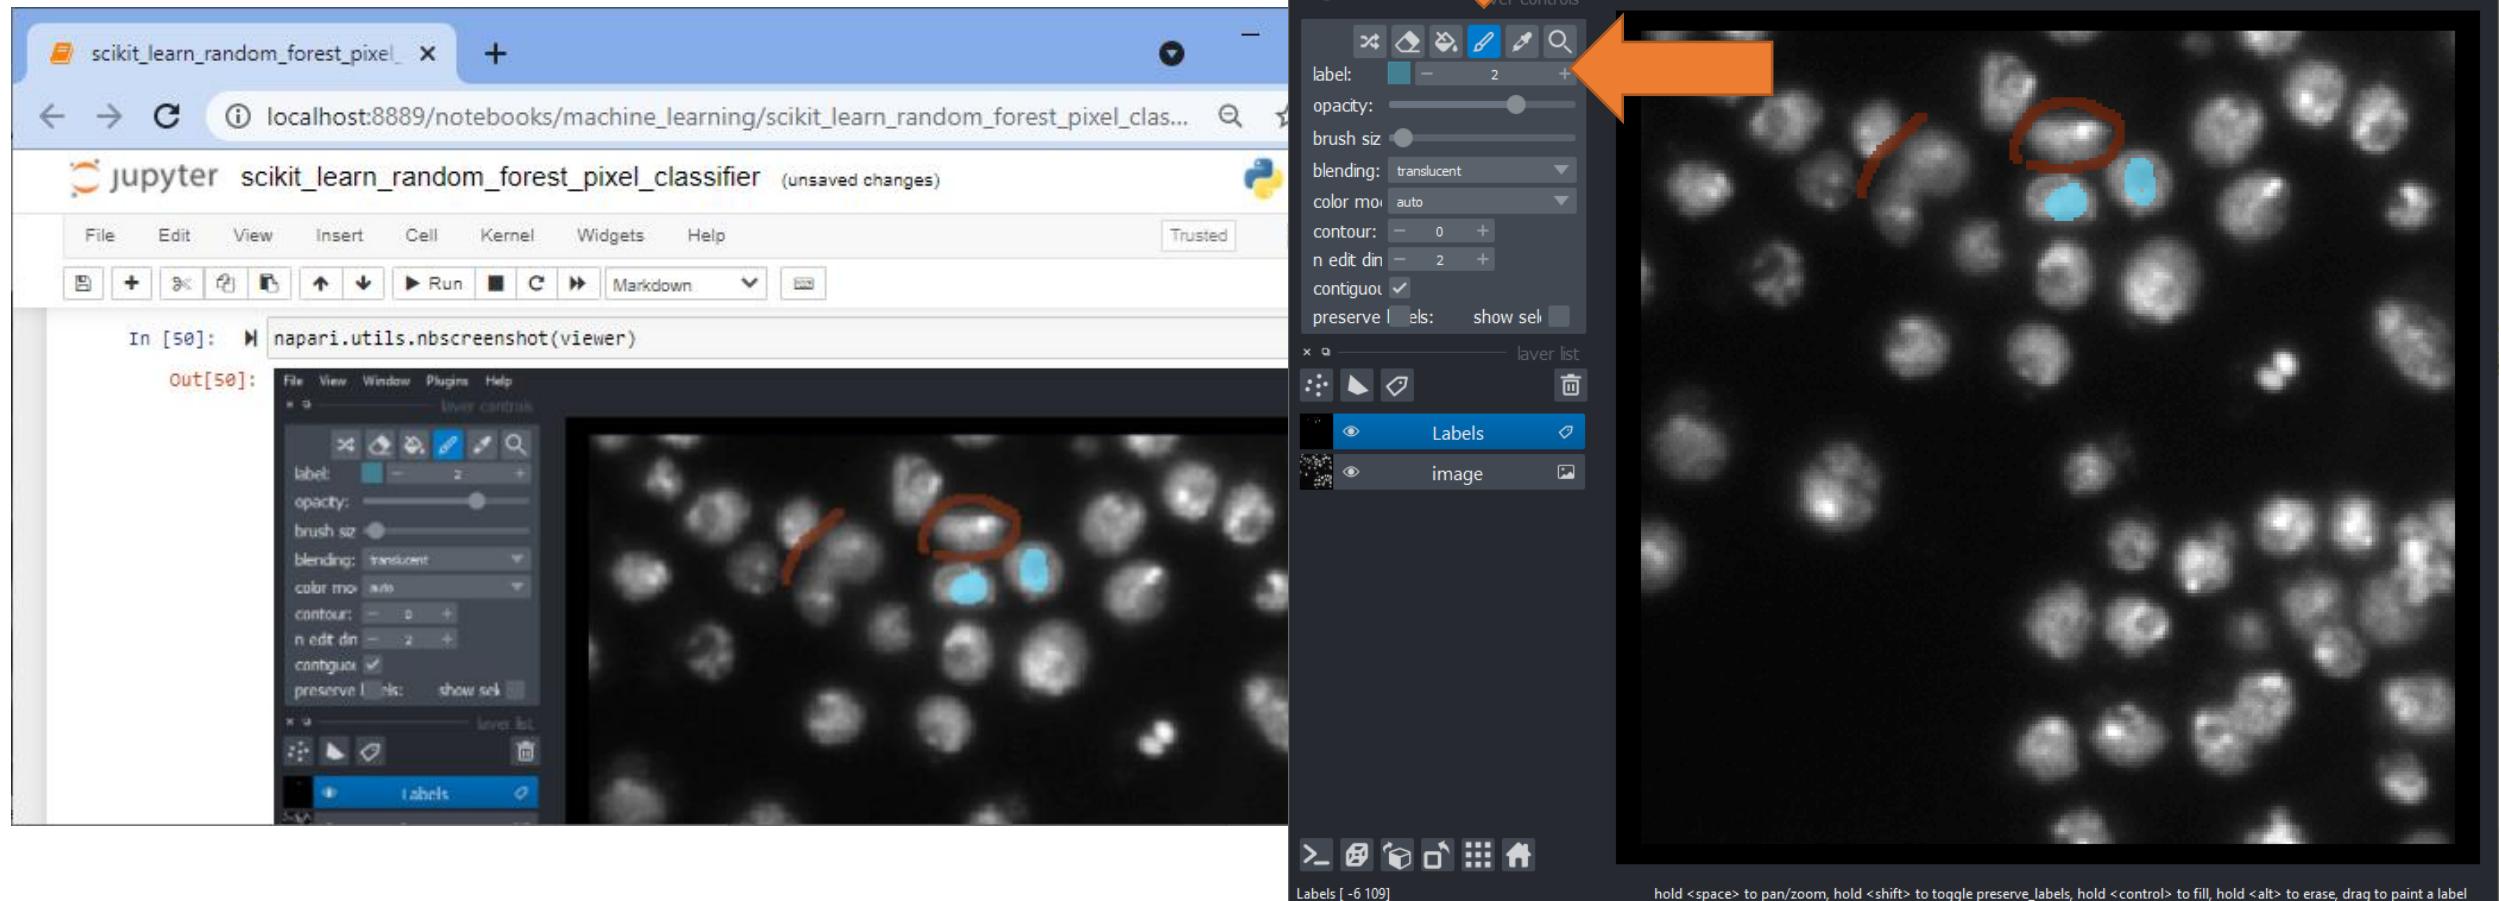
Go ahead after annotating at least two regions with labels 1 and 2.
Take a screenshot of the annotation:

napari Interface:

- Top bar: File, View, Window, Plugins, Help
- Left sidebar:
 - Layer controls: label (dropdown), opacity (slider), brush size (slider), blending (dropdown), color mode (dropdown), contour (slider), n edit dist (slider), contiguous (checkbox), preserve labels (checkbox).
 - Layer list: Labels (selected), image.
- Right panel: A grayscale image of cells with bright spots, representing the input image for segmentation.
- Bottom status bar: Labels [-6 -7], enter paint or fill mode to edit labels.

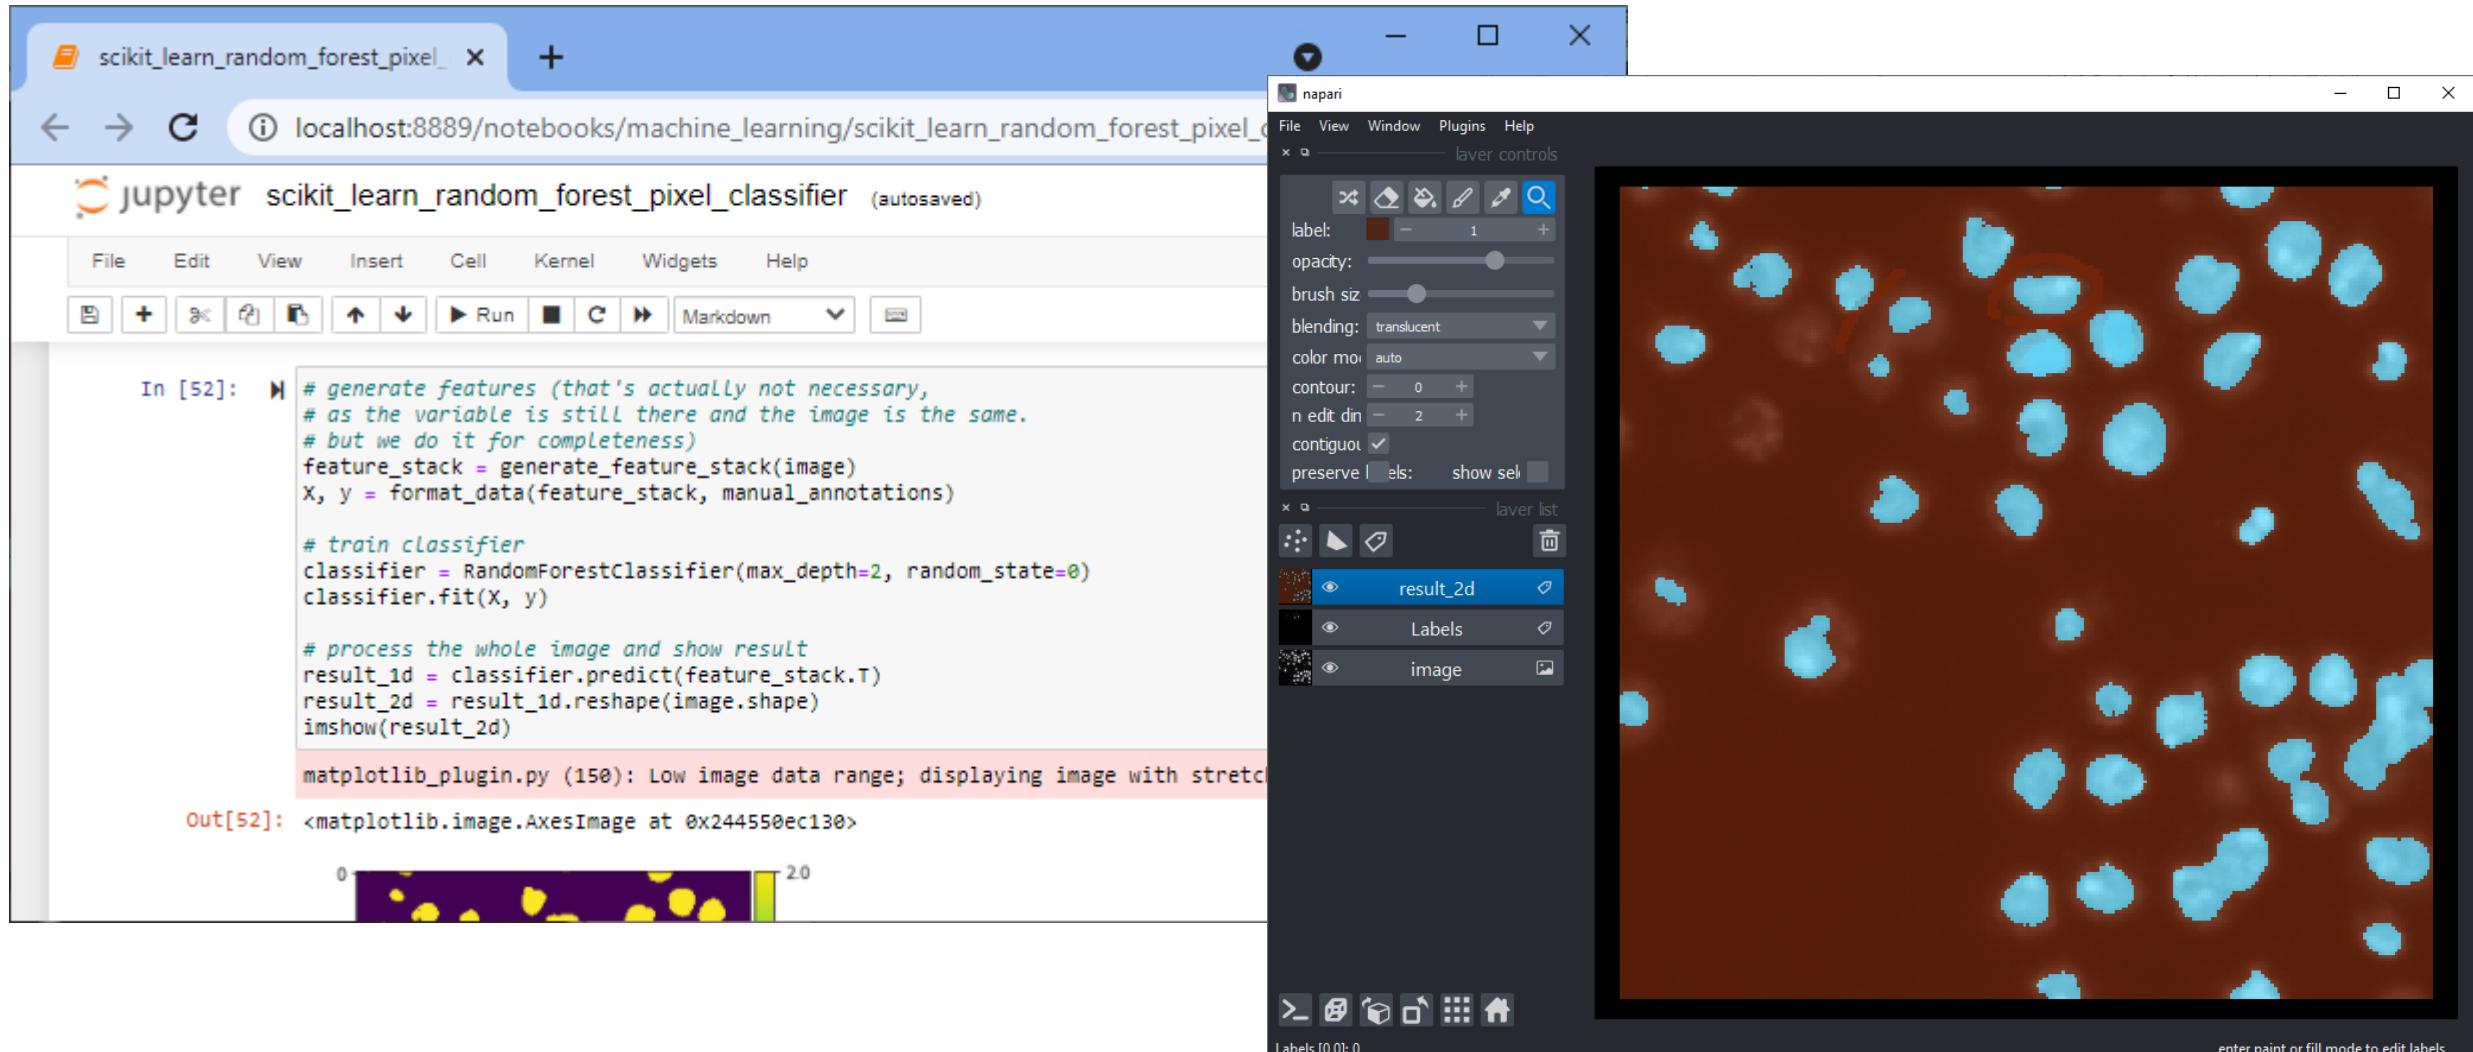
Interactive pixel classification

- Jupyter notebooks and napari side-by-side



Interactive pixel classification

- Jupyter notebooks and napari side-by-side



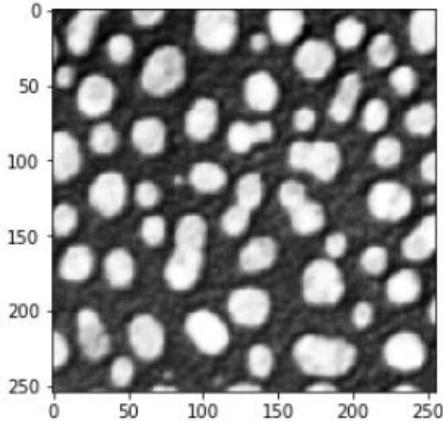
Accelerated pixel and object classification (APOC)

Robert Haase

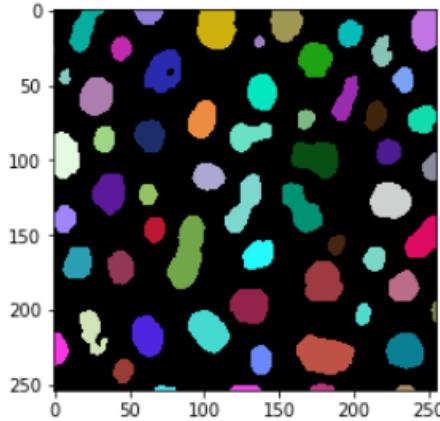
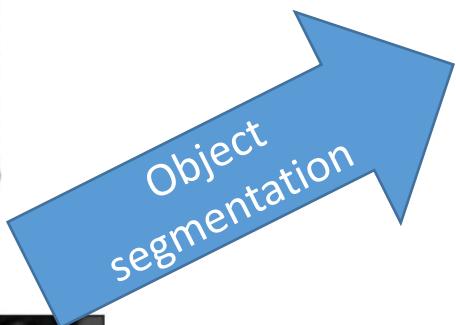
October 2022

Accelerated pixel and object classification

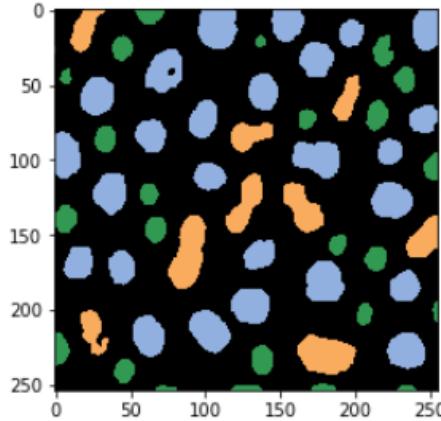
- APOC is a python library that makes use of OpenCL-compatible Graphics Cards to accelerate pixel and object classification



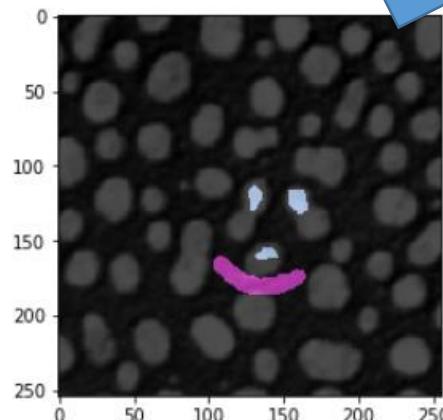
Raw image



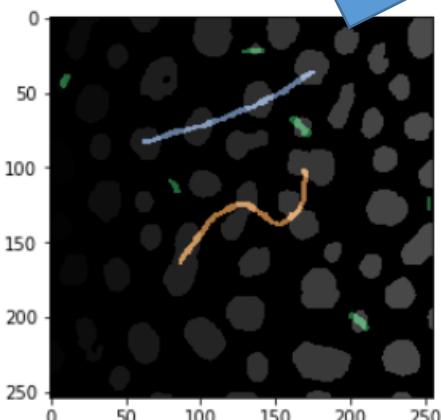
Object label image



Class label image



Pixel annotation



Object annotation

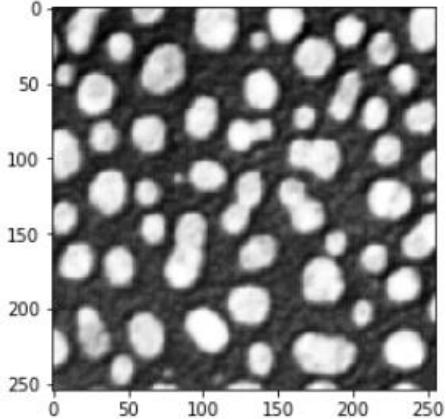
June 2023



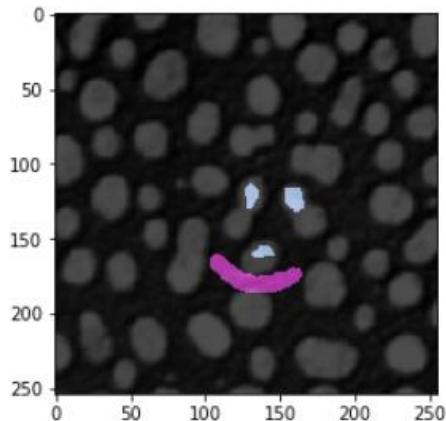
@haesleinhuepf

Object segmentation

- Pixel classification + connected component labeling



Raw image



Pixel annotation

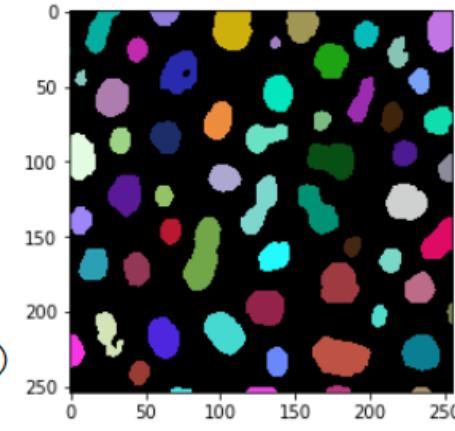
```
# define features
features = "gaussian_blur=1 gaussian_blur=5 sobel_of_gaussian_blur=1"

# this is where the model will be saved
cl_filename = 'my_object_segmenter.cl'

# delete classifier in case the file exists already
apoc.erase_classifier(cl_filename)

# train classifier
clf = apoc.ObjectSegmenter(opencl_filename=cl_filename, positive_class_identifier=2)
clf.train(features, manual_annotations, image)

segmentation_result = clf.predict(features=features, image=image)
cle.imshow(segmentation_result, labels=True)
```

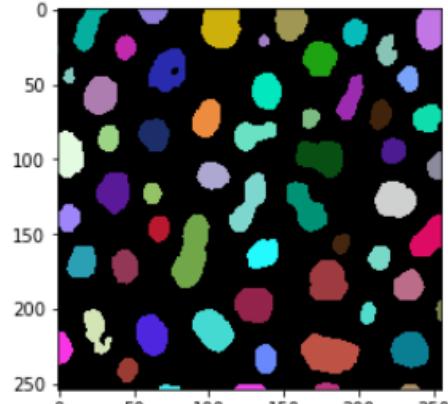


Object label image

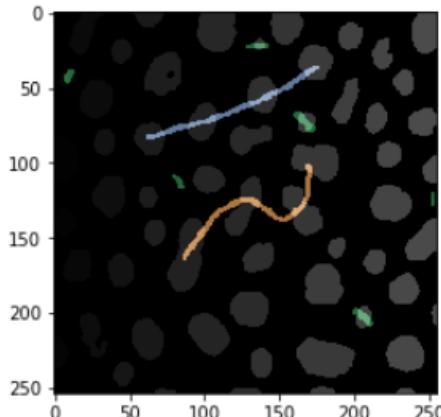
Object segmentation

Object classification

- Feature extraction + tabular classification



Object label image



Object annotation

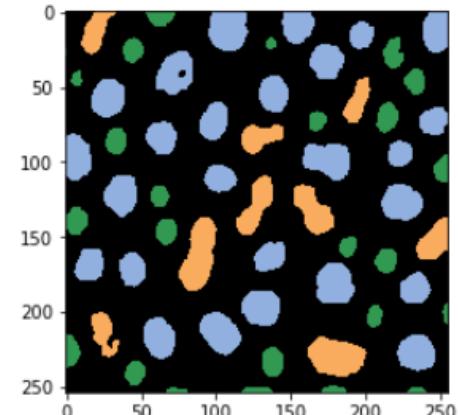
```
# for the classification we define size and shape as criteria
features = 'area mean_max_distance_to_centroid_ratio'

# This is where the model will be saved
cl_filename_object_classifier = "my_object_classifier.cl"

# delete classifier in case the file exists already
apoc.erase_classifier(cl_filename_object_classifier)

# train the classifier
classifier = apoc.ObjectClassifier(cl_filename_object_classifier)
classifier.train(features, segmentation_result, annotation, image)

# determine object classification
classification_result = classifier.predict(segmentation_result, image)
cle.imshow(classification_result, labels=True)
```

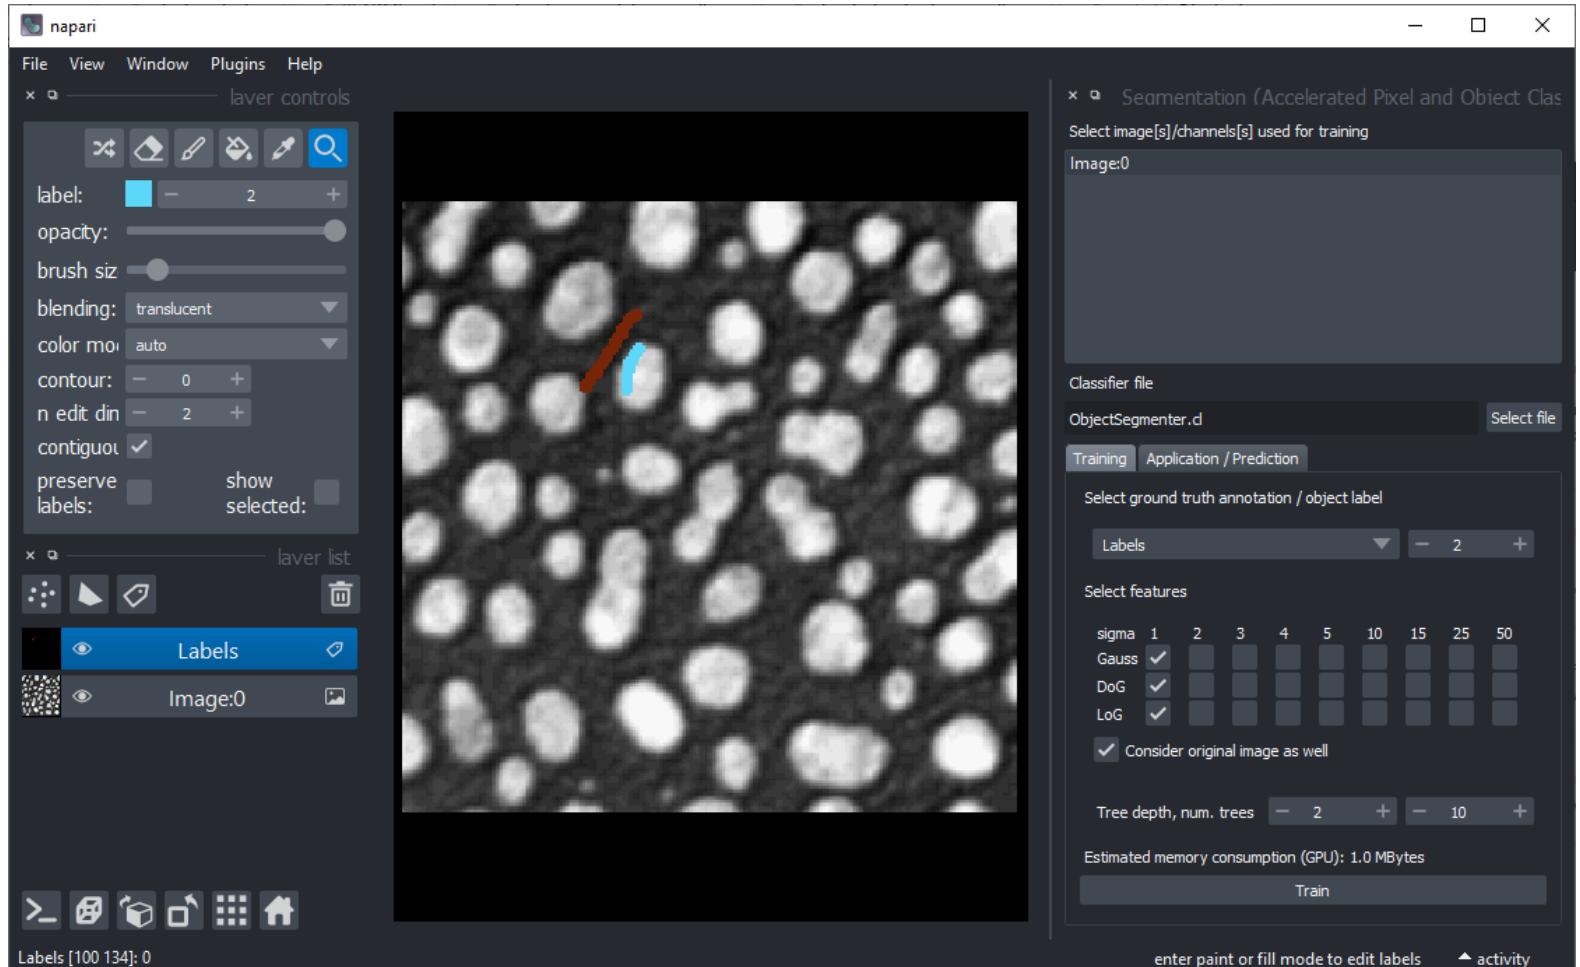


Class label image

Object classification

Graphical user interface

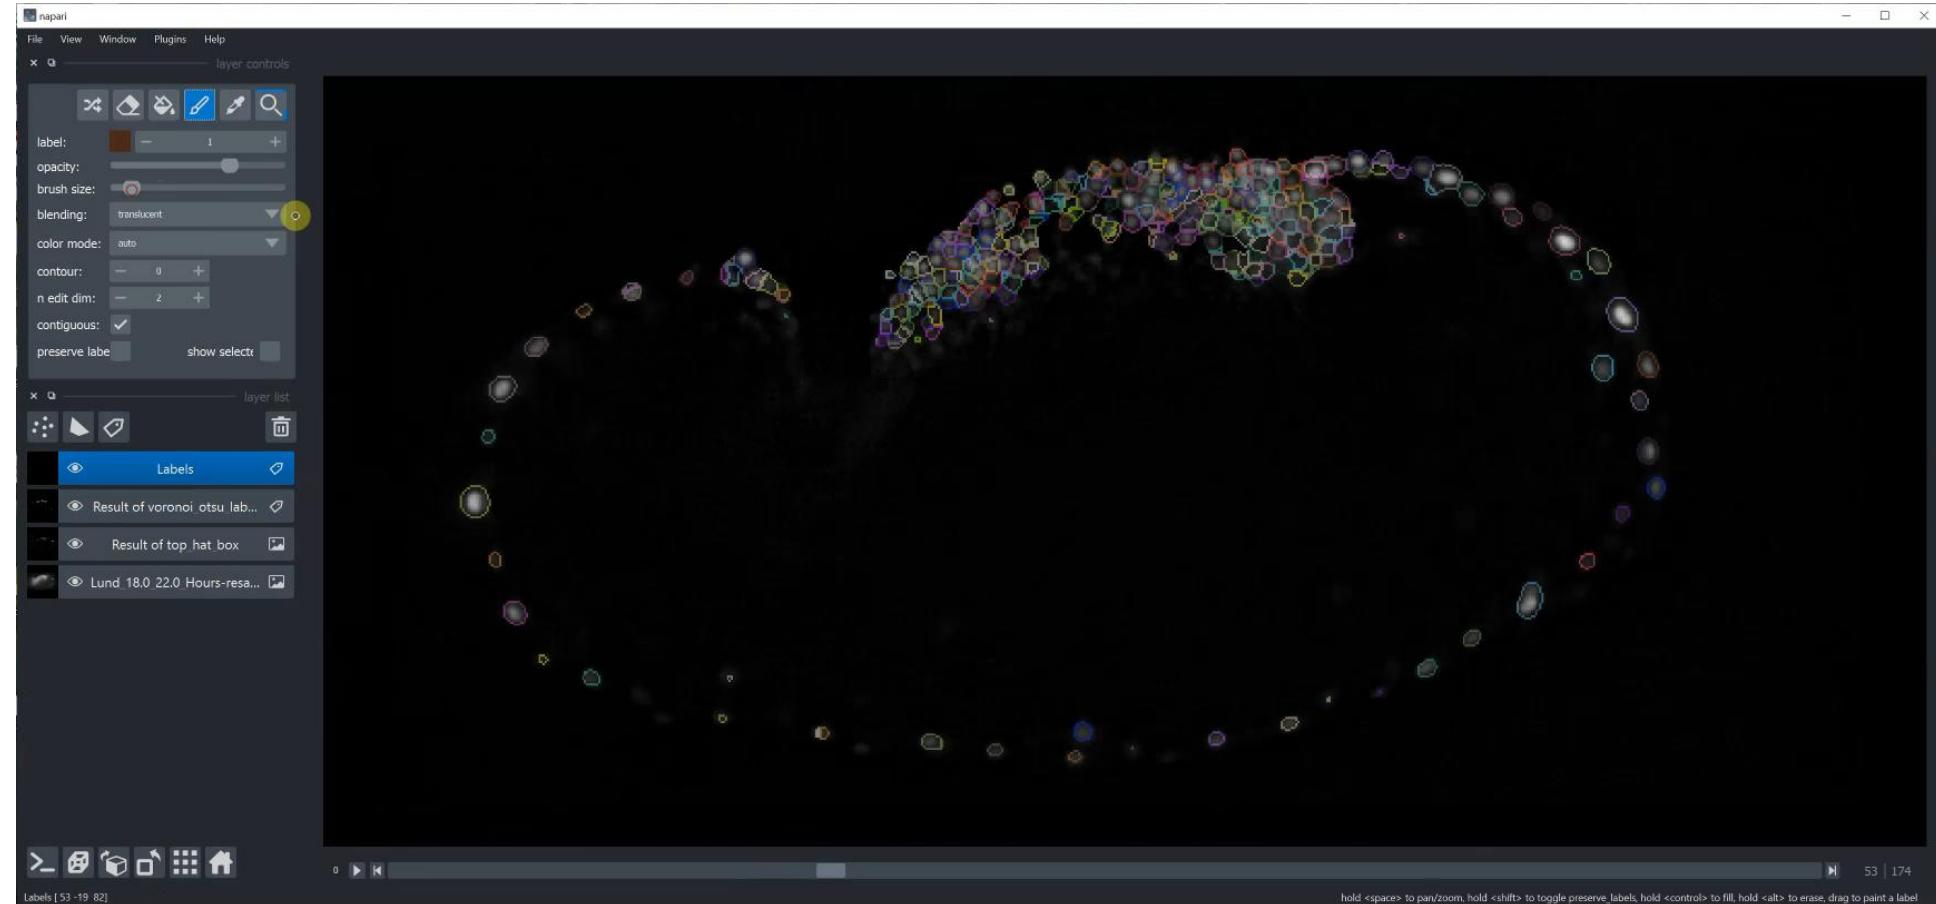
- Object segmentation
- <https://github.com/haesleinhuepf/napari-accelerated-pixel-and-object-classification#object-and-semantic-segmentation>



Supervised machine learning for tissue classification

Random Forest
Classifiers based on

- scikit-learn and
- clesperanto



@haesleinhuepf
@PoLDresden

<https://github.com/haesleinhuepf/napari-accelerated-pixel-and-object-classification>

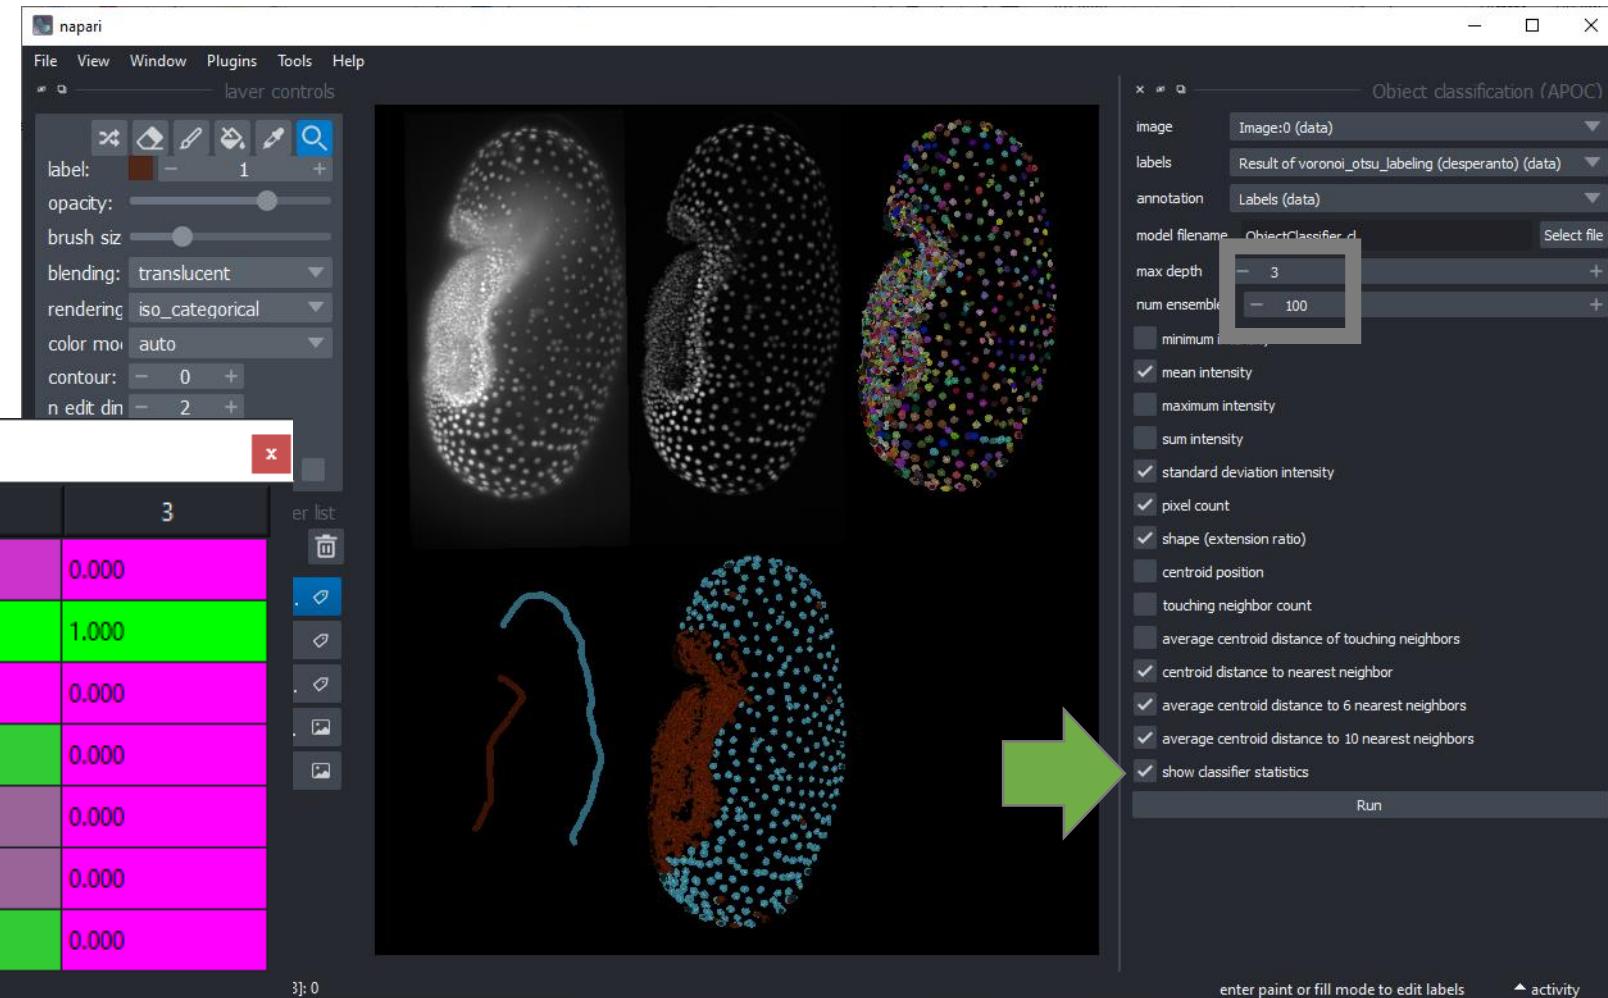
Image data source: Daniela Vorkel, Myers lab, MPI-CBG/CSBD

Data exploration / supervised machine learning

- Inspect how the random forest classifier makes decisions
- Note: Beware of correlated parameters!

Dock widget 1

	1	2	3
area	0.010	0.056	0.000
mean_intensity	0.200	0.278	1.000
standard_deviation_intensity	0.030	0.000	0.000
mean_max_distance_to_centroid_ratio	0.270	0.222	0.000
average_distance_of_n_nearest_neighbors=1	0.120	0.111	0.000
average_distance_of_n_nearest_neighbors=6	0.170	0.111	0.000
average_distance_of_n_nearest_neighbors=10	0.200	0.222	0.000



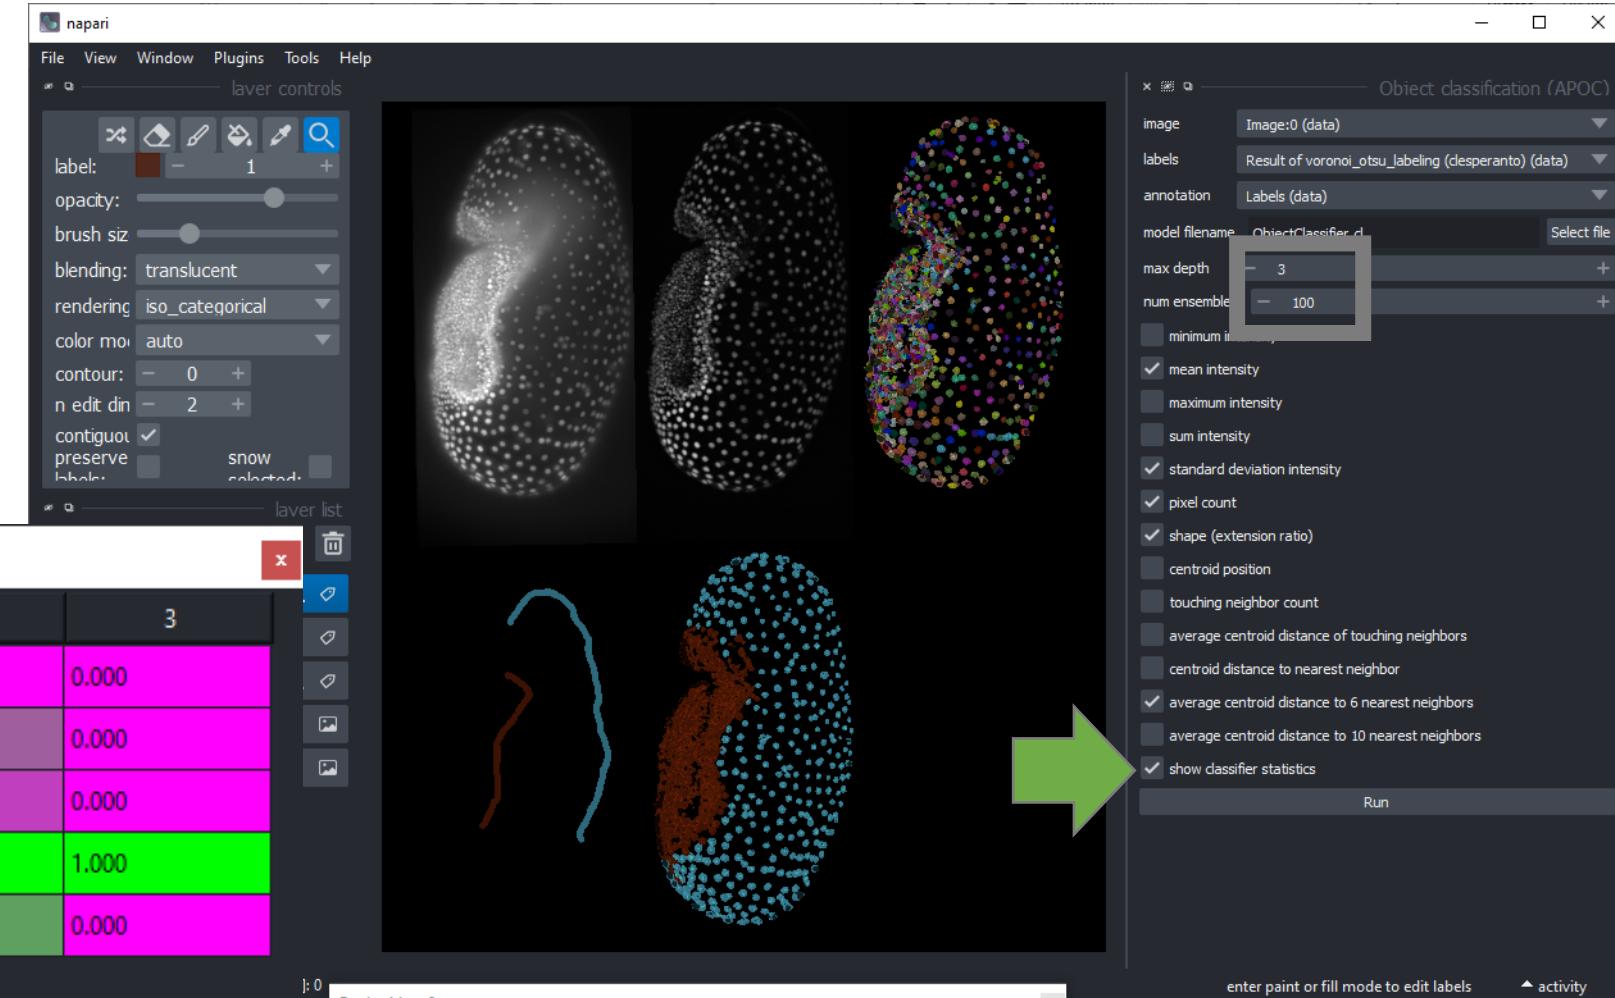
@haesleinhuepf
@PoLDresden

<https://github.com/haesleinhuepf/napari-accelerated-pixel-and-object-classification>

Image data source: Daniela Vorkel, Myers lab, MPI-CBG/CSBD

Data exploration / supervised machine learning

- Inspect how the random forest classifier makes decisions
- Note: Beware of correlated parameters!



@haesleinhuepf
@PoLDresden

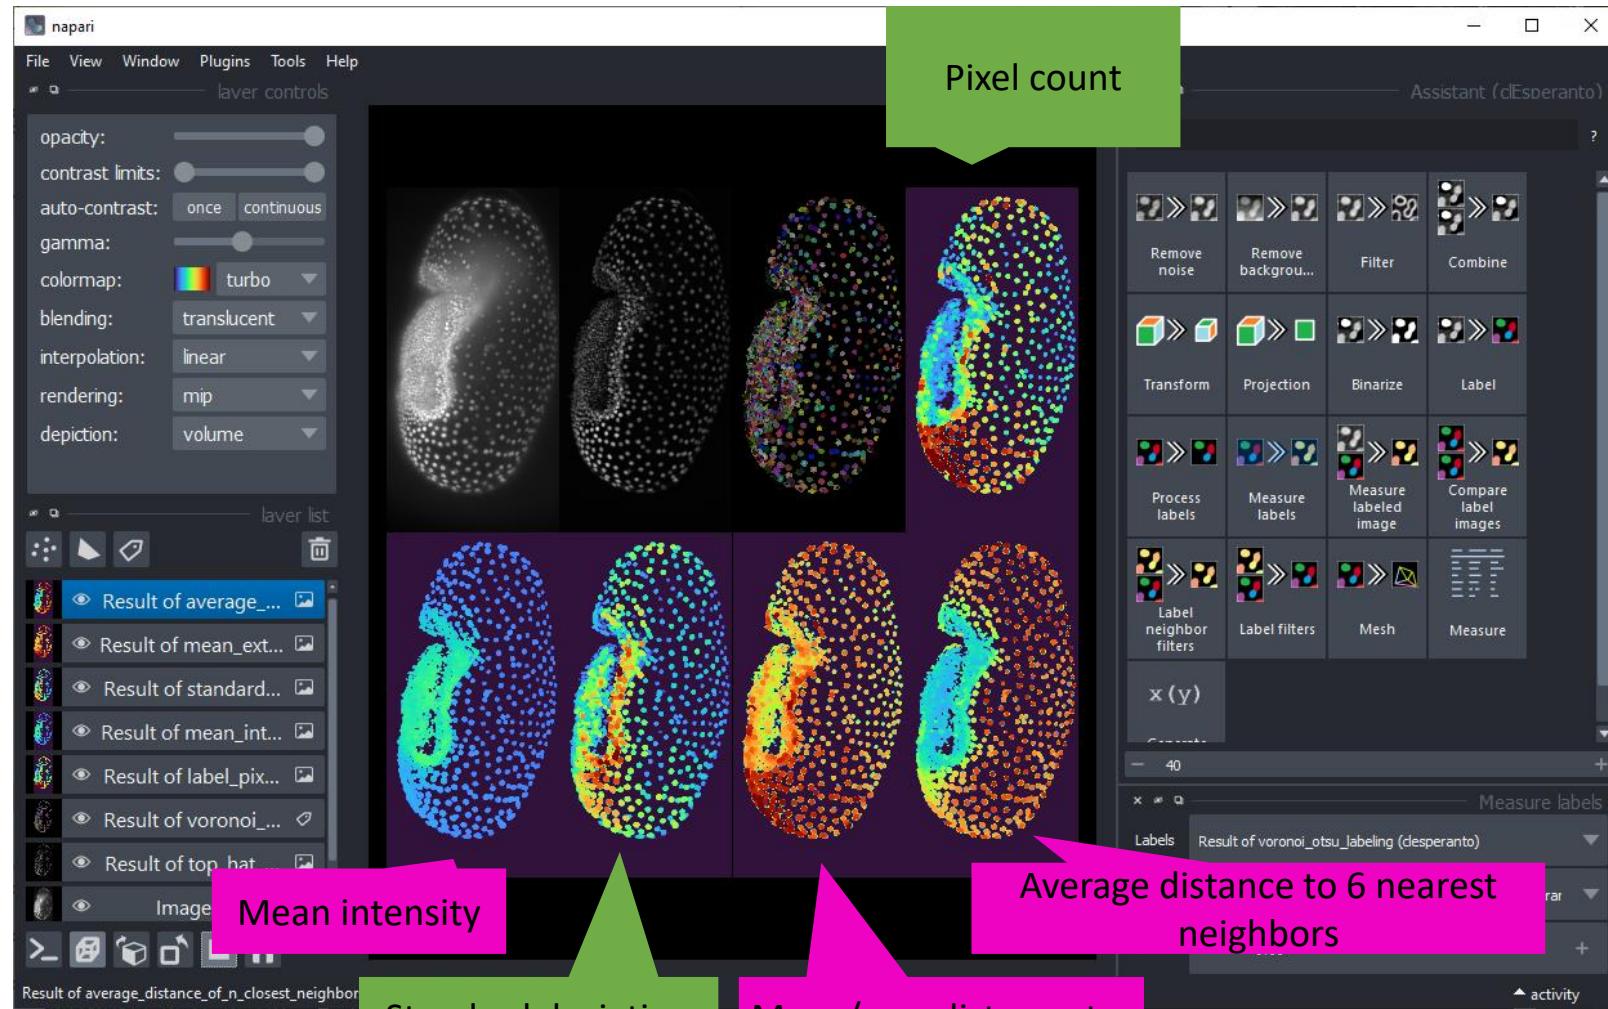
<https://github.com/haesleinhuepf/napari-accelerated-pixel-and-object-classification>

Image data source: Daniela Vorkel, Myers lab, MPI-CBG/CSBD

Data exploration / supervised machine learning

- Inspect how the random forest classifier makes decisions
- Note: Beware of correlated parameters!

	1	2	3
area	0.060	0.000	0.000
mean_intensity	0.330	0.167	0.000
standard_deviation_intensity	0.040	0.111	0.000
mean_max_distance_to_centroid_ratio	0.260	0.444	1.000
average_distance_of_n_nearest_neighbors=6	0.310	0.278	0.000



https://github.com/clEsperanto/napari_pyclesperanto_assistant

Image data source: Daniela Vorkel, Myers lab, MPI-CBG/CSBD

Exercises

Pixel classification / object segmentation

- Use Napari to segment objects

Interactive pixel classification and object segmentation in Napari

In this exercise we will train a [Random Forest Classifier](#) for pixel classification and convert the result in an instance segmentation. We will use the napari plugin [napari-accelerated-pixel-and-object-classification](#).

Getting started

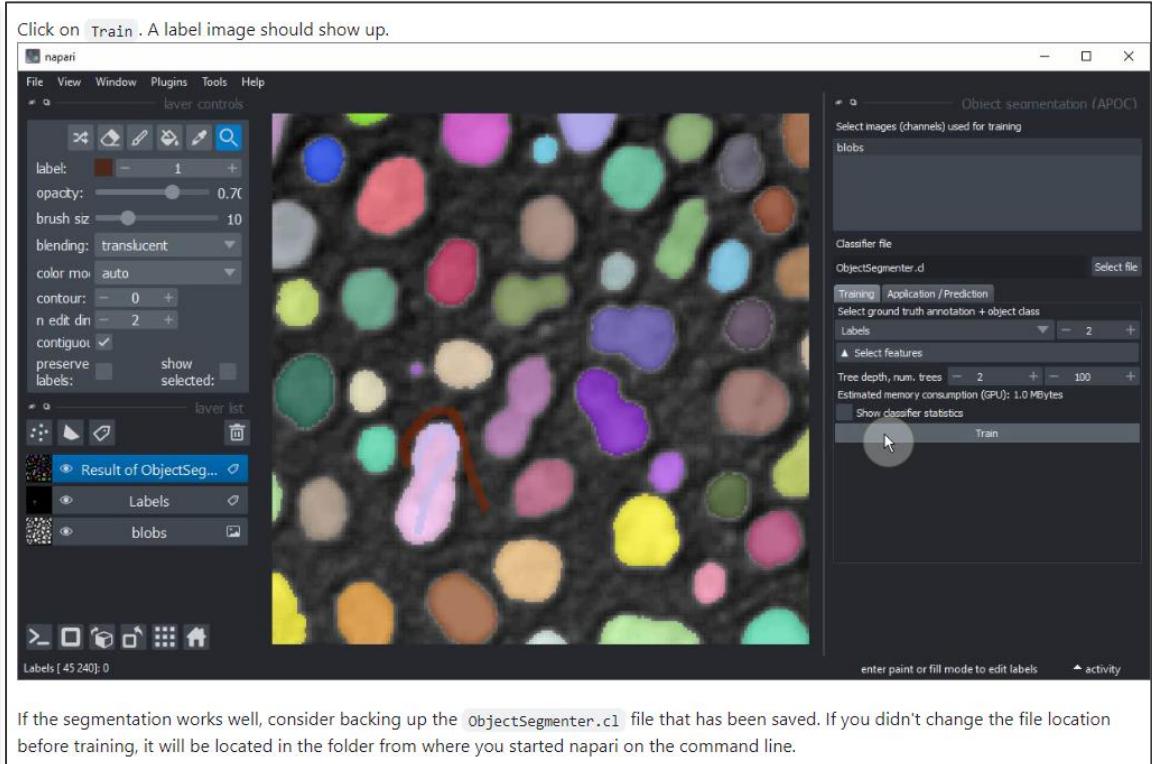
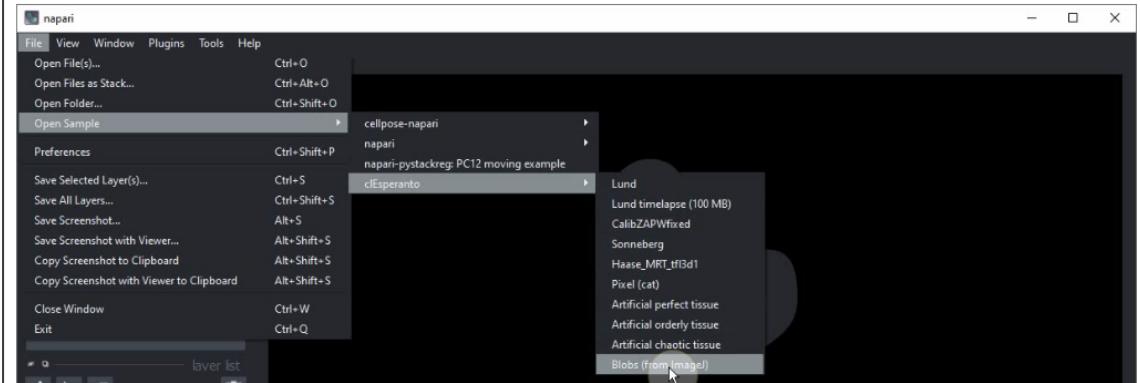
Open a terminal window and activate your conda environment:

```
conda activate devbio-napari-env
```

Afterwards, start up Napari:

```
napari
```

Load the "Blobs" example dataset from the menu `File > Open Sample > c1Esperanto > Blobs (from Image)`



https://github.com/BiAPoL/Bio-image_Analysis_with_Python/blob/main/09_machine_learning/interactive_pixel_classification/readme.md

Object classification

- Use Napari to group round and elongated objects

Interactive object classification in Napari

In this exercise we will train a [Random Forest Classifiers](#) for classifying segmented objects. We will use the napari plugin [napari-accelerated-pixel-and-object-classification](#).

Getting started

Open a terminal window and activate your conda environment:

```
conda activate devbio-napari-env
```

Afterwards, start up Napari:

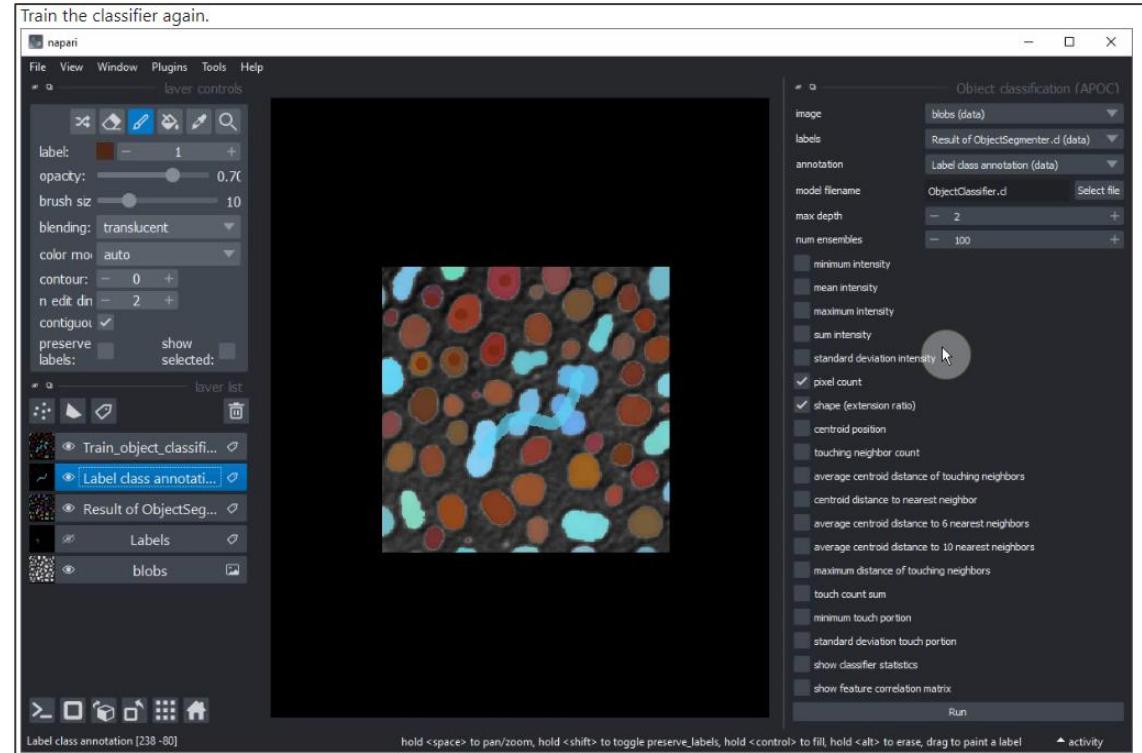
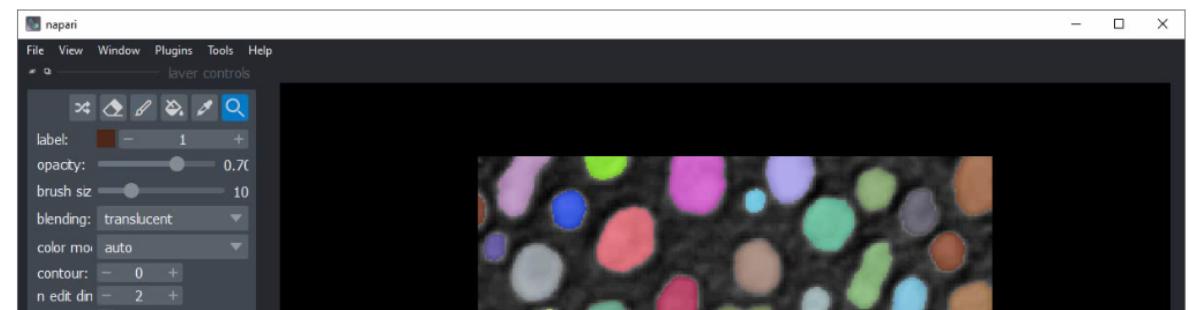
```
napari
```

Load the "Blobs" example dataset from the menu [File > Open Sample > c1Esperanto > Blobs \(from Image\)](#)

We furthermore need a label image. You can create it using the pixel classifier trained earlier or using the menu [Tools > Segmentation / labeling > Gauss-Otsu Labeling \(clesperanto\)](#).

Object classification

Our starting point is a loaded image and a label image with segmented objects. The following procedure is also shown in [this video](#).



If you are happy with the trained classifier, copy the file to a safe place. When training the next classifier this one might be overwritten.

Extra exercise

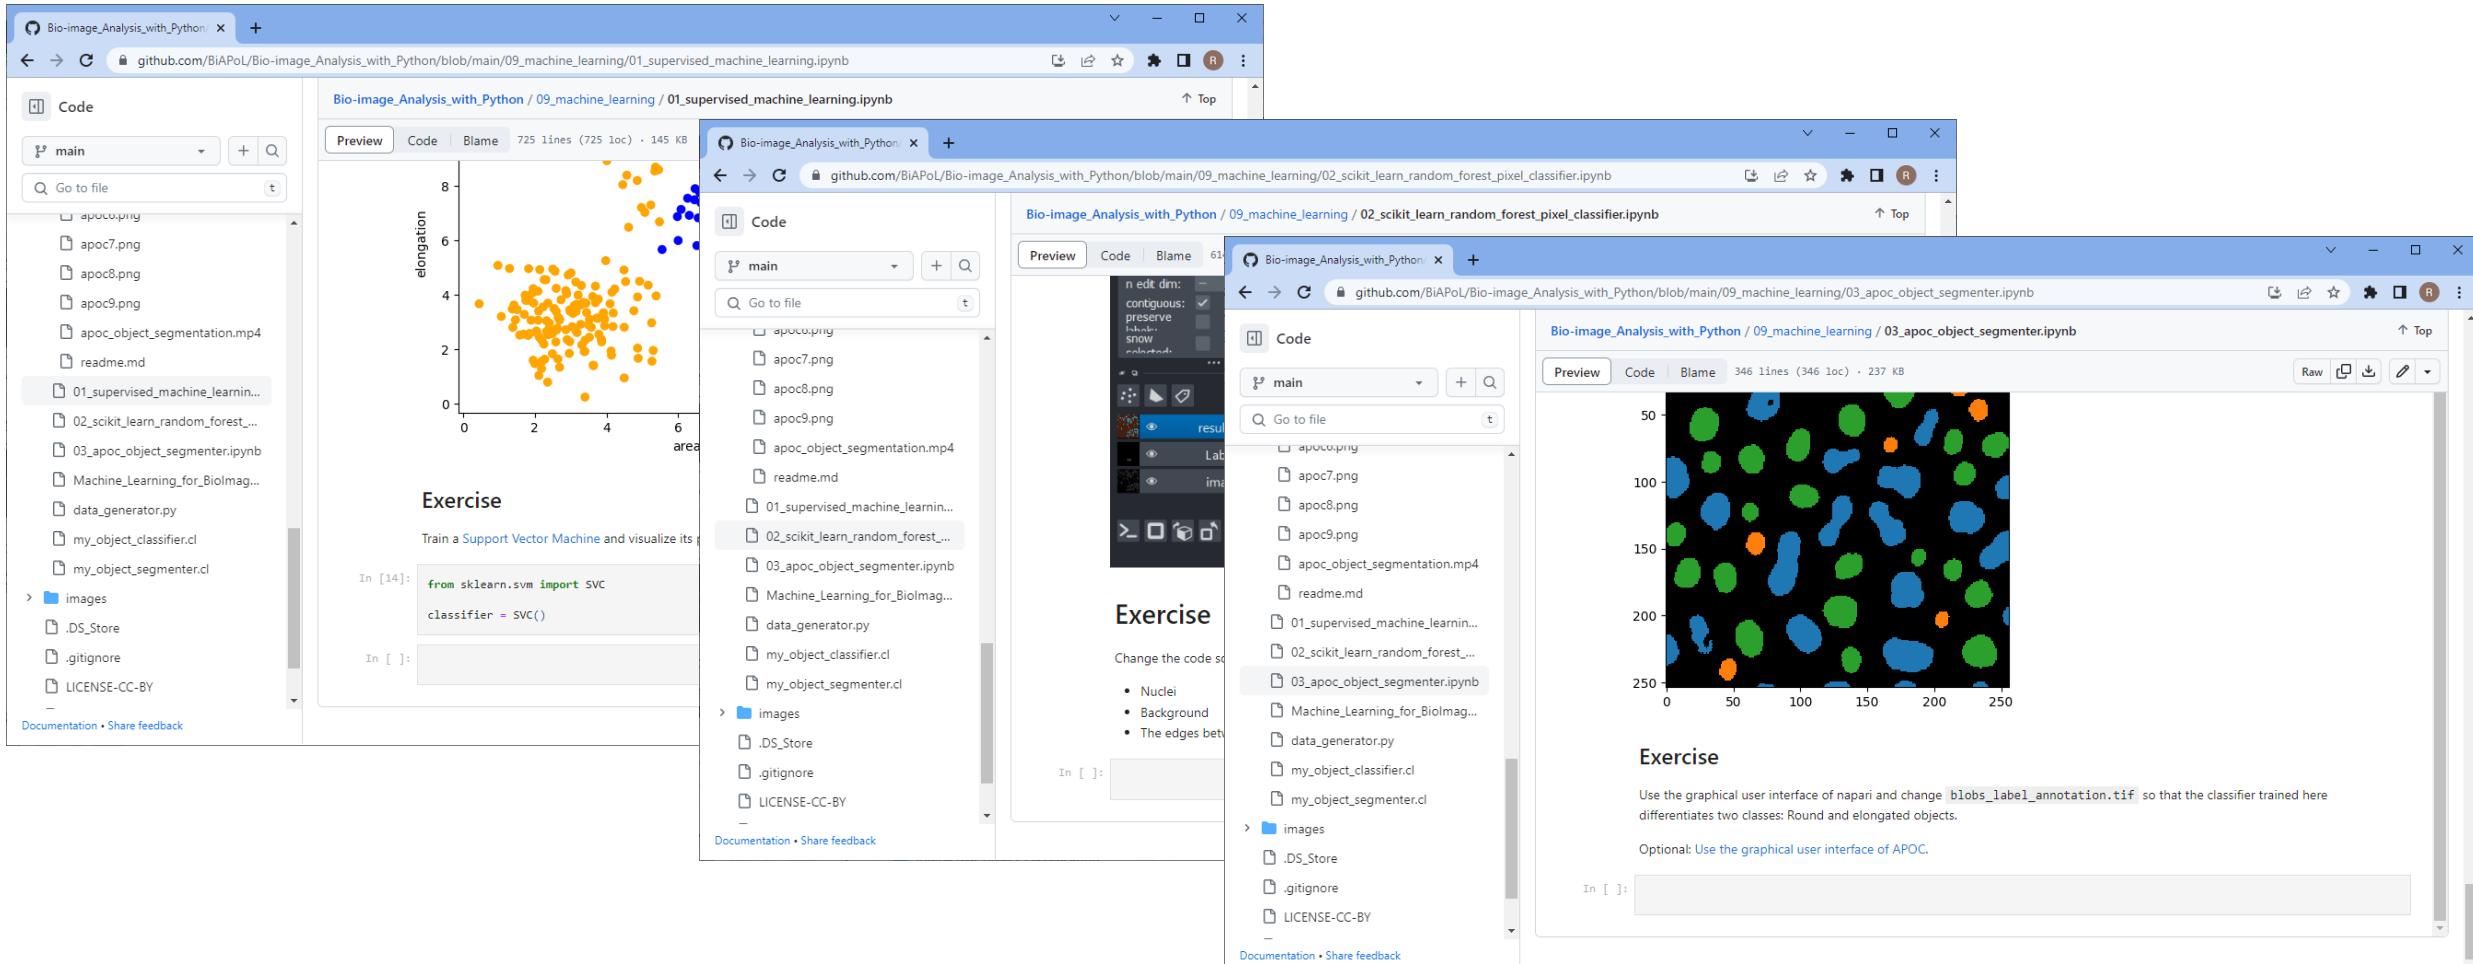
Retrain the classifier so that it can differentiate three different classes:

- Small round objects
- Large round objects
- Large elongated objects

https://github.com/BiAPoL/Bio-image_Analysis_with_Python/blob/main/09_machine_learning/interactive_object_classification/readme.md

Machine learning using Python

- Use scikit-learn and apoc in Jupyter Notebooks to train and apply Random Forest Classifiers and Support Vector Machines



The image shows three Jupyter Notebook windows side-by-side, each displaying a different aspect of machine learning analysis:

- Left Notebook:** Titled "Bio-image_Analysis_with_Python / 09_machine_learning / 01_supervised_machine_learning.ipynb". It contains a scatter plot of elongation vs area for two classes of objects (orange and blue), and a code cell for training a Support Vector Machine (SVC) classifier.
- Middle Notebook:** Titled "Bio-image_Analysis_with_Python / 09_machine_learning / 02_scikit_learn_random_forest_pixel_classifier.ipynb". It shows a visualization of a random forest classifier's decision boundaries on a 2D feature space.
- Right Notebook:** Titled "Bio-image_Analysis_with_Python / 09_machine_learning / 03_apoc_object_segmenter.ipynb". It displays a binary segmentation mask where objects are labeled as Nuclei (green), Background (blue), or Edges (orange).

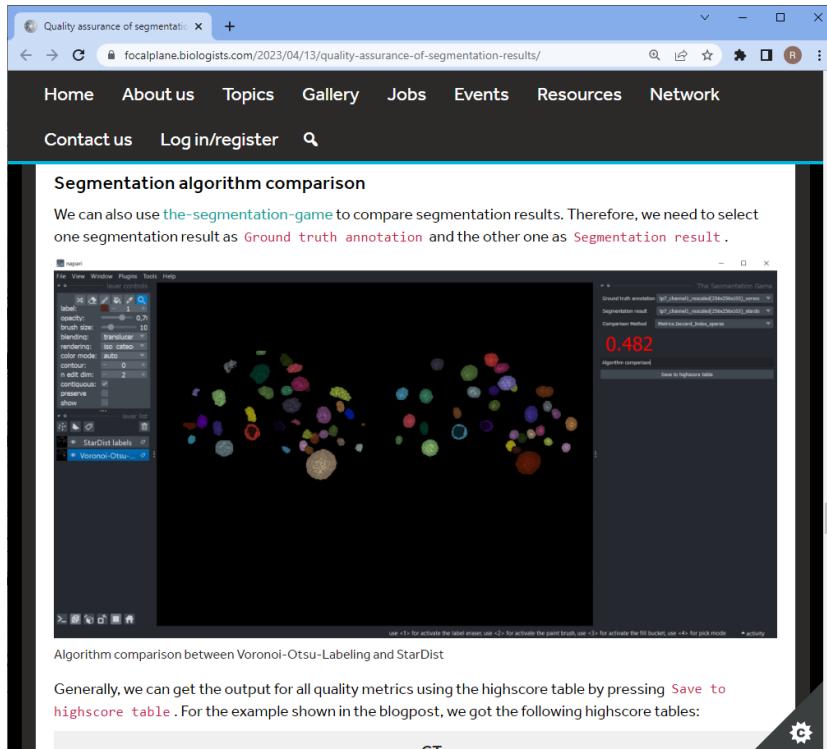
https://github.com/BiAPoL/Bio-image_Analysis_with_Python/blob/main/09_machine_learning/

Summary & outlook

Further reading

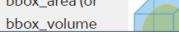
- Blog posts on the Focalplane about...

Segmentation Quality Assurance



<https://focalplane.biologists.com/2023/04/13/quality-assurance-of-segmentation-results/>

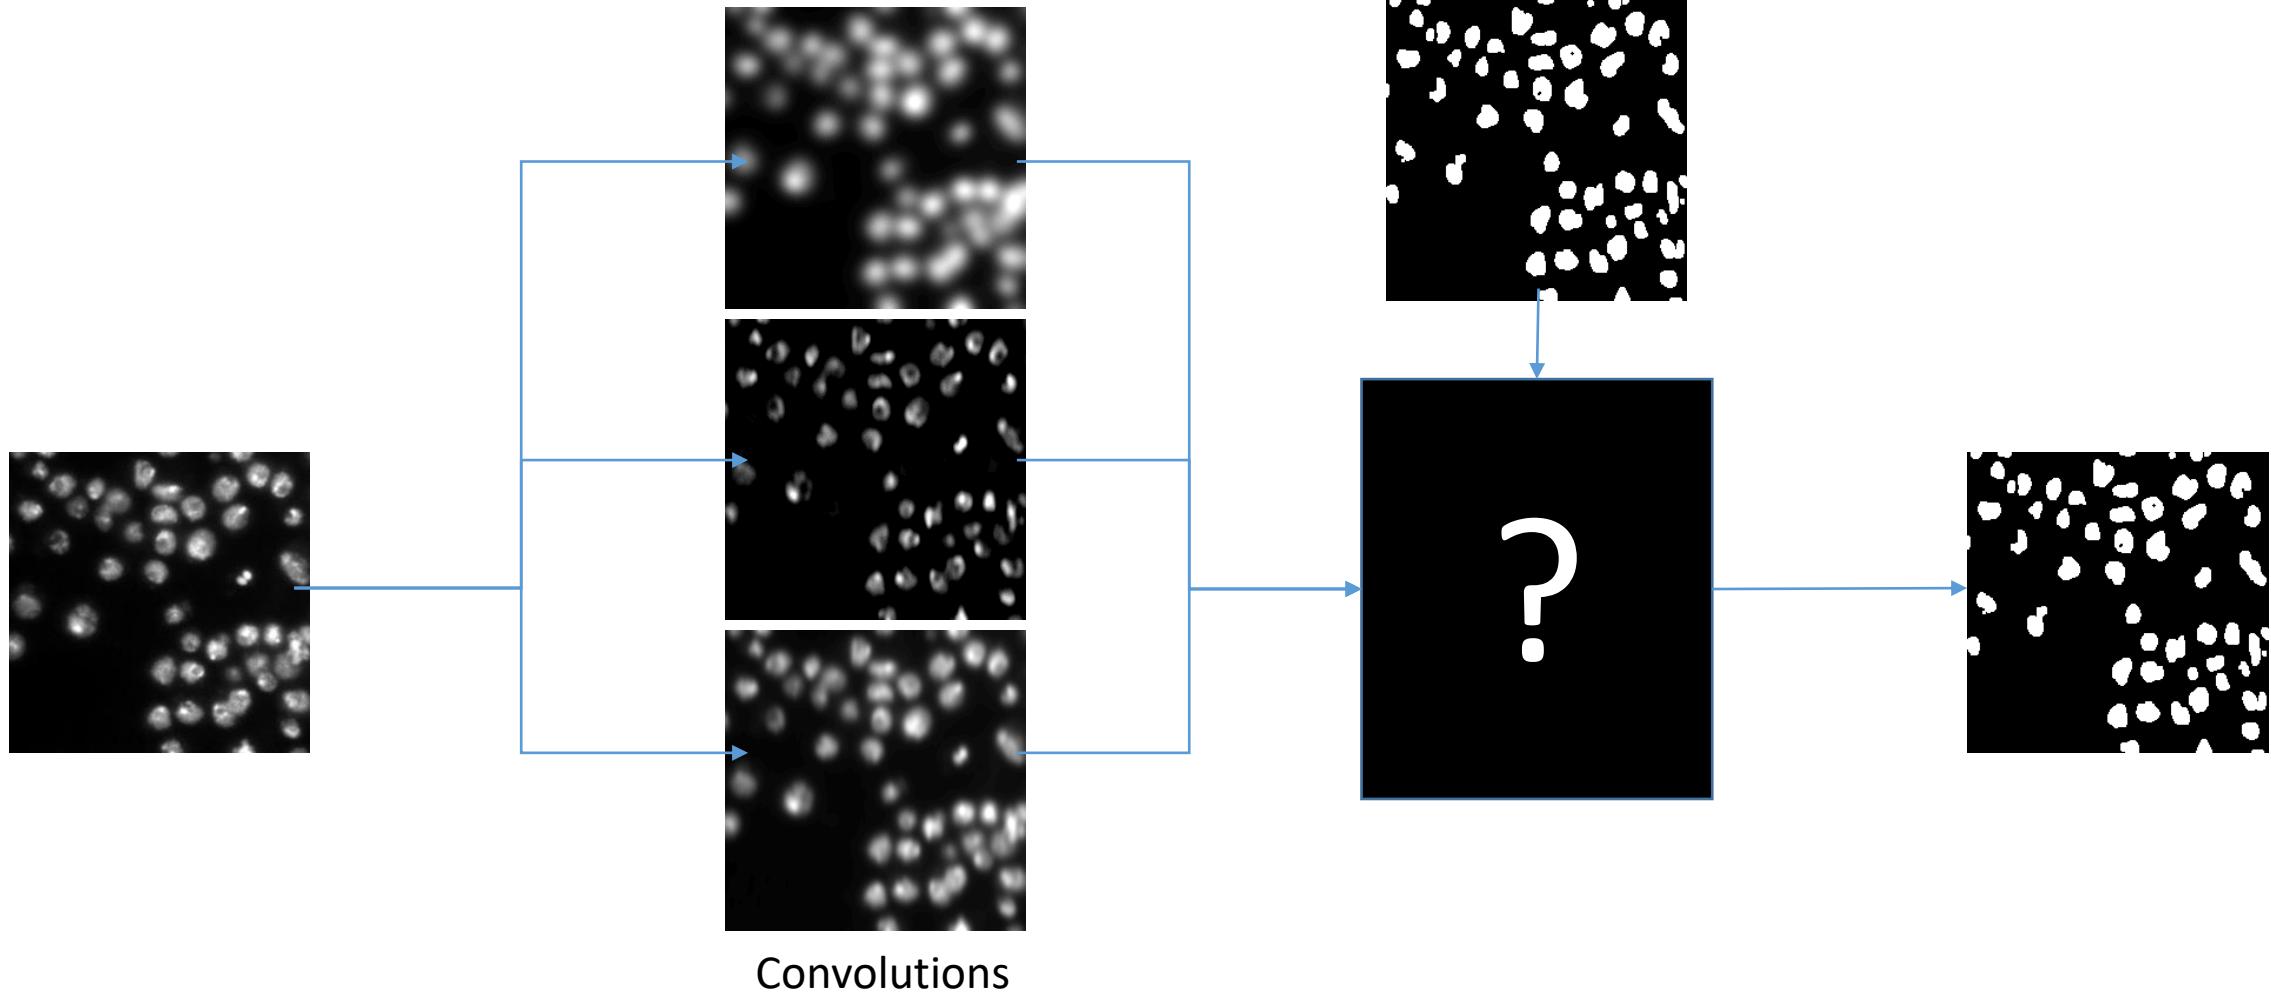
Feature Extraction

feature	schematics	category	explanation	formula	implementation
area (or volume in 3D)		siz	= sum of pixels/ voxels of the region scaled by pixel-area/ voxel-volume.	$\text{Area} = \sum \text{pixels} \times \text{pixel_area}$ $\text{Volume} = \sum \text{voxels} \times \text{voxel_area}$	nsr
aspect_ratio		sh	= the ratio between object width and object length.	$\text{Aspect ratio} = \frac{\text{major axis length}}{\text{minor axis length}}$	nsr
bbox		po	= minimum range in each spatial dimension	$\text{width} = (x_2 - x_1)$ $\text{height} = (y_2 - y_1)$ $\text{depth} = (z_2 - z_1)$	nsr
bbox_area (or bbox_volume)		siz	= number of pixels/ voxels of bounding box scaled by pixel-area/ voxel-volume	$\text{bbox_area} = \text{width} \times \text{height} \times \text{depth}$	nsr

<https://focalplane.biologists.com/2023/05/03/feature-extraction-in-napari/>

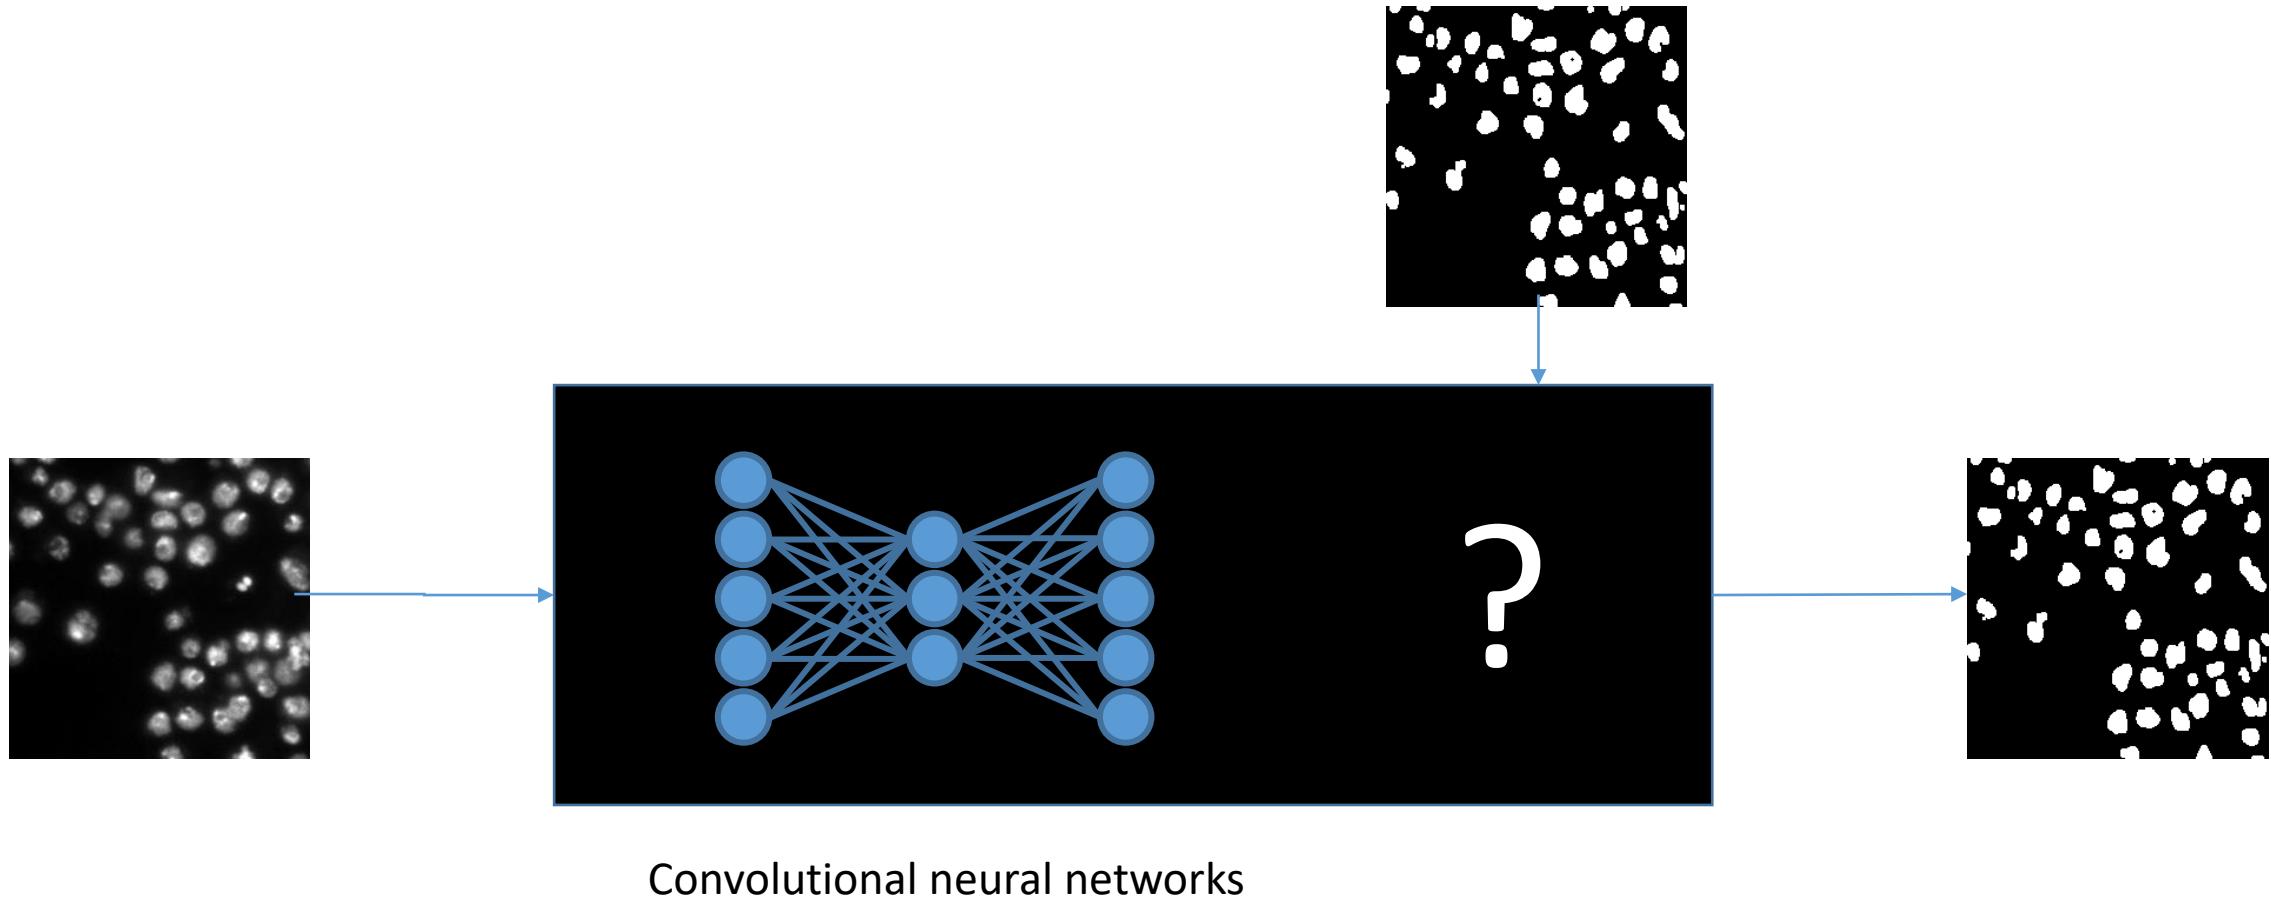
Machine learning for image analysis

- In classical machine learning, we typically select features for training our classifier



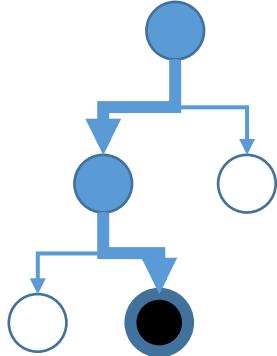
Outlook: Deep learning for image analysis

- In deep learning, this selection becomes part of the black box



Today, you learned

- Machine learning for Pixel and Object segmentation
- Python
 - Scikit-learn / napari
 - Accelerated pixel and object classifiers (APOC)



Coming up next:

- Unsupervised machine learning
- Deep learning

