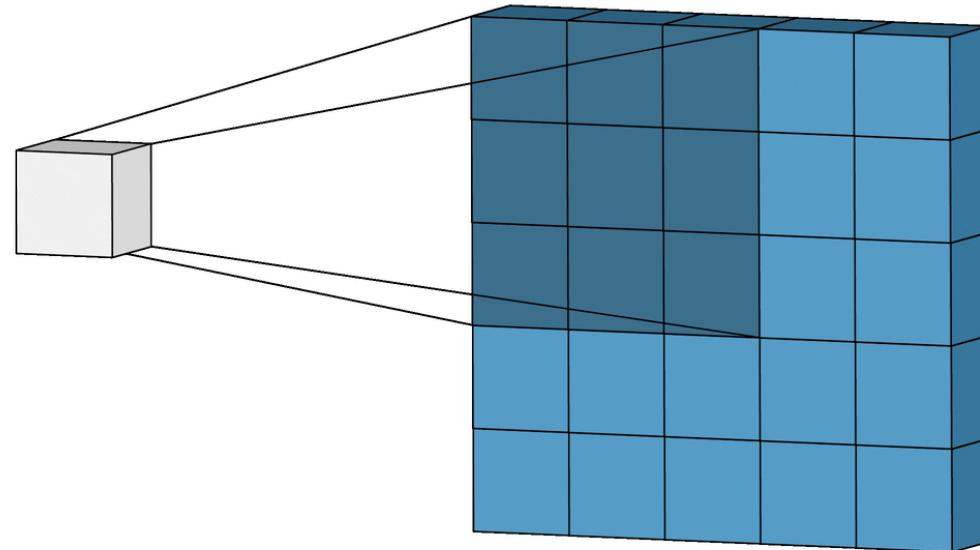


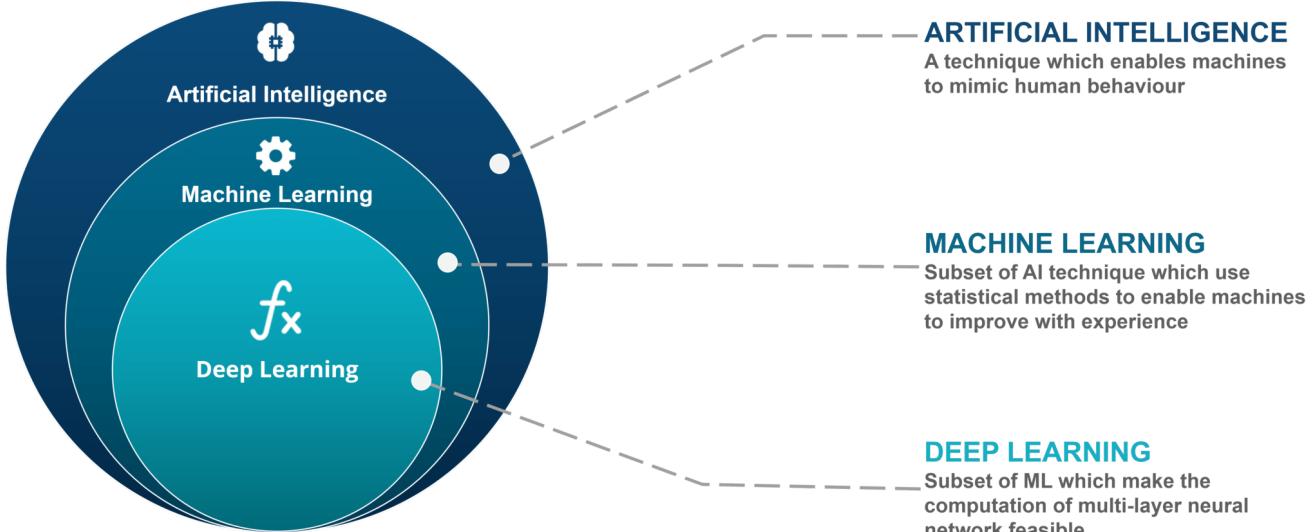
# An Introduction to Deep Learning

*Manan Lalit*

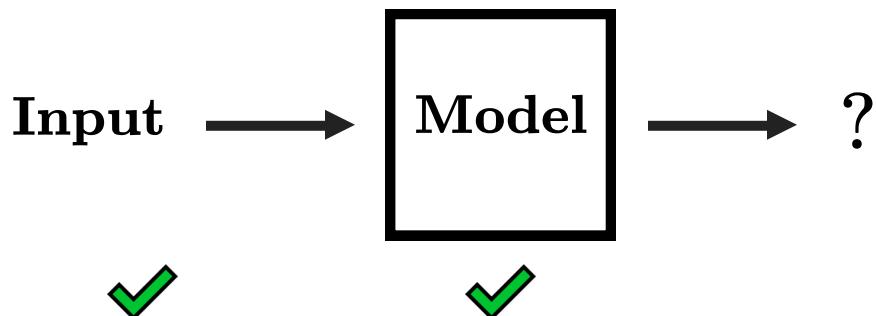
July 5, 2022



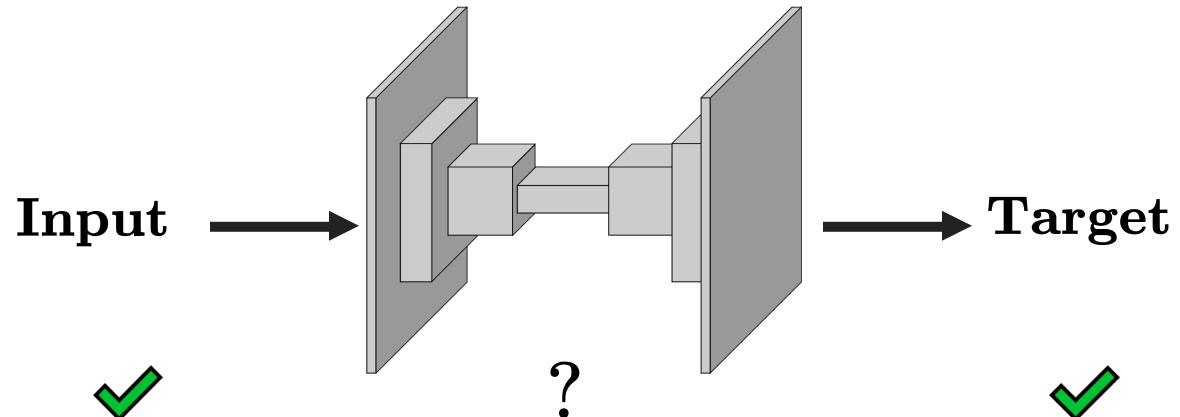
# Deep Learning Vs AI, Learning Vs Non-Learning



## Non-Learning Approaches



## (Deep) Learning-based Approaches



Training Phase: *Context*-aware learning of model weights or parameters

# History of Deep Learning

## Learning representations by back-propagating errors

David E. Rumelhart\*, Geoffrey E. Hinton†  
& Ronald J. Williams\*

\* Institute for Cognitive Science, C-015, University of California  
San Diego, La Jolla, California 92093, USA

† Department of Computer Science, Carnegie-Mellon University  
Pittsburgh, Philadelphia 15213, USA

(1986) Rumelhart ... Hinton showing  
back-propagation can be used to train networks

## ImageNet Classification with Deep Convolutional Neural Networks

Alex Krizhevsky      Ilya Sutskever      Geoffrey E. Hinton  
University of Toronto    University of Toronto    University of Toronto  
kriz@cs.utoronto.ca    ilya@cs.utoronto.ca    hinton@cs.utoronto.ca

(2012) AlexNet won the ImageNet challenge  
First network to show that deep networks  
can be trained using GPUs

## Multilayer Feedforward Networks are Universal Approximators

KURT HORNIK

Technische Universität Wien

MAXWELL STINCHCOMBE AND HALBERT WHITE

University of California, San Diego

(Received 16 September 1988; revised and accepted 9 March 1989)

(1989) Hornik et al showed that multi-layer  
networks are capable of approximating any function

## U-Net: Convolutional Networks for Biomedical Image Segmentation

Olaf Ronneberger, Philipp Fischer, and Thomas Brox

Computer Science Department and BIOSS Centre for Biological Signalling Studies,  
University of Freiburg, Germany  
ronneber@informatik.uni-freiburg.de,  
WWW home page: <http://lmb.informatik.uni-freiburg.de/>

(2015) Ronneberger et al show skip connections  
architecture enables better training with fewer examples



(1993) Yann Le Cunn showing performance of  
convolutional nets on hand-written digits

## Attention Is All You Need

Ashish Vaswani\*      Noam Shazeer\*      Niki Parmar\*      Jakob Uszkoreit\*  
Google Brain            noam@google.com       nikip@google.com       usz@google.com

Llion Jones\*      Aidan N. Gomez\* †      Lukasz Kaiser\*  
Google Research       University of Toronto       lukaszkaiser@google.com  
llion@google.com       aidan@cs.toronto.edu

Illia Polosukhin\* †  
illia.polosukhin@gmail.com

(2017) Vaswani et al propose a new  
non-convolutional architecture

# Vocabulary of Deep Learning

Optimizer

Scheduler

Loss fn.

Epoch

U-Net

Skip  
Connections

Learning  
Rate

Batch  
Norm

Gradient  
Descent

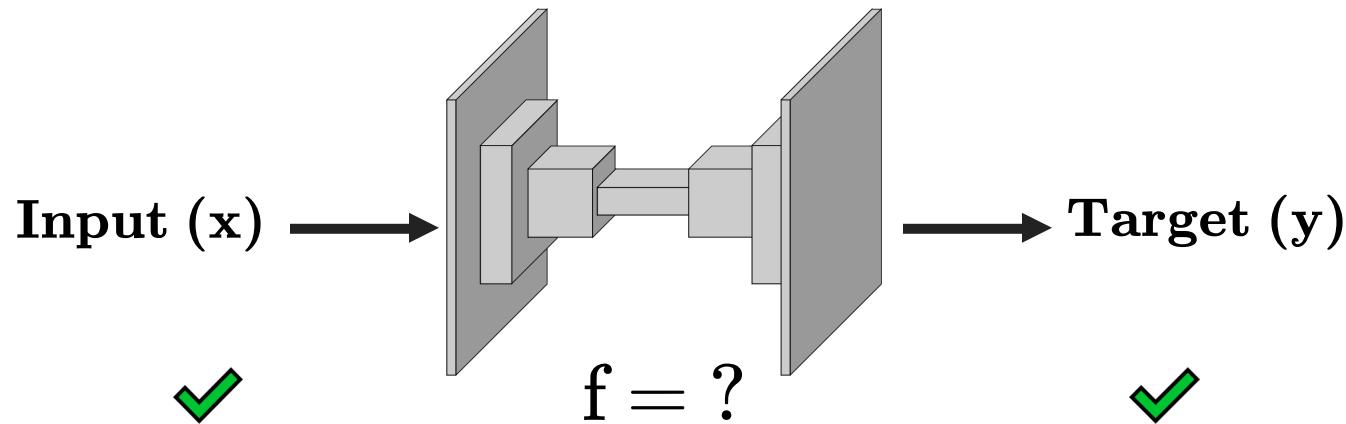
Mini  
Batch

Res-Net

Convolutions

# Loss Functions and Gradient Descent

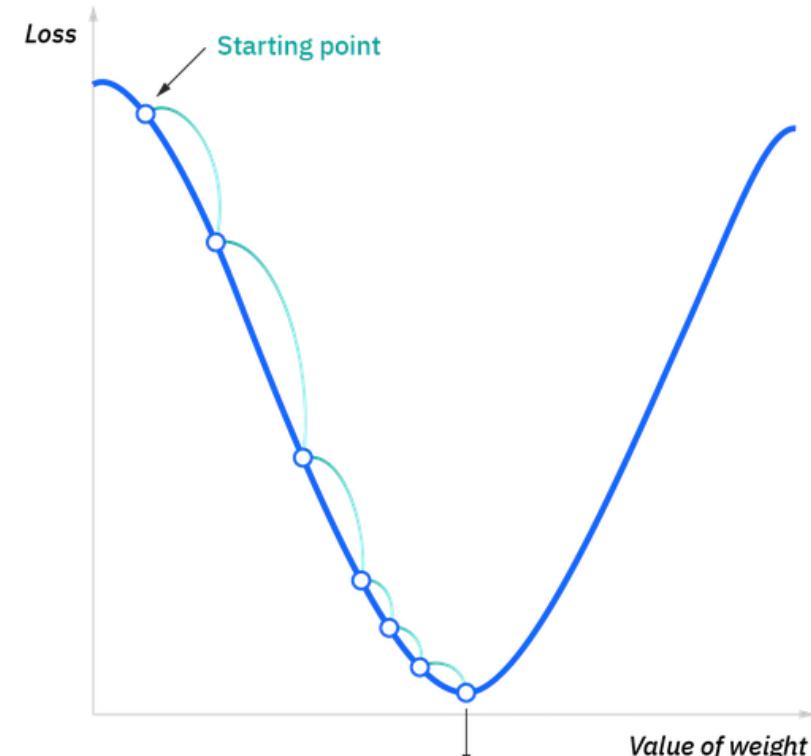
## (Deep) Learning-based Approaches



Training Phase: *Context*-aware learning of model weights or parameters

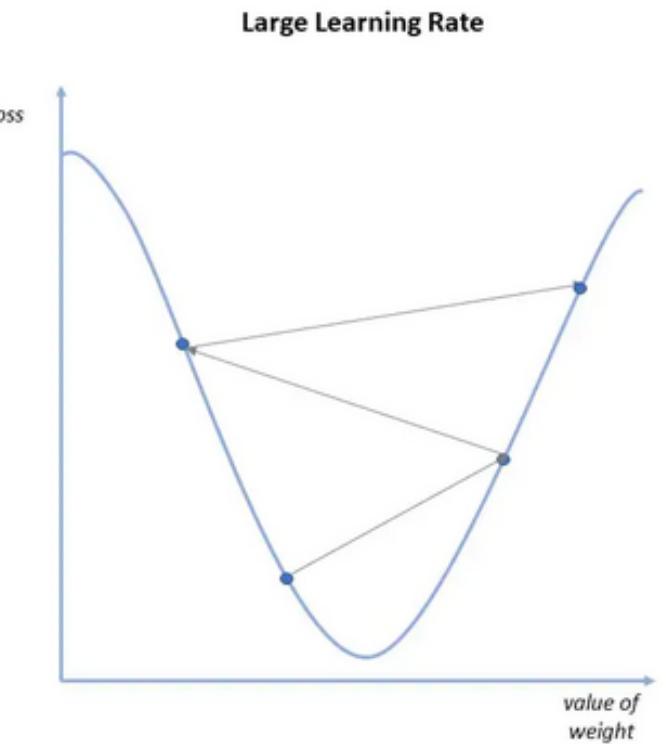
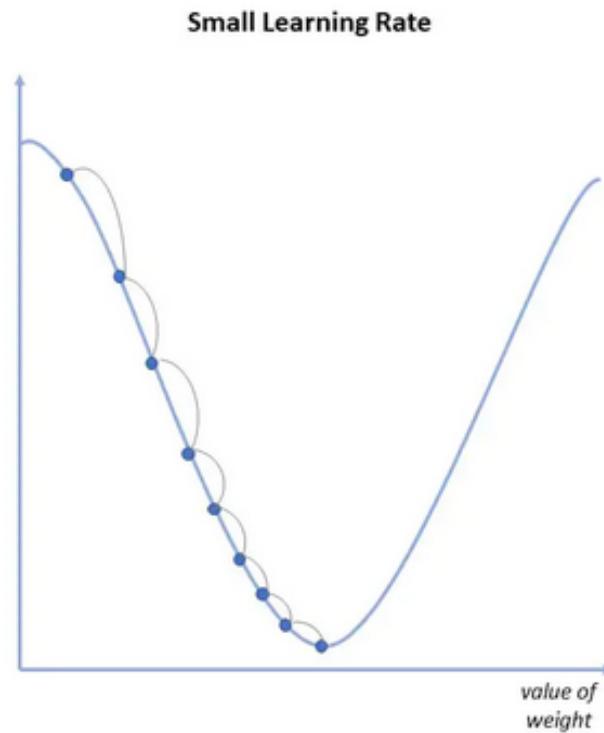
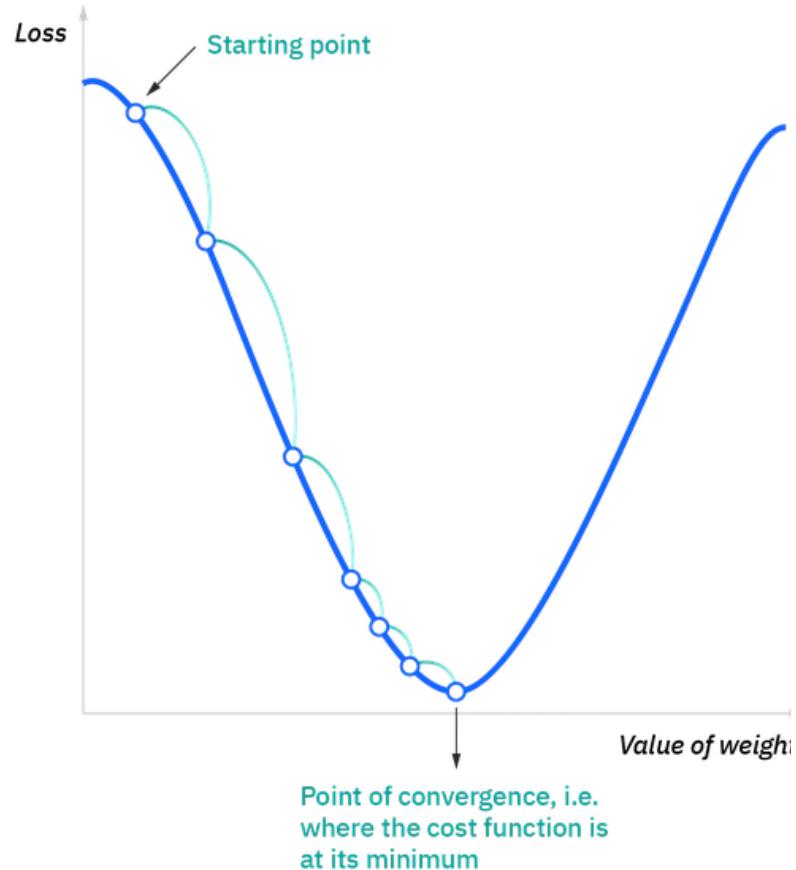
$$L = \frac{1}{M} \sum_i^M \| \hat{f}_\theta(x_i), y_i \|_2$$

$$\theta_{\text{next}} = \theta_{\text{prev}} - \alpha \frac{\partial L}{\partial \theta} |_{\theta_{\text{prev}}}$$



Point of convergence, i.e.  
where the cost function is  
at its minimum

# Gradient Descent, Learning Rate and Scheduling



$$\theta_{\text{next}} = \theta_{\text{prev}} - \alpha \frac{\partial L}{\partial \theta} |_{\theta_{\text{prev}}}$$

Reduce learning rate / take smaller steps as training progresses

# Convolutional Layers enable fewer parameters

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input

1	0	1
0	1	0
1	0	1

Filter / Kernel

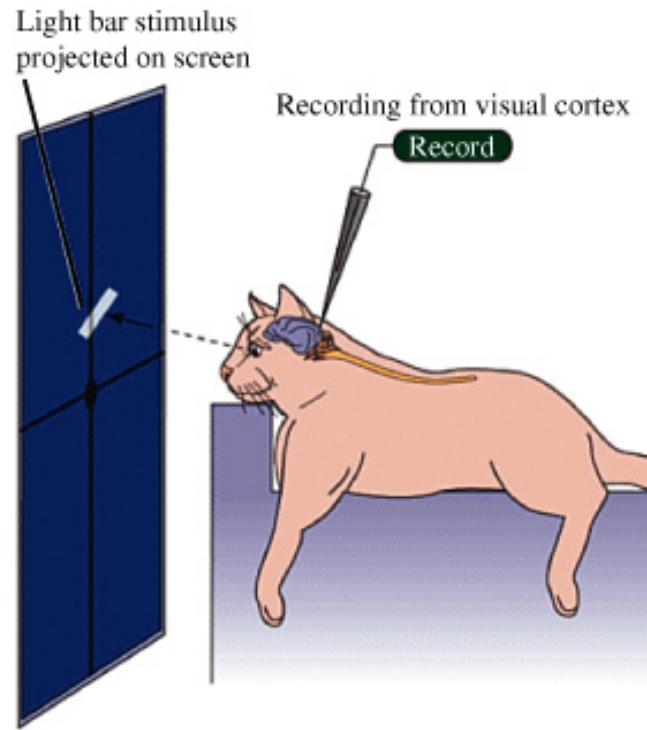
1x1	1x0	1x1	0	0
0x0	1x1	1x0	1	0
0x1	0x0	1x1	1	1
0	0	1	1	0
0	1	1	0	0

4		

These are the weights which are learnt during the training procedure

# Convolutions and Effective Receptive Fields

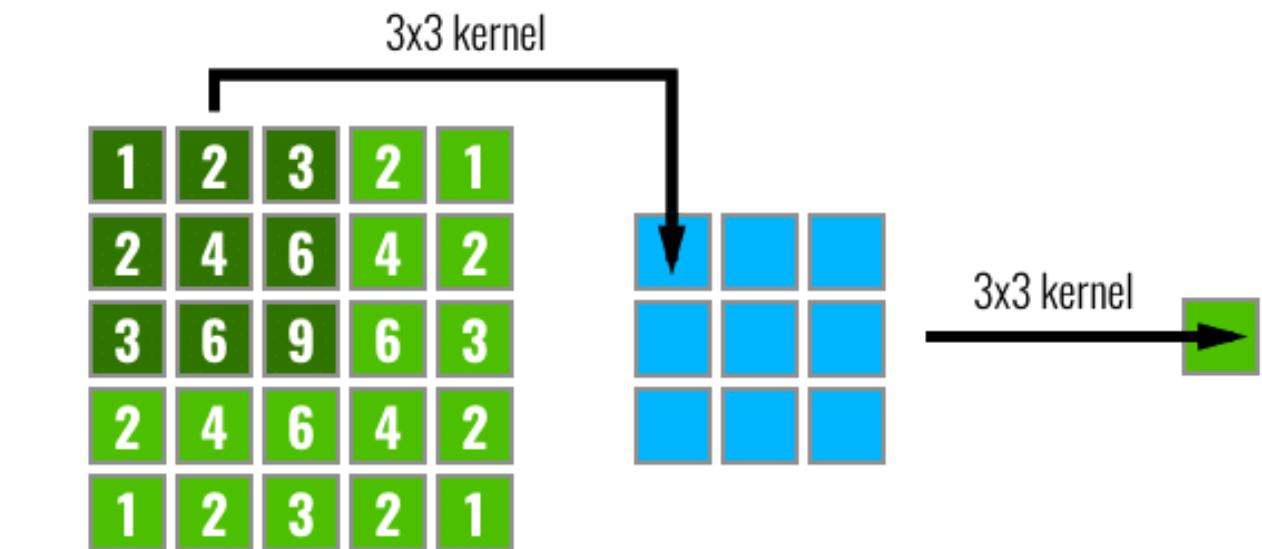
## A Experimental setup



Source: <https://www.informati.com/articles/article.aspx?p=1431818>

*Hubel and Weisel's experiments showed that visual cortical neurons were encoding the properties of objects in the world*

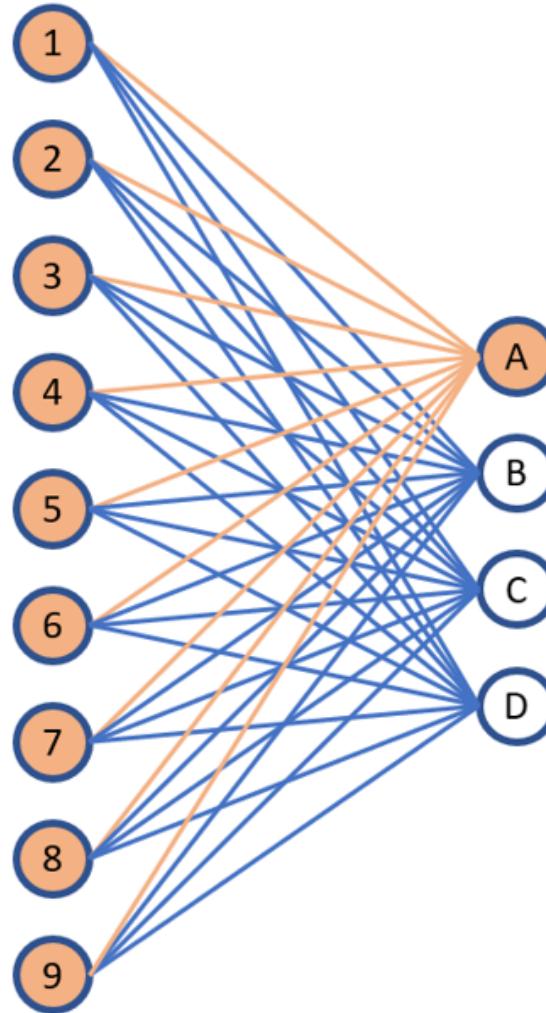
Also, a typical V1 neuron responded to stimulus at specific location  
(defines the neuron's receptive field)



Source: <https://blog.christianperone.com/2017/11/the-effective-receptive-field-on-cnns>

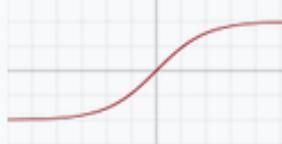
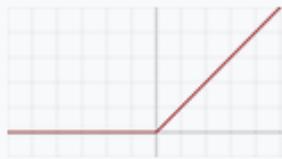
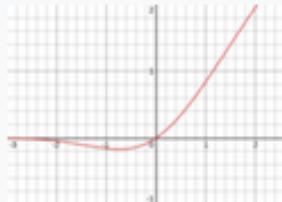
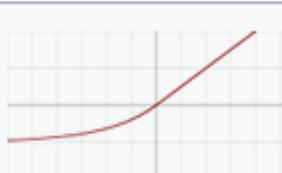
Effective Receptive Field of 5

# Fully Connected Layers require more parameters



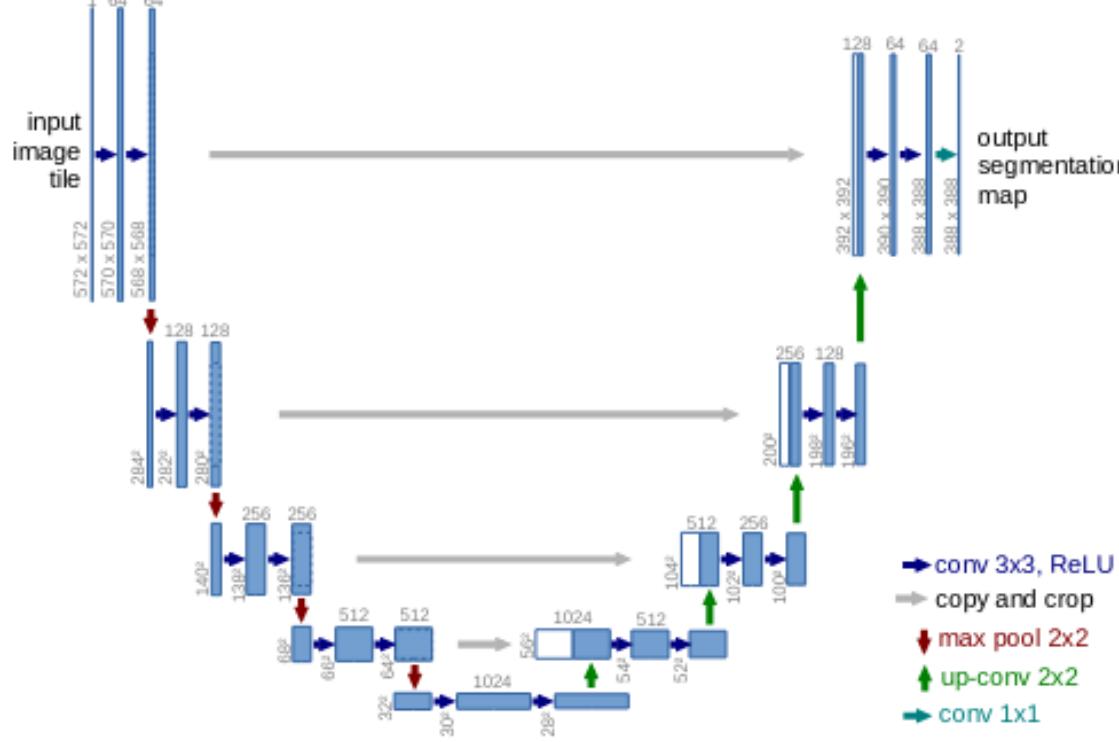
Source: <https://towardsdatascience.com/convolutional-layers-vs-fully-connected-layers-364f05ab460b>

## Between each pair of consecutive layers, Non-Linear Activations are used

Logistic, sigmoid, or soft step		$\sigma(x) = \frac{1}{1 + e^{-x}}$
Hyperbolic tangent (tanh)		$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$
Rectified linear unit (ReLU) <sup>[8]</sup>		$\begin{cases} 0 & \text{if } x \le 0 \\ x & \text{if } x > 0 \end{cases} = \max\{0, x\} = x \mathbf{1}_{x>0}$
Gaussian Error Linear Unit (GELU) <sup>[5]</sup>		$\frac{1}{2}x \left( 1 + \operatorname{erf}\left(\frac{x}{\sqrt{2}}\right) \right) = x\Phi(x)$
Softplus <sup>[9]</sup>		$\ln(1 + e^x)$
Exponential linear unit (ELU) <sup>[10]</sup>		$\begin{cases} \alpha(e^x - 1) & \text{if } x \le 0 \\ x & \text{if } x > 0 \end{cases}$ with parameter $\alpha$

Source: [https://en.wikipedia.org/wiki/Activation\\_function](https://en.wikipedia.org/wiki/Activation_function)

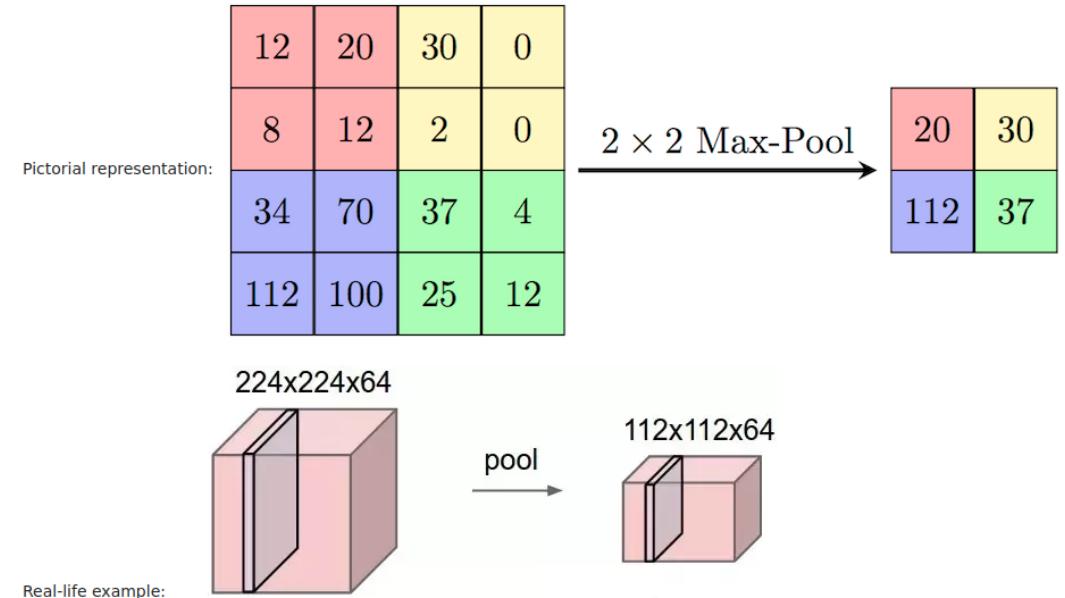
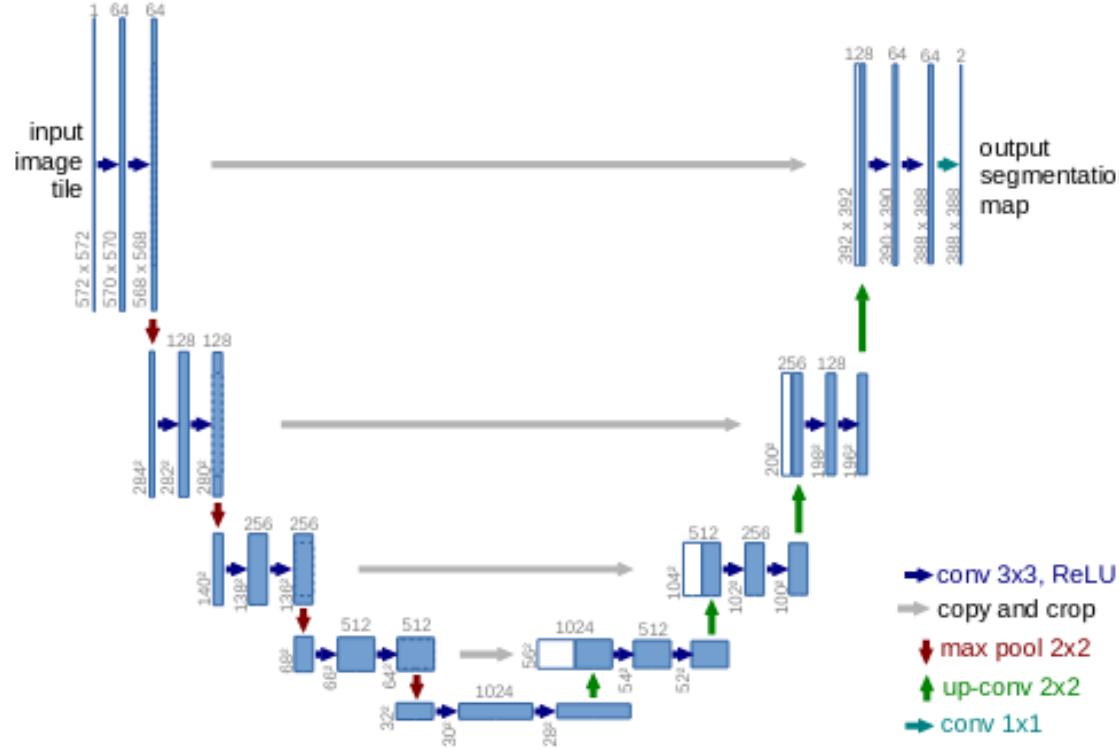
# U-Net is a Fully Convolutional Network



“U-Net: Convolutional Networks for Biomedical Image Segmentation.”, Ronneberger, O. et al, 2015

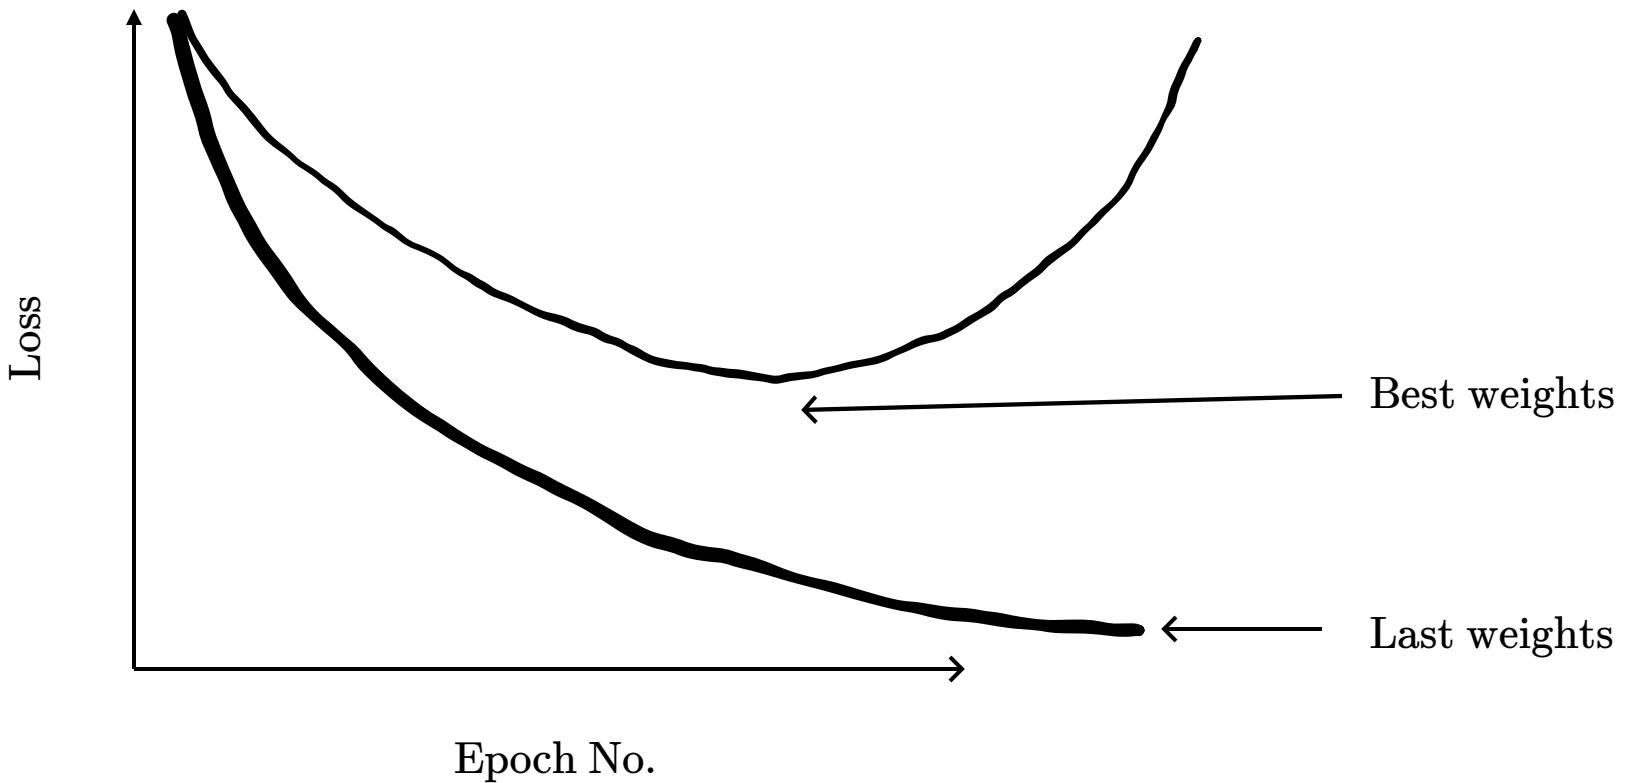
# U-Net is a Fully Convolutional Network

It achieves downsampling through Max Pool and upsampling through Up Convolutions



“U-Net: Convolutional Networks for Biomedical Image Segmentation.”, Ronneberger, O. et al, 2015

# Overfitting and Validation



Training on small dataset can lead to overfitting!

One solution: Use the best performing model state on validation subset instead of last model state

# Supervised vs Self-Supervised Training

## Supervised Training

Both input ( $x$ ) and targets ( $y$ ) are available

Examples include CARE (denoising), StarDist (instance segmentation) etc

## Self-Supervised Training

Only input ( $x$ ) is available. Target ( $y$ ) is constructed through a *proxy* task

Examples in NLP include given a sentence, predict the next word ( $y$ ) in the sentence given all the preceding words ( $x$ )

Examples in CV include given an image, predict a masked pixel given all the neighboring pixels

# Typical Lines of Code in DL training pipelines

```
1 # set device
2 device = torch.device("cuda:0" if args['cuda'] else "cpu")
```

```
1 # train dataloader
2 train_dataset = get_dataset(
3     args['train_dataset']['name'], args['train_dataset']['kwargs'])
4
5 train_dataset_it = torch.utils.data.DataLoader(
6     train_dataset, batch_size=args['train_dataset']['batch_size'], shuffle=True, drop_last=True,
7     num_workers=args['train_dataset']['workers'], pin_memory=True if args['cuda'] else False)
```

```
1 # val dataloader
2 val_dataset = get_dataset(
3     args['val_dataset']['name'], args['val_dataset']['kwargs'])
4
5 val_dataset_it = torch.utils.data.DataLoader(
6     val_dataset, batch_size=args['val_dataset']['batch_size'], shuffle=True, drop_last=True,
7     num_workers=args['train_dataset']['workers'], pin_memory=True if args['cuda'] else False)
```

```
1 # set model
2 model = get_model(args['model']['name'], args['model']['kwargs'])
3 model = torch.nn.DataParallel(model).to(device)
```

```
1 # set optimizer
2 optimizer = torch.optim.Adam(
3     model.parameters(), lr=args['lr'], weight_decay=1e-4)
```

# Typical Lines of Code in DL training pipelines

```
1 # set scheduler
2 scheduler = torch.optim.lr_scheduler.LambdaLR(
3     optimizer, lr_lambda=lambda_, )
```

```
1 # create criterion
2 criterion = nn.CrossEntropyLoss(weight=torch.tensor([1.0, 1.0, 10.0]).cuda())
3 criterion = torch.nn.DataParallel(criterion).to(device)
```

```
1 def train(epoch):
2     model.train()
3     for i, sample in enumerate(tqdm(train_dataset_it)):
4         images = sample['image']
5         semantic_masks = sample['semantic_mask']
6         semantic_masks.squeeze_(1)
7         instance = sample['instance_mask']
8         output = model(images)
9         loss = criterion(output, semantic_masks.long())
10        loss = loss.mean()
11        optimizer.zero_grad()
12        loss.backward()
13        optimizer.step()
14        loss_meter.update(loss.item())
15
16    return loss_meter.avg
```

# Tools can hide away or suggest boilerplate code

WHAT IS PYTORCH LIGHTNING?

Lightning makes coding complex networks simple.

Spend more time on research, less on engineering. It is fully flexible to fit any use case and built on pure PyTorch so there is no need to learn a new language. A quick refactor will allow you to:

- Run your code on any hardware
- Performance & bottleneck profiler
- Model checkpointing
- 16-bit precision
- Run distributed training
- Logging
- Metrics
- Visualization
- Early stopping
- ... and many more!

```
# Lightning Module
import torch
from torch import nn
from torch.nn import functional as F
from torch.utils.data import Dataset
from torch.utils.data import random_split
from torchvision.datasets import MNIST
from torchvision import transforms
import pytorch_lightning as pl

class LitAutoEncoder(pl.LightningModule):
    def __init__(self):
        super().__init__()
        self.encoder = nn.Sequential(
            nn.Linear(28 * 28, 64),
            nn.ReLU(),
            nn.Linear(64, 3))
        self.decoder = nn.Sequential(
            nn.Linear(3, 64),
            nn.ReLU(),
            nn.Linear(64, 28 * 28))

    def forward(self, x):
        embedding = self.encoder(x)
        return embedding

    def configure_optimizers(self):
        optimizer = torch.optim.Adam(self.parameters(), lr=1e-3)
        return optimizer
```

Pytorch Lightning

# Your AI pair programmer

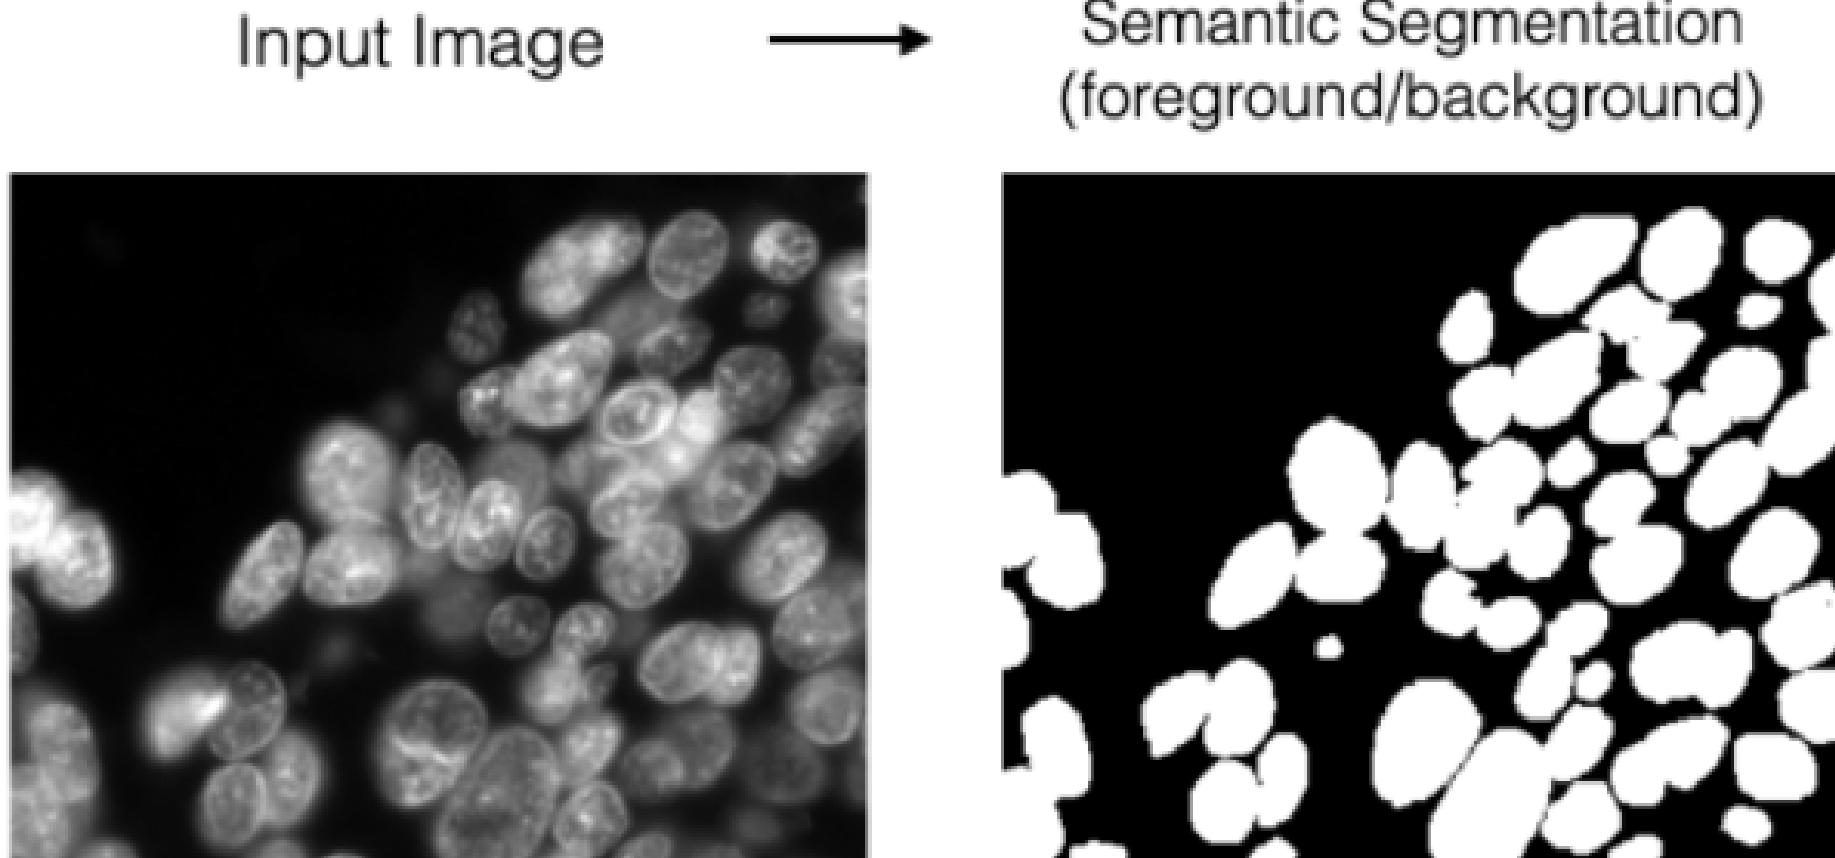
GitHub Copilot uses the OpenAI Codex to suggest code and entire functions in real-time, right from your editor.

[Explore docs >](#)

Github Copilot

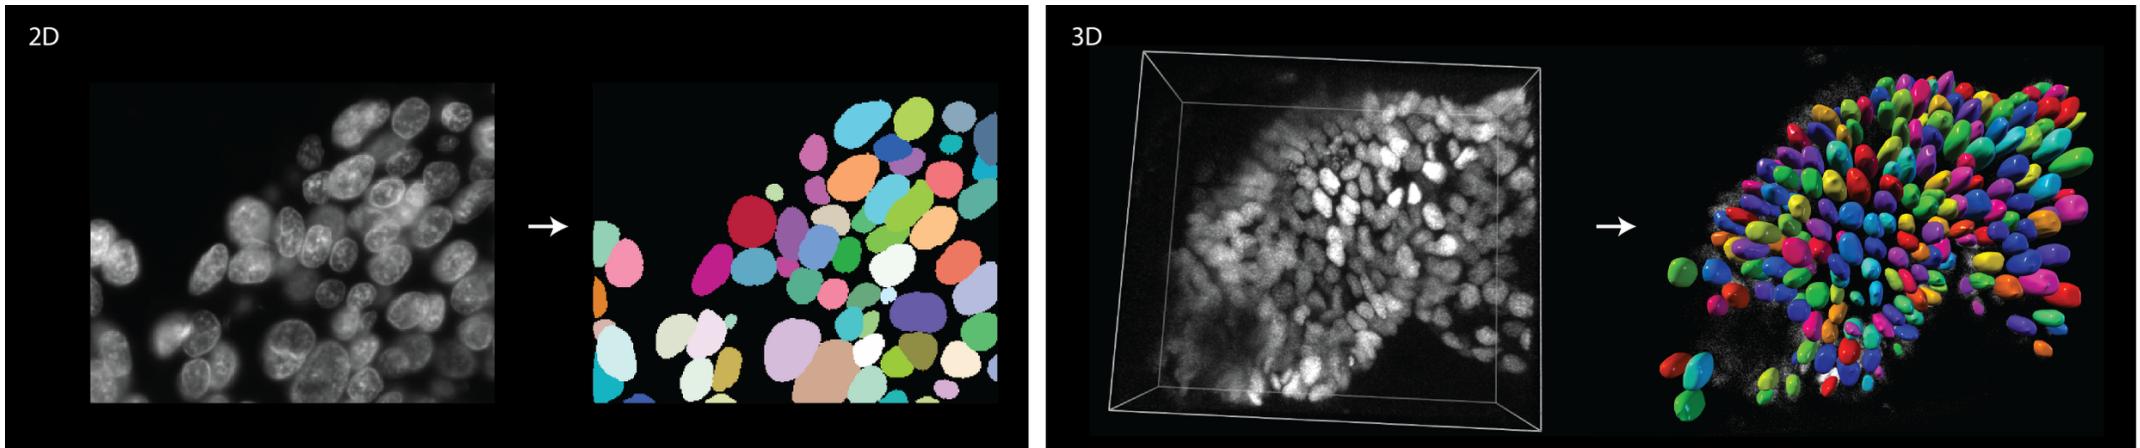
# Deep Learning in Bio-Image Analysis

## Semantic Segmentation



# Deep Learning in Bio-Image Analysis

## Instance Segmentation



“Cell Detection with Star-convex Polygons.”, Schmidt, U. and Weigert, M. et al, 2018

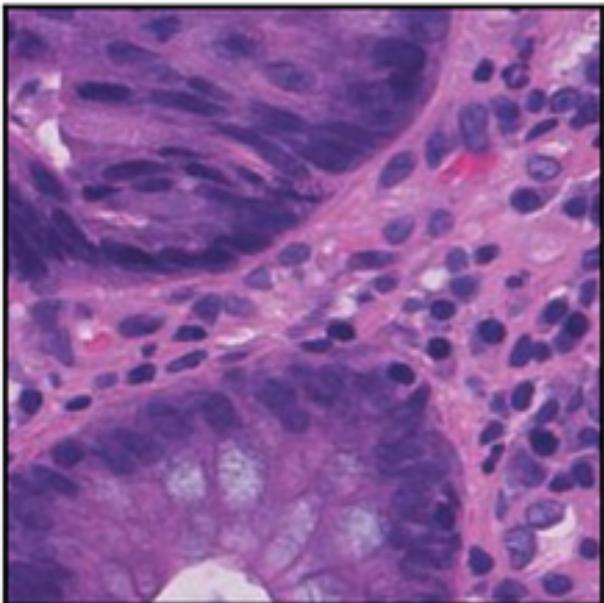
“Star-convex Polyhedra for 3D Object Detection and Segmentation in Microscopy.”, Weigert, M. and Schmidt, U. et al, 2020

“Cellpose: a generalist algorithm for cellular segmentation.”, Stringer, C. et al, 2021

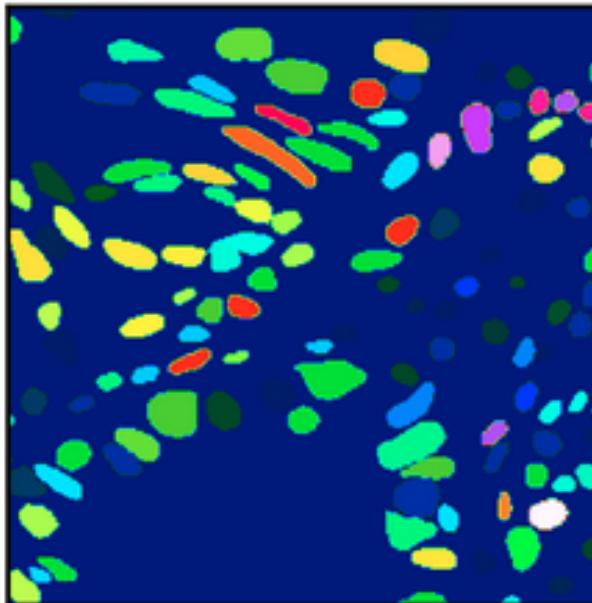
“Accurate and versatile 3D segmentation of plant tissues at cellular resolution.”, Wolny, A. et al, 2020

# Deep Learning in Bio-Image Analysis

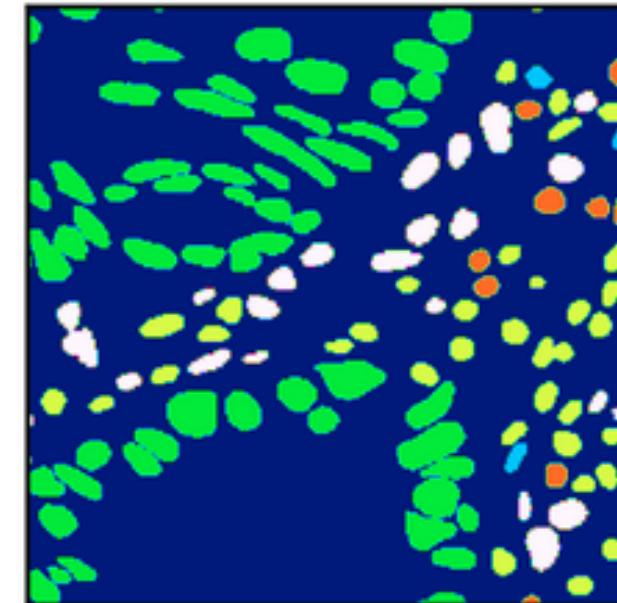
## Joint Instance Segmentation & Classification



Image



Instance map



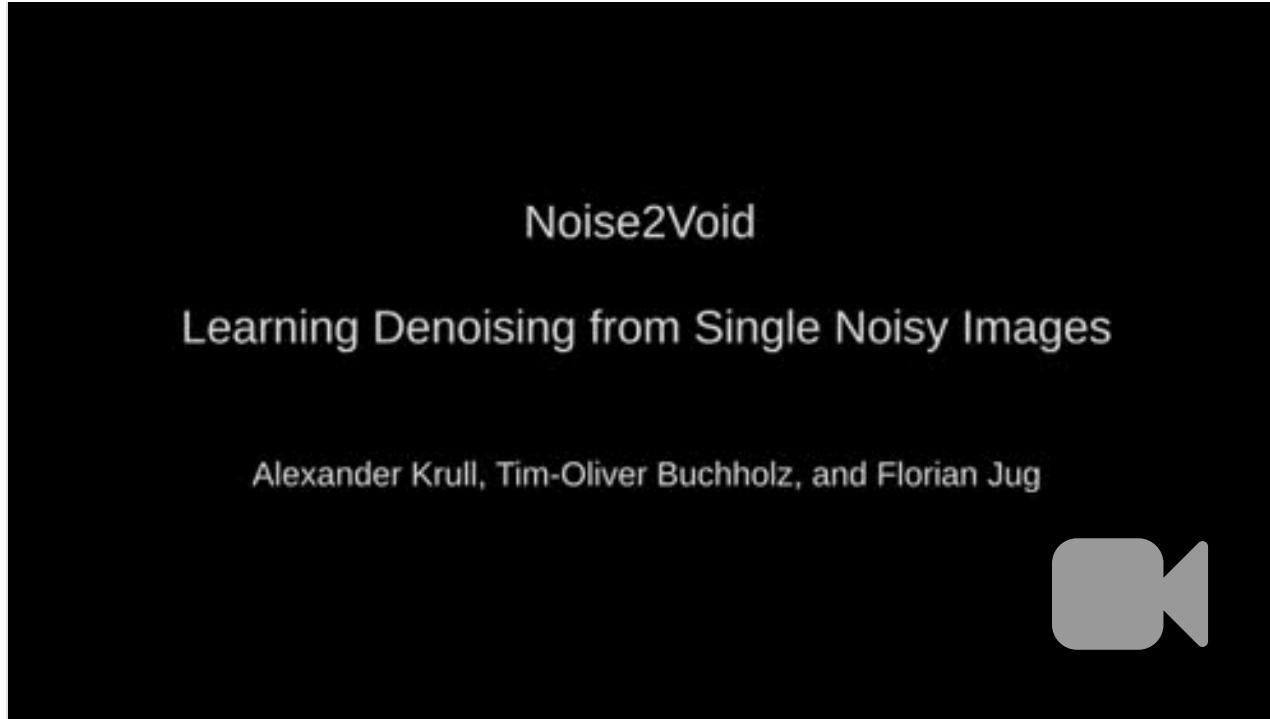
Class map

“Nuclei Instance Segmentation & Classification in Histopathology Images with StarDist”, Weigert, M. and Schmidt, U., 2022

“CoNIC: Colon Nuclei Identification and Counting Challenge 2022.”, Graham, S. et al.

# Deep Learning in Bio-Image Analysis

## Image Denoising and Restoration

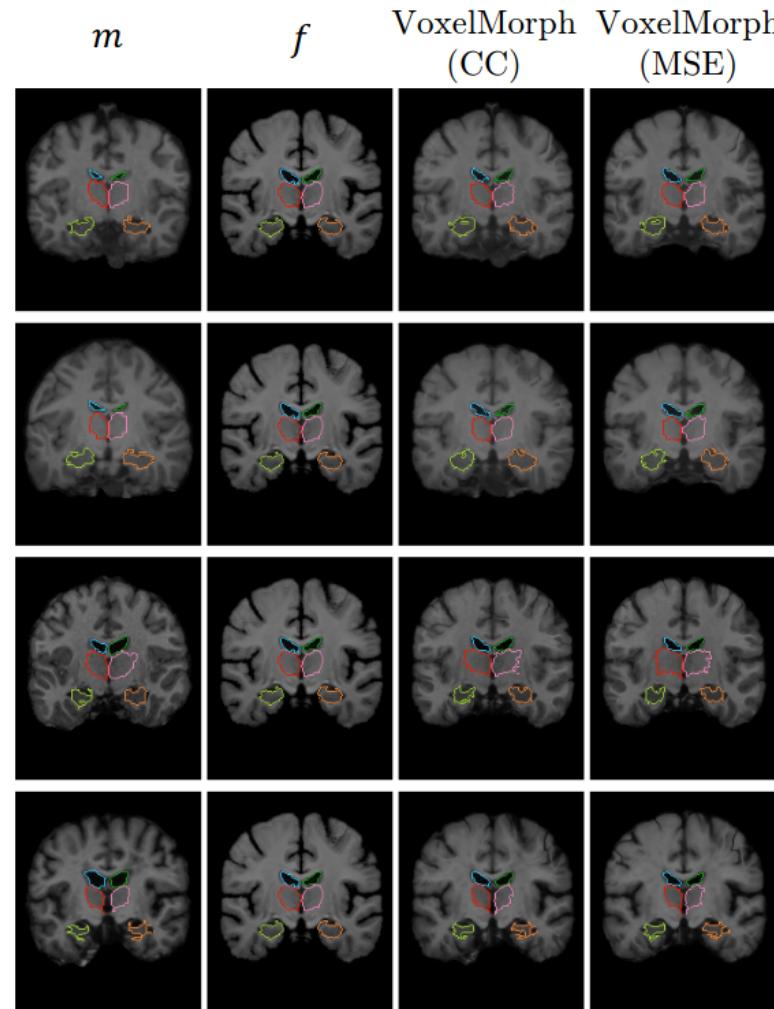


“Noise2Void - Learning Denoising from Single Noisy Images”, Krull, A. and Buchholz, T-O et al, 2018

“Content-aware image restoration: pushing the limits of fluorescence microscopy”, Weigert, M. et al, 2018

# Deep Learning in Bio-Image Analysis

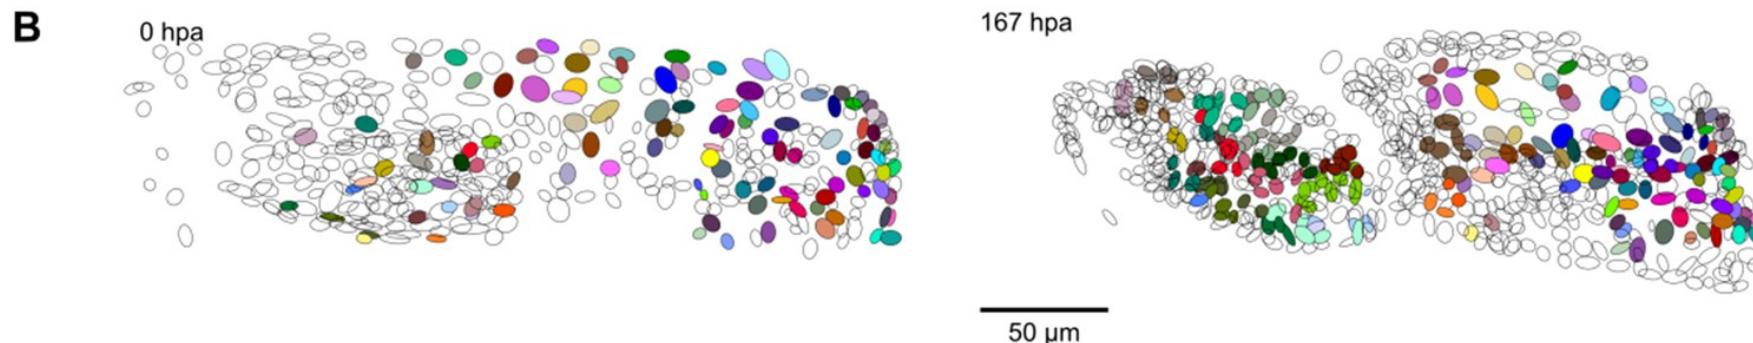
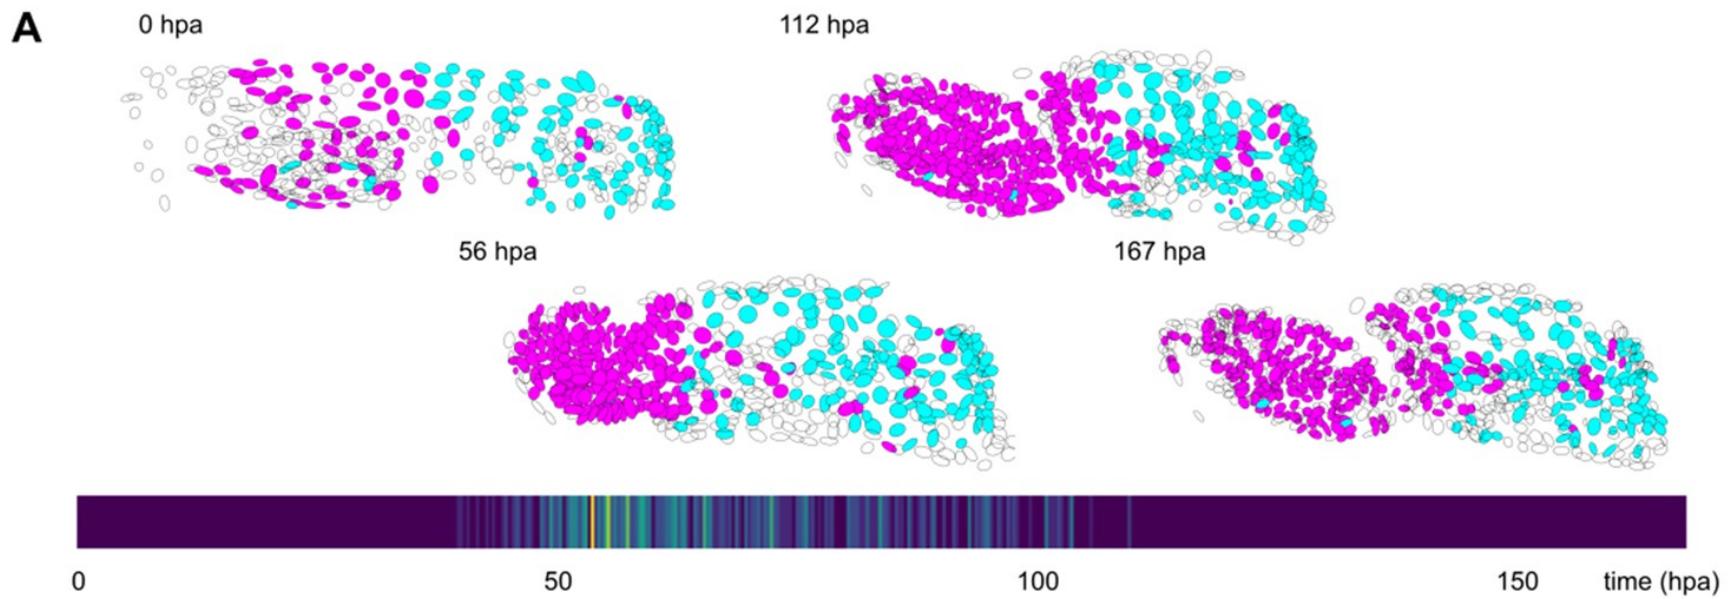
## Image Registration



“VoxelMorph: A Learning Framework for Deformable Medical Image Registration”, Balakrishnan, G. et al, 2018

# Deep Learning in Bio-Image Analysis

## Tracking in Time Lapse Movies



“Tracking cell lineages in 3D by incremental deep learning”, Sugawara, K. et al, 2022

# Exercises for Today

1 - Image Denoising using Noise2Void

2 - Image Semantic Segmentation using Standard U-Nets

